



CS 103 -06

Mid-Term Review and AI Platform Introduction

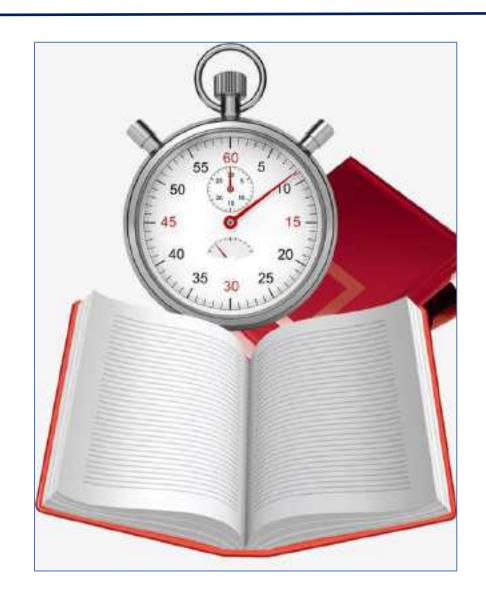
Jimmy Liu 刘江

TA: 章晓庆, 张洋, 肖尊杰, 杨冰

2020-10-23



Lectures Review





Learn and Study

Active learning: It is about how much you think and learn

Collective study: Let us study together



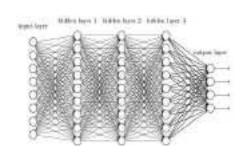
CS 103 Module Coverage

What CS 103 will cover?

Al Concepts Al Algorithms

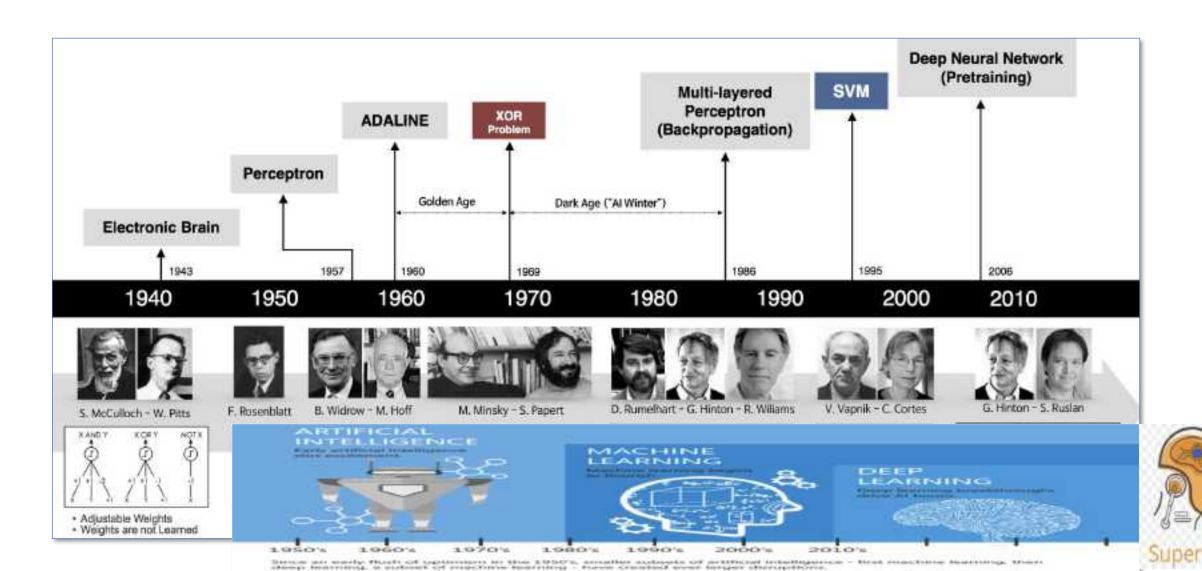
Al Application







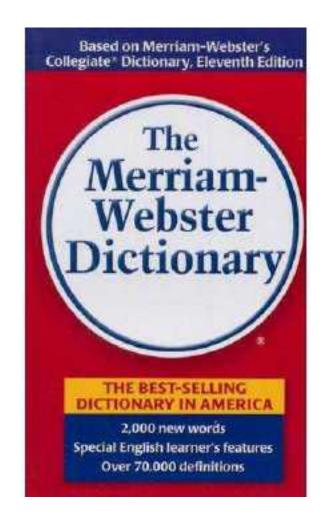
Computer Algorithm and Al algorithm Development Stages and Future Direction





"Intelligence" from Dictionaries

The ability to Learn **Understand** Deal with Try new situations

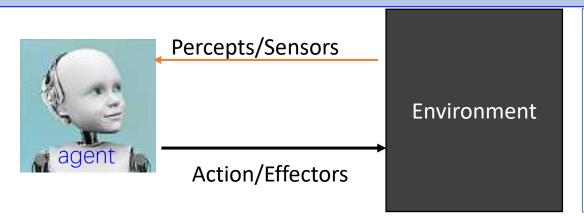




Al from Computer Science - Agent

• An agent is anything that can perceive its environment through sensors and acts upon that environment through effectors. Abstractly, an agent is a function from percept histories to actions:


- A human agent has sensory organs such as eyes, ears, nose, tongue and skin parallel to the sensors, and other organs such as hands, legs, mouth, for effectors.
- A robotic agent replaces cameras and infrared range finders for the sensors, and various motors and actuators for effectors.
- A software agent has encoded bit strings as its programs and actions.

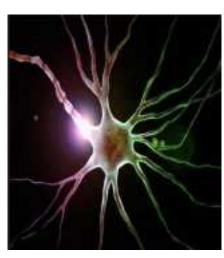


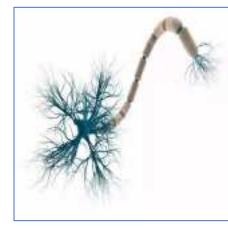




Single Neuron

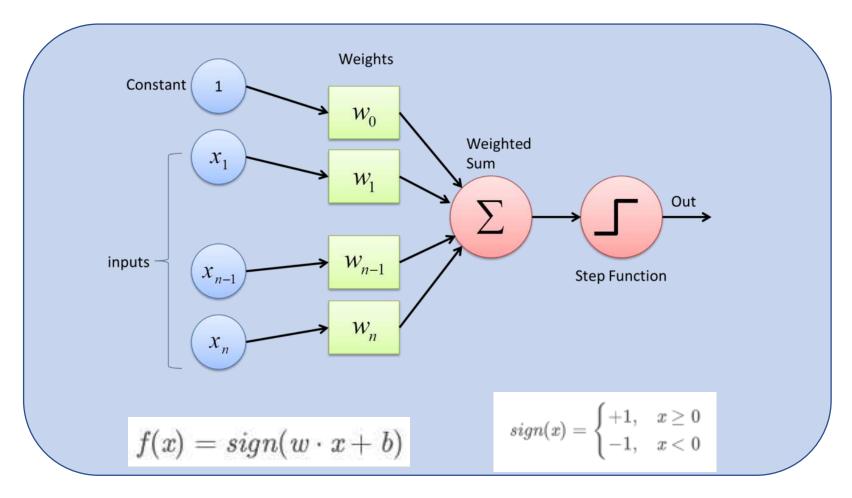









From MCP Neuron to Perceptron







Perceptron Learning Rule 1: PLR

- 1. Randomly choose the weights in the range 0 and 1.
- 2. Training examples are presented to perceptron one by one from the beginning, and its output is observed for each training example.
- 3. If the output is correct then the next training example is presented to perceptron.
- 4. If the output is incorrect then the weights are modified as per the following Perceptron Learning Rule (PLR).

New Wi = Wi +
$$(\eta * Xi * E)$$
.

Change in Weight i = Learning Rate × Current Value of Input i × E (Expected Output, Current Output).

- 5. A simple form of E = (Expected Output Current Output) or SIGN (Expected Output Current Output).
- 6. In PLR, output is 1/0 (or -1), and the transfer is Threshold Step Function



Any Question?





背景:阿尔茨海默症多发于老年群体。起病隐匿,病因尚未明确,且无法治愈。故阿尔茨海默症的早期 诊断尤为重要。

调研方向:

- 1.基于深度学习的AD分类的相关研究
- 2.AD患者的头部MRI形态特征
- 3.成熟的2D图像分类CNN模型及其变体
- 4.使用深度学习基于3D医学影像进行其它疾病诊断的相关研究

目的: 我们希望使用深度学习基于头部MRI(T2w)对AD患者和正常受试者进行分类。以辅助阿尔茨海默症的临床诊断。

组长注: 我看了《基于磁共振图像和改进的UNet++模型区分阿尔茨海默症患者和健康人群》 (doi:10.11938/cjmr20192769) 的引言部分, Introduction还需要总结前人的工作以及提出自己的方法, 但是我们小组目前还在做数据预处理。故暂时无法提供更多内容。

1. Introduction

Gestures are useful in many situations. For example, people communicate with each other conveniently with simple gestures. It is a good way to use gestures to deliver information. Therefore, we can build a gesture recognition system to achieve a Human-computer interaction system with much efficiency and conveniency. In fact, it has been an active research for many years. There're many valuable applications such as, the interaction and communication in VR/AR, smart homes, etc.

We want to construct a system to control computers with simple gestures instead of mouse, which is sometimes inconvenient to use. However, a large portion of previous works require extra hardware such as depth camera[1], motion sensor[2], multiple cameras and etc.. While the model's accuracy is improving, the requirements for hardware are also increasing, making it even more inconvenient for people to use.

Fortunately, a recent study shows that it's possible to get the landmarks of a hand with great accuracy within acceptable time[3]. Thus, the biggest problem now is how to classify the gestures since everyone does one gesture differently. Having located movements of the skeleton, it's still difficult for the model to recognize the gestures because the movements are usually in 3D space, but only 2D pictures are available. In this paper we'll focus on those problems and propose a model that perform in real-time on most devices without any more hardware except for a normal camera.



一、项目名称: 我的五子棋 AI 果然有问题:

二、团队组成:

Leader: 肺样辰**

Followers:陈茜、金冬阳、夏星晨、张坤龙、周野玮、赵云龙-

三、項目目标:

AI 支持五子棋人人对战、人机对战,机机对战。+

包含两种模式。智能模式、智障模式(即最优策略模式与最劣策略模式)。并可以在两种模式之间随意切换。·

四、创意来源

近年來, AlphaGo 在概坛上打遍天下无桩手, 甚至向选军电子竞技行业。人工智能在发展到今天, 人类在竞技体育领域可能越来越不是他们的对于, 所以我们认为可以尝试做一个简单的五子供的 Al。但是。显然光对胜利的渴求并不新颖, 因为人工智能现在越来越多的在各个领域聪明, 从以前的人工智障变成了人工智能, 在去年, 日本一个公司开发了一款人工智能, 号称史上最高人工智能, 这个人工智能在几百万次的游戏对战中只获取了1000次的匪利。无论人类如何放水, 这个人工智能反倒越来越强。于是这个消息便成为了我们组的设计思路, 放弃填有的老畜人工智能思路, 没为设计"人工智障"。两者的差距或许仅仅是调参的一点点区别。但是这个组意总能让大家废他一亮。\$P\$

五、预期使用的工具: -

Python 编程语言

Tensorflow I.A.

CUDA 加速 (暂定, 视情况优化) =

六、预期的算法

费初神经网络设计。cuda 加速实现矩阵运算的优化~

图像着色

CARRACTICADADO COCOMISTACIONALEZA PARA CARRACTURA DE CARRACTURA PARA POR PORCESA PARA SER PARA CARRACTURA DE CARRA

作品を表上を作用を上をがっていた。

・選手を選集を選択できたから選択できたが成功がある。とはご覧覧が必要が選挙を登録された。 で、選手を選択させ続くようできたからは運動をおりが開催しませんがはます。またものから を行うに対し、これをは当時をははあるとはなり、それないのできたとなった。

DO NIE YK. INDERDOMALENERSE MALORMOTERSENER ERMANN DETENLEMELENE 1884 INCOMMODITE MENER BETTAN FRANCISCHE TENNERSENERSENERSENERSEN INDAM BETTAN FRANCISCHE TENNERSEN MENERSEN MENARENTEN MALE MENARENDEM AUCHTENFERNMENN MEN

ATTEMPT TO ANALYSIS (ACTORNO TO ANALYSIS AT AN ARTHUR ATTEMPT AND ANALYSIS ATTEMPT ANAL

BUT *, because the restaurance of the decreasing of all methods are acceptant to the second of the s

RESERVANC, EXITINGENERALISMS, THE STRAIGH, STORE STALL T

あたけ者が世界を定理を行っててやするべまへともにある。 しこのでは発生をみませまへれてか





479/95 Theory 30, Indian's curl lade.

Campus Bus Route Optimization by Artificial Intelligence

FAN GINGYUAN', FANG GIHAN' HE HONGJIE' WANG XIANGCHEN' WU ZIYU' YUAN Place Discount of States and States and

ARSTRACT

INDEX TERMS

L INTRODUCTION

A. BACKBROLND

W ITH the increasing size of large, functionally specific puries, there is an increasing need to move people quickly between two pixins and to transport people within the park and to extend butipett. For this resum, many farce computes have introduced abundle but systems to meet the commuting needs of their internal staff. Usually there shuttle buses have fixed steps in the park and travel back and booth from both ends of the park at regular intervals. throughout the day to transport people. Through multiple interviews and daily observations, we found that there are still some areas that could be improved: the scheduling of bussubedules and dispatching is not automated and will refer on the experience of the drivers towalved, which does not allow for rooce detailed scheduling according to the changing flow of people in the park at different times of the day, resulting in a certain degree of enposity shortage or wanted expands: Employees lack a selection means of knowing the location of bases and their estimated arrival times. At present, staff cononly rely on the departure schedule of bases to entrace the arrival time of hoom, which will easily result in a large want-

W. PRINKAPICH CONTRACT

This project mainly focuses on the bus scheduling operation on campus, which is based on the sudents' class schedule, the general life and rest time of the bus riders, the researcher's research schedule, and the density of the bus tiuffle at a speeffic time, etc. Based on the existing but scheduling system. the project uses the deep tearning and artificial useffigence method to preclas the buy indens' behavior at a specific time.

In this project, the location of campus bases is predicted by the existing schedules of compan personnel, and the logical flow of people. By predicting the estitting schedules of the



FRAME 1. A Specificance to a re-current of SOSTICE

students, we quantify and compare the route of each cumpus bus that circles around the campus, and simulate me location of people based on his day and neural network learning. The eyesem to capable of predicting the number of passangers waiting at a bus stop to improve passenger experience and

At the name time, the scheduling system can inform the busdriver of the stops to be passed by the bus after scheduling is completed in the system, so to as improve the pursenger expendence and compart. Multiple information such as bus location, schedule time and stops can be displayed on the personal terminal to improve the efficiency of his pessengers and achieve the goal of rational optimization of compastus emeralisan

PERFERENCES



我们小组研究方向是AI+虚拟主播,其主要内容是识别收到的文本,将 其转化为信息,再通过语音的形式输出。我们会致力于日常的对话, 使之更加智能化,人性化,以致可以进行基本没有逻辑错误的日常对 话和使用。

1

Start to Write the Introduction of Your Group Project The research topic of our group is the user preferences of big data computing, and the research on big data in an era of large amount of information



Start to Write the Introduction of Your Group Project

The whole project is based on the MVC model, which consist four parts Front end, control, business logic and data base

Outline









Python



Scikit-learn



Pytorch



Tensorflow

Background



- **Python** is an interpreted, high-level and general-purpose programming language. Created by Guido van Rossum and first released in 1991. Python's design philosophy emphasizes code readability with its notable use of significant whitespace. Its language constructs and object-oriented approach aim to help programmers write clear, logical code for large-scale projects.
- Python 3.0, released in 2008, was a major revision of the language that is not completely backward-compatible, and <u>much Python 2 code does not run unmodified on Python 3</u>.
- The Python 2 language was officially discontinued in 2020 (first planned for 2015), and "Python 2.7.18 is the last Python 2.7 release and therefore the last Python 2 release." No more security patches or other improvements will be released for it. With Python 2's end-of-life, only Python 3.6.x and later are supported.

Background



Basic Python construct:

Focused on Python 3.0

Help you build a foundation

Get ready to follow online tutorials https://www.runoob.com/python/python-tutorial.html

Background



Python is an interpreted, high-level, object-oriented programming language.

- Flexible programming language
- Designed to be human readable Related applications:
- Machine learning models
- Artificial intelligence projects
- Web applications



print displays output to your console

print('Hello world')

Hello world



Enclose strings in single or double quotes

```
print('Hello world single quotes')
print("Hello world double quotes")
```

Hello world single quotes Hello world double quotes



Getting information from the user

```
name = input('Please enter your name: ')
print(name)
```

Please enter your name: Susan Susan



Printing blank lines can improve readability

```
print('Hello world')
print()
print('Did you see that blank line?')
print('Blank line \nin the middle of string')
```

```
Hello world

Did you see that blank line?

Blank line

in the middle of string
```

Debug



Debugging with print

```
print('Adding numbers')
x = 42 + 206
print('Performing division')
y = x / 0
print('Math complete')
```

```
Adding numbers

Performing division

Traceback (most recent call last):

File "demo.py", line 4, in <module>

y = x / 0

ZeroDivisionError: float division by zero
```

Comment (注释)



 Comments document your code so you and other programmers can understand the code

```
# Using double quotes for this string because # the string itself contains a single quote print("It's a small world after all")
```

It's a small world after all



• Strings can be stored in variables

```
first_name = 'Susan'
print(first_name)
```

Susan



You can combine strings with +

```
first_name = 'Susan'
last_name = 'Ibach'
print(first_name + last_name)
print('Hello' + first_name + '' + last_name)
```

```
SusanIbach
Hello Susan Ibach
```



You can use functions to modify strings

```
sentence = 'The dog is named Sammy'
print(sentence.upper())
print(sentence.lower())
print(sentence.capitalize())
print(sentence.count('a'))
```

```
THE DOG IS NAMED SAMMY
the dog is named sammy
The dog is named sammy
2
```



Numbers can be stored in variables

```
pi = 3.14159
print(pi)
```

3.14159



You can do math with numbers

```
first_num = 6
second_num = 2
print(first_num + second_num)
print(first_num ** second_num)
```

Symbol	Operation
+	Addition
-	Subtraction
*	Multiplication
/	Division
**	Exponent

8

36



- If you combine strings with numbers, Python gets confused
- When displaying a string that contains numbers you must convert the numbers into strings

```
days in feb = 28
print(days in feb + ' days in February')
days in feb = 28
print(str(days in feb) + ' days in February')
File "February.py", line 2, in <module>
    print(days_in_feb + ' days in February')
TypeError: unsupported operand type(s) for +: 'int' and 'str'
28 days in February
```

Date



 We often need current date and time when logging errors and saving data

```
# To get current date and time
# we need to use the datetime library
from datetime import datetime

current_date = datetime.now()
# the now function returns a datetime object
print('Today is: ' + str(current_date))
```

Today is: 2019-06-06 16:17:18.694511

Date



 There are functions you can use with datetime objects to manipulate dates

```
from datetime import datetime, timedelta
today = datetime.now()
print('Today is: ' + str(today))

# timedelta is used to define a period of time
one_day = timedelta(days=1)
yesterday = today - one_day
print('Yesterday was: ' + str(yesterday))
```

```
Today is: 2019-06-06 16:14:24.615495
Yesterday was: 2019-06-05 16:14:24.615495
```

Date



Use date functions to control date formatting

```
from datetime import datetime
current_date = datetime.now()

print('Day: ' + str(current_date.day))
print('Month: ' + str(current_date.month))
print('Year: ' + str(current_date.year))
```

Day: 6 Month: 6

Year: 2019

Error



Syntax errors

```
# This code won't run at all
x = 42
y = 206
if x == y
    print('Success!!')
```

SyntaxError: invalid syntax

Error



Runtime errors

```
# This code will fail when run
x = 42
y = 0
print(x / y)
```

```
Traceback (most recent call last):
   File "runtime.py", line 3, in <module>
     print(x / y)
ZeroDivisionError: division by zero
```

Error



Catching runtime errors

```
try:
    print(x / y)
except ZeroDivisionError as e:
    # Optionally, log e somewhere
    print('Sorry, something went wrong')
except:
    print('Something really went wrong')
finally:
    print('This always runs on success or failure')
```

Sorry, something went wrong



 Your code needs the ability to take different actions based on different conditions

```
if price >= 1.00:
    tax = .07
    print(tax)
```

Symbol	Operation
>	Greater than
<	Less than
>=	Greater than or equal to
<=	Less than or equal to
==	is equal to
!=	is not equal to





You can add a default action using else

```
if price >= 1.00:
    tax = .07
    print(tax)
else:
    tax = 0
    print(tax)
```



Be careful when comparing strings

```
country = 'CANADA'
if country == 'canada':
    print('Oh look a Canadian')
else:
    print('You are not from Canada')
String comparisons
are case sensitive
are case sensitive
country == 'CANADA'
String comparisons
are case sensitive
are case sensitive
case sensitiv
```

You are not from Canada



Use string functions to make case insensitive comparisons

```
country = 'CANADA'
if country.lower() == 'canada':
    print('Oh look a Canadian')
else:
    print('You are not from Canada')
```



You may need to check multiple conditions to determine the correct action


```
if province == 'Alberta':
    tax = 0.05
if province == 'Nunavut':
    tax = 0.05
if province == 'Ontario':
    tax = 0.13
```





 If only one of the conditions will ever occur you can use a single if statement with elif


```
if province == 'Alberta':
    tax = 0.05
elif province == 'Nunavut':
    tax = 0.05
elif province == 'Ontario':
    tax = 0.13
```





 When you use elif instead of multiple if statements you can add a default action


```
if province == 'Alberta':
    tax = 0.05
elif province == 'Nunavut':
    tax = 0.05
elif province == 'Ontario':
    tax = 0.13
else:
    tax = 0.15
```






 If multiple conditions cause the same action they can be combined into a single condition

```
if province == 'Alberta' \
    or province == 'Nunavut':
    tax = 0.05
elif province == 'Ontario':
    tax = 0.13
else:
    tax = 0.15
```






If you have a list of possible values to check, you can use the IN operator





 If an action depends on a combination of conditions you can nest if statements

```
if country == 'Canada':
   if province in('Alberta',\
       'Nunavut','Yukon'):
       tax = 0.05
   elif province == 'Ontario':
       tax = 0.13
    else:
       tax = 0.15
    tax = 0.0
```





 Sometimes you can combine conditions with AND instead of nesting if statements

```
if gpa >= .85:
    if lowest_grade >= .70:
        print('Well done')
```

```
if gpa >= .85 and lowest_grade >= .70:
    print('Well done')
```

Requirements for honour roll

- Minimum 85% grade point average
- Lowest grade is at least 70%



 If you need to remember the results of a condition check later in your code, use Boolean variables as flags

```
if gpa >= .85 and lowest_grade >= .70:
    honour_roll = True
else:
    honour_roll = False
# Somewhere later in your code
if honour_roll:
    print('Well done')
```

Med

Collections

Lists are collections of items

```
names = ['Christopher', 'Susan']
scores = []
scores.append(98) # Add new item to the end
scores.append(99)
print(names)
print(scores)
print(scores[1]) # Collections are zero-indexed
```

```
['Christopher', 'Susan']
[98, 99]
99
```



Arrays are also collections of items

```
from array import array
scores = array('d')
scores.append(97)
scores.append(98)
print(scores)
print(scores[1])
```

```
array('d', [97.0, 98.0])
98.0
```



Common operations

```
names = ['Susan', 'Christopher']
print(len(names)) # Get the number of items
names.insert(0, 'Bill') # Insert before index
print(names)
names.sort()
print(names)
```

```
2
['Bill', 'Susan', 'Christopher']
['Bill', 'Christopher', 'Susan']
```



Retrieving ranges

```
names = ['Susan', 'Christopher', 'Bill']
presenters = names[0:2] # Get the first two items
# Starting index and number of items to retrieve
print(names)
print(presenters)
```

```
['Susan', 'Christopher', 'Bill']
['Susan', 'Christopher']
```



Dictionaries

```
person = {'first': 'Christopher'}
person['last'] = 'Harrison'
print(person)
print(person['first'])
```

```
{'first': 'Christopher', 'last': 'Harrison'}
Christopher
```

Loop



Loop through a collection

```
for name in ['Christopher', 'Susan']: print(name)
```

Christopher Susan



Loop

Looping a number of times

```
for index in range(0, 2):
print(index)
```



Loop

 Looping with a condition names = ['Christopher', 'Susan'] index = 0 while index < len(names): print(names[index]) # Change the condition!! index = index + 1

Christopher Susan





Sometimes we copy and paste our code

```
import datetime
# print timestamps to see how long sections of code
# take to run
first_name = 'Susan'
print('task completed')
print(datetime.datetime.now())
print()
for x in range(0,10):
    print(x)
print('task completed')
print(datetime.datetime.now())
print()
```

```
task completed
2019-05-30 16:55:01.815327
task completed
2019-05-30 16:55:01.817263
```





Use functions instead of repeating code

```
import datetime
# Print the current time
def print_time():
    print('task completed')
    print(datetime.datetime.now())
    print()
first_name = 'Susan'
print_time()
for x in range(0,10):
    print(x)
print_time()
```

```
task completed
2019-05-30 16:55:45.397319
0
8
task completed
2019-05-30 16:55:45.399314
```



_oop

 By moving the code to a function, you reduce rework and the chance of introducing bugs when you change the code you had copied

```
# Import the datetime class from datetime
library
from datetime import datetime
# Print the current time
def print time():
   print('task completed')
   # Now I don't need the extra datetime
prefix
   print(datetime.now())
   print()
```





What if I want a different message displayed?

```
from datetime import datetime
# print timestamps to see how long sections of code
# take to run
first_name = 'Susan'
print('first name assigned')
print(datetime.now())
print()
for x in range(0,10):
    print(x)
print('loop completed')
print(datetime.now())
print()
```

```
first name assigned
2019-05-31 10:18:53.419754
0
loop completed
2019-05-31 10:18:53.422748
```





Pass the task name as a parameter

```
from datetime import datetime
# Print the current time and task name
def print_time(task_name):
    print(task_name)
    print(datetime.now())
    print()
first_name = 'Susan'
print_time('first name assigned')
for x in range(0,10):
  print(x)
print_time('loop completed')
```

```
first name assigned
2019-05-31 10:18:53.419754
0
loop completed
2019-05-31 10:18:53.422748
```



 Here's another example where the code looks different but we are doing the same logic over and over

```
Enter your first name: Susan
Enter your last name: Ibach
Your initials are: SI
```



• I can still use a function, but this time my function returns a value

```
Enter your first name: susan
Enter your last name: ibach
Your initials are: si
```



 If you need to change something you only have to change it in one place!

```
def get_initial(name):
    initial = name[0:1].upper()
    return initial

first_name = input('Enter your first name: ')
first_name_initial = get_initial(first_name)

last_name = input('Enter your last name: ')
last_name_initial = get_initial(last_name)
```

```
Enter your first name: susan
Enter your last name: ibach
Your initials are: SI
```



 Functions that return values allow clever code, but you might trade readability for less code

```
def get_initial(name):
    initial = name[0:1].upper()
    return initial

first_name = input('Enter your first name: ')
last_name = input('Enter your last name: ')

print('Your initials are: '\
    + get_initial(first_name) \
    + get_initial(last_name))
```

```
Enter your first name: susan
Enter your last name: ibach
Your initials are: SI
```



 We already learned to create functions which accept a parameter and return values

```
def get_initial(name):
    initial = name[0:1].upper()
    return initial

first_name = input('Enter your first name: ')
first_name_initial = get_initial(first_name)

print('Your initial is: ' + first_name_initial)
```

```
Enter your first name: adam
Your initial is: A
```

Function



Functions can accept multiple parameters

```
def get initial (name, force_uppercase):
  if force_uppercase:
    initial = name[0:1].upper()
  else:
    initial = name[0:1]
  return initial
first_name = input('Enter your first name: ')
first_name_initial = get_initial(first_name, False)
print('Your initial is: ' + first name initial)
```

Pass the parameters in the same order they are listed in the function declaration

Enter your first name: adam
Your initial is: a

Function



You can specify a default value for a parameter

```
def get_initial(name, force_uppercase=True):
    if force_uppercase:
        initial = name[0:1].upper()
    else:
        initial = name[0:1]
    return initial

first_name = input('Enter your first name: ')
first_name_initial = get_initial(first_name)

print('Your initial is: ' + first_name_initial)
```

```
Enter your first name: adam
Your initial is: A
```

Function



You can also assign the values to parameters by name when you call the function

```
Enter your first name: adam
Your initial is: A
```

Scikit-learn





What is scikit-learn?


Scikit-learn is a Python module for **machine learning** built on top of SciPy and is distributed under the 3-Clause BSD license. The project was started in 2007 by David Cournapeau as a Google Summer of Code project, and since then many volunteers have contributed. (https://github.com/scikit-learn)

Scikit-learn (formerly **scikits.learn** and also known as **sklearn**) is a <u>free</u> <u>software machine learning library</u> for the <u>Python programming language</u>. It features various <u>classification</u>, <u>regression</u> and <u>clustering</u> algorithms including <u>support vector machines</u>, <u>random forests</u>, <u>gradient boosting</u>, <u>k-means</u> and <u>DBSCAN</u>, and is designed to interoperate with the Python numerical and scientific libraries <u>NumPy</u> and <u>SciPy</u>. (**Wikipedia**)



Cearn Scikit-learn







Dimensionality reduction Reducing the number of random variables to consider. Applications: Visualization, Increased efficiency Algorithms: in-Means, Feature selection, non-negative matrix factorization, and more.

Down ples

Regression Predicting a continuous-valued attribute associated with an object. Applications: Drug response, Stock prices. Algorithms: SVR, hearest rieigibors, random forest, and more...











Installing scikit-learn

There are different ways to install scikit-learn:

- •<u>Install the latest official release</u>. This is the best approach for most users. It will provide a stable version and pre-built packages are available for most platforms.
- •Install the version of scikit-learn provided by your <u>operating system or Python distribution</u>. This is a quick option for those who have operating systems or Python distributions that distribute scikit-learn. It might not provide the latest release version.
- •<u>Building the package from source</u>. This is best for users who want the latest-and-greatest features and aren't afraid of running brand-new code. This is also needed for users who wish to contribute to the project.

Installing the latest release Operating System Install the 64bit version of Python 3, for instance from https://www.python.org \$ pip install -O scikit-luarn In order to check your installation you can use 5 orthor -m pig show scikit-learn # to see which version and where scikit-learn is installed \$ gython -m pip freeze # to see all packages installed in the active virtualeny \$ gythom ~c "import sklearn; sklearn;show versions()"

https://scikit-learn.org/stable/install.html





Installing scikit-learn

There are different ways to install scikit-learn:

- •<u>Install the latest official release</u>. This is the best approach for most users. It will provide a stable version and pre-built packages are available for most platforms.
- •Install the version of scikit-learn provided by your <u>operating system or Python distribution</u>. This is a quick option for those who have operating systems or Python distributions tha distribute scikit-learn. It might not provide the latest release version.
- •<u>Building the package from source</u>. This is best for users who want the latest-and-greatest features and aren't afraid of running brand-new code. This is also needed for users who wish to contribute to the project.



URL: https://scikit-learn.org/stable/install.html





Dependencies

scikit-learn requires:

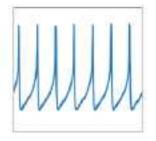
- Python
- •NumPy
- •SciPy
- •joblib
- •threadpoolctl
- Matplotlib
- Seaborn
- Jupyter


http://seaborn.pydata.org/index.html


https://matplotlib.org/

Matplotlib: Visualization with Python

Matplotlib is a comprehensive library for creating static, animated, and interactive visualizations in Python.

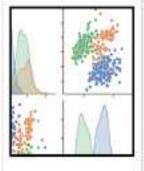


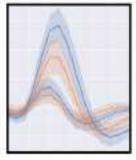


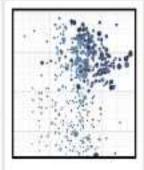


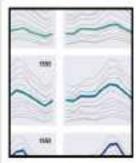



seaborn: statistical data visualization


















User Guide

- 1. Supervised learning
- 2. Unsupervised learning
- Model selection and evaluation
- 4. Inspection
- 5. Visualizations
- 6. Dataset transformations
- 7. Dataset loading utilities
- 8. Computing with scikit-learn

https://scikit-learn.org/stable/user_guide.html

3. Model selection and evaluation

3.1. Cross-validation: evaluating estimator performance

- 3.1.1 Computing cross-validated metrics
- · 3.1.2. Cross validation iterators
- . 3.1.3. A note on shuffling
- · 3.1.4. Cross validation and model selection

3.2. Tuning the hyper-parameters of an estimator

- . 3.2.1, Exhaustive Grid Search
- 3.2.2. Randomized Parameter Optimization
- 3.2.3. Tips for parameter search
- · 3.2.4. Alternatives to brute force parameter search

3.3. Metrics and scoring: quantifying the quality of predictions

- · 3.3.1. The exercise parameter: defining model evaluation rules
- 3.3.2. Classification metrics.
- 3.3.3 Multilabel ranking metrics
- · 3.3.4. Regression metrics
- 3.3.5. Qustering metrics
- 3.3.6. Dummy estimators

3.4. Model persistence

- · 3.4.1. Persistence example.
- 3.4.2. Security & maintainability limitations

3.5. Validation curves: plotting scores to evaluate models

- · 3.5.1. Validation curve
- 3.5.2. Learning curve





User Guide

- 1. Supervised learning
- 2. Unsupervised learning
- Model selection and evaluation
- 4. Inspection
- 5. Visualizations
- 6. Dataset transformations
- 7. Dataset loading utilities
- 8. Computing with scikit-learn

https://scikit-learn.org/stable/user_guide.html

3. Model selection and evaluation

3.1. Cross-validation: evaluating estimator performance

- 3.1.1 Computing cross-validated metrics
- 3.1.2 Cross validation iterators
- . 3.1.3. A note on shuffling
- 3.1.4 Cross validation and model selection

3.2. Tuning the hyper-parameters of an estimator

- . 3.2.1. Exhaustive Grid Search
- 3.2.2: Randomized Parameter Optimization
- 3.2.3. Tips for parameter search
- 3.2.4. Alternatives to brute force parameter search

3.3. Metrics and scoring: quantifying the quality of predictions

- . 3.3.1. The according parameter: defining model evaluation rules
- 3.3.2. Classification metrics.
- 3.3.3. Multilabel ranking metrics
- · 3.3.4. Regression metrics
- 3.3.5. Quitering metrics
- · 3.3.6. Dummy estimators

3.4. Model persistence

- · 3.4.1. Persistence example
- 3.4.2. Security & maintainability limitations

3.5. Validation curves: plotting scores to evaluate models


- · 3.5.1. Validation curve
- . 3.5.2. Learning curve



Cross-validation: evaluating estimator Med performance




Learning the parameters of a prediction function and testing it on the same data is a methodological mistake: a model that would just repeat the labels of the samples that it has just seen would have a perfect score but would fail to predict anything useful on yetunseen data. This situation is called overfitting. To avoid it, it is common practice when performing a (supervised) machine learning experiment to hold out part of the available data as a test set X_test, y_test. Note that the word "experiment" is not intended to denote academic use only, because even in commercial settings machine learning usually starts out experimentally. Here is a flowchart of typical cross validation workflow in model training.











k-fold CV

https://scikitlearn.org/stable/modules/cross_validation.html#c omputing-cross-validated-metrics







Any Question?



Metrics and scoring: quantifying the Med quality of predictions



There are 3 different APIs for evaluating the quality of a model's predictions:

- > Estimator score method: Estimators have a score method providing a default evaluation criterion for the problem they are designed to solve. This is not discussed on this page, but in each estimator's documentation.
- \triangleright Scoring parameter: Model-evaluation tools using cross-validation (such as model_selection. cross_val_score and model_selection.GridSearchCV) rely on an internal scoring strategy. This is discussed in the section The scoring parameter: defining model evaluation rules.
- > Metric functions: The metrics module implements functions assessing prediction error for specific purposes. These metrics are detailed in sections on Classification metrics, Multilabel ranking metrics, Regression metrics and Clustering metrics.

https://scikit-learn.org/stable/modules/model_evaluation.html





Scoring	Function	Comment
Classification		
'accuracy'	metrics.accuracy_score	
'balanced_accuracy'	motrics.balanced_accuracy_score	
'average_precision'	metrics.average_precision_score	
'neg_brier_score'	metrics.brier_score_loss	
'f1'	metrics.fl_score	for binary targets
'f1_micro'	metrics.fl_score	micro-averaged
f1_macro*	metrics.fl_score	macro-averaged
f1_weighted	metrics.fl_score	weighted average
f1_samples'	metrics.fl_score	by multilabel sample
neg_log_loss*	metrics.log_loss	requires predict_proba support
precision' etc.	metrics.precision_score	suffixes apply as with 'f1'
recall' etc.	metrics.recall_score	suffixes apply as with 'f1'
jaccard' etc.	metrics.jaccard_score	suffixes apply as with 'f1'
roc_auc'	metrics.roc_auc_score	
roc_auc_ovr'	metrics.roc_auc_score	
roc_auc_ovo'	metrics.roc_auc_score	
roc_auc_ovr_weighted'	metrics.roc_auc_score	





Scoring	Function	Comment
Classification		
'accuracy'	metrics.accuracy_score	
'balanced_accuracy'	metrics.balanced_accuracy_score	
'average_precision'	metrics.average_precision_score	
'neg_brier_score'	metrics.brier_score_loss	
'f1'	metrics.fl_score	for binary targets
'f1_micro'	metrics.fl_score	micro-averaged
f1_macro*	metrics.fl_score	macro-averaged
f1_weighted	metrics.fl_score	weighted average
f1_samples'	metrics.fl_score	by multilabel sample
neg_log_loss*	metrics.log_loss	requires predict_proba support
precision' etc.	metrics.precision_score	suffixes apply as with 'f1'
recall' etc.	metrics.recall_score	suffixes apply as with 'f1'
jaccard' etc.	metrics.jaccard_score	suffixes apply as with 'f1'
roc_auc'	metrics.roc_auc_score	
roc_auc_ovr'	metrics.roc_auc_score	
roc_auc_ovo'	metrics.roc_auc_score	
roc_auc_ovr_weighted'	metrics.roc_auc_score	





Some of these are restricted to the binary classification case:

```
precision_recall_curve(y_true, probas_pred, *) Compute precision-recall pairs for different probability thresholds

roc_curve(y_true, y_score, *[, pos_label, ...]) Compute Receiver operating characteristic (ROC)
```

Others also work in the multiclass case:

balanced_accuracy_score(y_true, y_pred, * [,])	Compute the balanced accuracy
cohen_kappa_score(y1, y2, *[, labels,])	Cohen's kappa: a statistic that measures inter-annotator agreement.
confusion_matrix(y_true, y_pred, *[,])	Compute confusion matrix to evaluate the accuracy of a classification.
hinge_loss(y_true, pred_decision, *[,])	Average hinge loss (non-regularized)
matthews_corrcoof(y_true, y_pred, *[,])	Compute the Matthews correlation coefficient (MCC)
roc_auc_score(y_true, y_score, * [, average,])	Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores.





Some also work in the multilabel case:

accuracy_score(y_true, y_pred, *[,])	Accuracy classification score.
classification_report(y_true, y_pred, *[,])	Build a text report showing the main classification metrics.
f1_score(y_true, y_pred, *[, labels,])	Compute the F1 score, also known as balanced F-score or F-measure
fbeta_score(y_true, y_pred, *, beta[,])	Compute the F-beta score
hamming_loss(y_true, y_pred, * [, sample_weight])	Compute the average Hamming loss.
jaccard_score(y_true, y_pred, *[, labels,])	Jaccard similarity coefficient score
log_loss(y_true, y_pred, *[, eps,])	Log loss, aka logistic loss or cross-entropy loss.
multilabel_confusion_matrix(y_true, y_pred, *)	Compute a confusion matrix for each class or sample
precision_recall_fscore_support(y_true,)	Compute precision, recall, F-measure and support for each class
precision_score(y_true, y_pred, *[, labels,])	Compute the precision
recall_score(y_true, y_pred, *[, labels,])	Compute the recall
roc_auc_score(y_true, y_score, *[, average,])	Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores.
zero_one_loss(y_true, y_pred, *[,])	Zero-one classification loss.

And some work with binary and multilabel (but not multiclass) problems:

average_precision_score(y_true, y_score, *) Compute average precision (AP) from prediction scores





In extending a binary metric to multiclass or multilabel problems, the data is treated as a collection of binary problems, one for each class. There are then a number of ways to average binary metric calculations across the set of classes, each of which may be useful in some scenario. Where available, you should select among these using the average parameter.

"macro" simply calculates the mean of the binary metrics, giving equal weight to each class. In problems where infrequent classes are nonetheless important, macro-averaging may be a means of highlighting their performance. On the other hand, the assumption that all classes are equally important is often untrue, such that macro-averaging will over-emphasize the typically low performance on an infrequent class.

"weighted" accounts for class imbalance by computing the average of binary metrics in which each class's score is weighted by its presence in the true data sample.

"micro" gives each sample-class pair an equal contribution to the overall metric (except as a result of sample-weight). Rather than summing the metric per class, this sums the dividends and divisors that make up the per-class metrics to calculate an overall quotient. Micro-averaging may be preferred in multilabel settings, including multiclass classification where a majority class is to be ignored.

"samples" applies only to multilabel problems. It does not calculate a per-class measure, instead calculating the metric over the true and predicted classes for each sample in the evaluation data, and returning their (sample_weight-weighted) average.

Selecting average=None will return an array with the score for each class.



Classification metrics: Confusion matrix

In the field of <u>machine learning</u> and specifically the problem of <u>statistical classification</u>, a **confusion matrix**, also known as an error matrix, is a specific table layout that allows visualization of the performance of an algorithm, typically a <u>supervised learning</u> one (in <u>unsupervised learning</u> it is usually called a <u>matching matrix</u>). Each row of the <u>matrix</u> represents the instances in a predicted class while each column represents the instances in an actual class (or vice versa). The name stems from the fact that it makes it easy to see if the system is confusing two classes (i.e. commonly mislabeling one as another).

混淆矩阵也称误差矩阵,是表示精度评价的一种标准格式,用n行n列的矩阵形式来表示。具体评价指标有总体精度、<u>制图精度</u>、<u>用户精度</u>等,这些精度指标从不同的侧面反映了图像分类的精度。在人工智能中,混淆矩阵(confusion matrix)是可视化工具,特别用于<u>监督学习</u>,在<u>无监督学习</u>一般叫做匹配矩阵。在图像精度评价中,主要用于比较分类结果和实际测得值,可以把分类结果的精度显示在一个混淆矩阵里面。混淆矩阵是通过将每个实测像元的位置和分类与分类图像中的相应位置和分类相比较计算的。

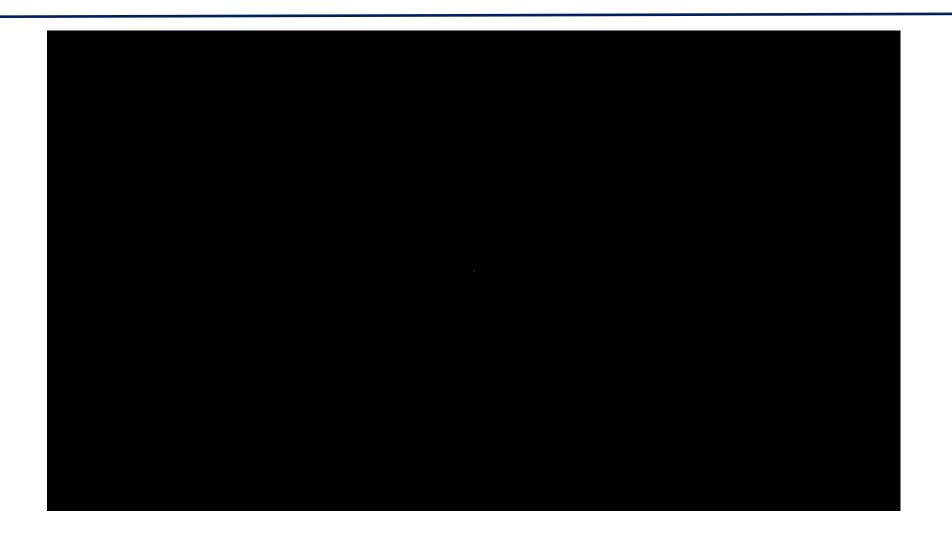


Classification metrics: Confusion matrix

		True condition				
	Total population	Condition positive	Condition negative	Prevalence = Σ Condition positive Σ Total population	Σ True posit	uracy (ACC) = ive + Σ True negative tal population
Predicted condition	Predicted condition positive	True positive	False positive, Type I error	Positive predictive value (PPV), Precision = Σ True positive Σ Predicted condition positive	False discovery rate (FDR) = Σ False positive Σ Predicted condition positive	
	Predicted condition negative	False negative, Type II error	True negative	False omission rate (FOR) = Σ False negative Σ Predicted condition negative	Negative predictive value (NPV) = Σ True negative Σ Predicted condition negative	
		True positive rate (TPR), Recall, Sensitivity, probability of detection, Power $= \frac{\Sigma \text{ True positive}}{\Sigma \text{ Condition positive}}$	False positive rate (FPR), Fall-out, probability of false alarm = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Condition negative}}$	Positive likelihood ratio (LR+) = TPR FPR	Diagnostic odds ratio (DOR)	F ₁ score = 2 · Precision · Recal
		False negative rate (FNR), Miss rate $= \frac{\Sigma \text{ False negative}}{\Sigma \text{ Condition positive}}$	Specificity (SPC), Selectivity, True negative rate (TNR) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Condition negative}}$	Negative likelihood ratio (LR-) = FNR TNR	$= \frac{LR+}{LR-}$	2 · Precision · Recal



Classification metrics: Confusion matrix





Any Question?





Supervised learning

- 1. Linear Models
- 2. Linear and Quadratic Discriminant Analysis
- 3. Kernel ridge regression
- 4. Support Vector Machines
- 5. Stochastic Gradient Descent
- 6. Nearest Neighbors
- 7. Gaussian Processes
- 8. Cross decomposition
- 9. Naive Bayes
- **10. Decision Trees**
- 11. Ensemble methods
- 12. Multiclass and multilabel algorithms
- 13. Feature selection
- 14. Semi-Supervised
- 15. Isotonic regression
- 16. Probability calibration
- 17. Neural network models (supervised)

1.1. Linear Models

- 1.1.1. Ordinary Least Squares
- 1.1.2. Ridge regression and classification
- 1.1.3. Lasso
- 1.1.4. Multi-task Lasso
- 1.1.5. Elastic-Net
- 1.1.6. Multi-task Elastic-Net
- 1.1.7. Least Angle Regression
- 1.1.8. LARS Lasso
- 1.1.9. Orthogonal Matching Pursuit (OMP)
- 1.1.10. Bayesian Regression
- 1.1.11. Logistic regression
- 1.1.12. Generalized Linear Regression
- 1.1.13, Stochastic Gradient Descent SGD
- 1.1.14. Perceptron
- 1.1.15. Passive Aggressive Algorithms
- 1.1.16. Robustness regression: outliers and modeling errors
- 1.1.17. Polynomial regression: extending linear models with basis functions



Supervised learning

1.4. Support Vector Machines

- 1.4.1 Classification
- 1.4.2. Regression
- 1.4.3. Density estimation, novelty detection
- 1.4.4. Complexity
- 1.4.5. Tips on Practical Use
- 1.4.6. Kernel functions
- 1.4.7. Mathematical formulation
- 1.4.8. Implementation details

1.5. Stochastic Gradient Descent

- 1.5.1. Classification
- 1.5.2. Regression
- 1.5.3. Stochastic Gradient Descent for sparse data
- 1.5.4. Complexity
- 1.5.5. Stopping criterion
- 1.5.6. Tips on Practical Use
- 1.5.7. Mathematical formulation
- 1.5.8. Implementation details

1.9. Naive Bayes

- 1.9.1. Gaussian Naive Bayes
- 1.9.2. Multinomial Naive Bayes
- 1.9.3. Complement Naive Bayes
- 1.9.4. Bernoulli Naive Bayes
- 1.9.5. Categorical Naive Bayes
- 1.9.6. Out-of-core naive Bayes model fitting

1.10. Decision Trees

- 1.10.1. Classification
- 1.10.2 Regression
- 1.10.3, Multi-output problems
- 1.10.4. Complexity
- 1.10.5: Trps on practical use
- 1.10.6, Tree algorithms: ID3, C4.5, C5.0 and CART
- 1.10.7. Mathematical formulation
- 1.10.8, Minimal Cost-Complexity Pruning.

1.11. Ensemble methods

- 1.11.1. Bagging meta-estimator
- 1.11.2. Forests of randomized trees
- 1.11.3. AdaBoost
- 1.11.4. Gradient Tree Boosting
- 1.11.5. Histogram-Based Gradient Boosting
- 1.11.6, Voting Classifier
- 1.11.7. Voting Regressor
- 1.11.8. Stacked generalization



Supervised Learning

1.4. Support Vector Machines

- 1.4.1 Classification
- 1.4.2. Regression
- 1.4.3. Density estimation, novelty detection
- 1.4.4. Complexity
- 1.4.5. Tips on Practical Use
- 1.4.6. Kernel functions
- 1.4.7. Mathematical formulation
- 1.4.8. Implementation details

1.5. Stochastic Gradient Descent

- 1.5.1. Classification
- 1.5.2. Regression
- 1.5.3. Stochastic Gradient Descent for sparse data
- 1.5.4. Complexity
- 1.5.5. Stopping criterion
- 1.5.6. Tips on Practical Use
- 1.5.7. Mathematical formulation
- 1.5.8. Implementation details

1.9. Naive Bayes

- 1.9.1. Gaussian Naive Bayes
- 1.9.2. Multinomial Naive Bayes
- 1.9.3. Complement Naive Bayes
- 1.9.4. Bernoulli Naive Bayes
- 1.9.5. Categorical Naive Bayes
- 1.9.6. Out-of-core naive Bayes model fitting

1.10. Decision Trees

- 1.10.1. Classification
- 1.10.2 Regression
- 1.10.3, Multi-output problems
- 1.10.4. Complexity
- 1.10.5: Trps on practical use
- 1.10.6, Tree algorithms: ID3, C4.5, C5.0 and CART
- 1.10.7. Mathematical formulation
- 1.10.8, Minimal Cost-Complexity Pruning.

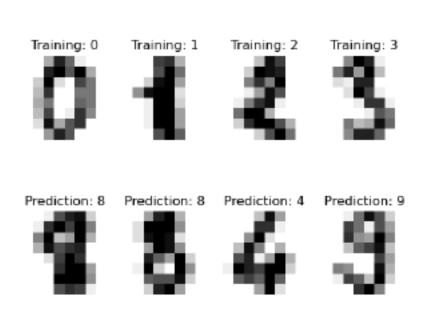
1.11. Ensemble methods

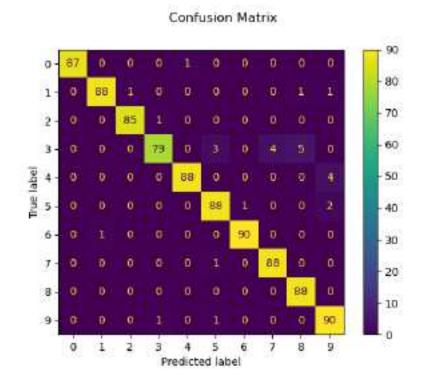
- 1.11.1. Bagging meta-estimator
- 1.11.2. Forests of randomized trees
- 1.11.3. AdaBoost
- 1.11.4. Gradient Tree Boosting
- 1.11.5. Histogram-Based Gradient Boosting
- 1.11.6, Voting Classifier
- 1.11.7. Voting Regressor
- 1.11.8. Stacked generalization



Exercise: Plot the confusion matrix

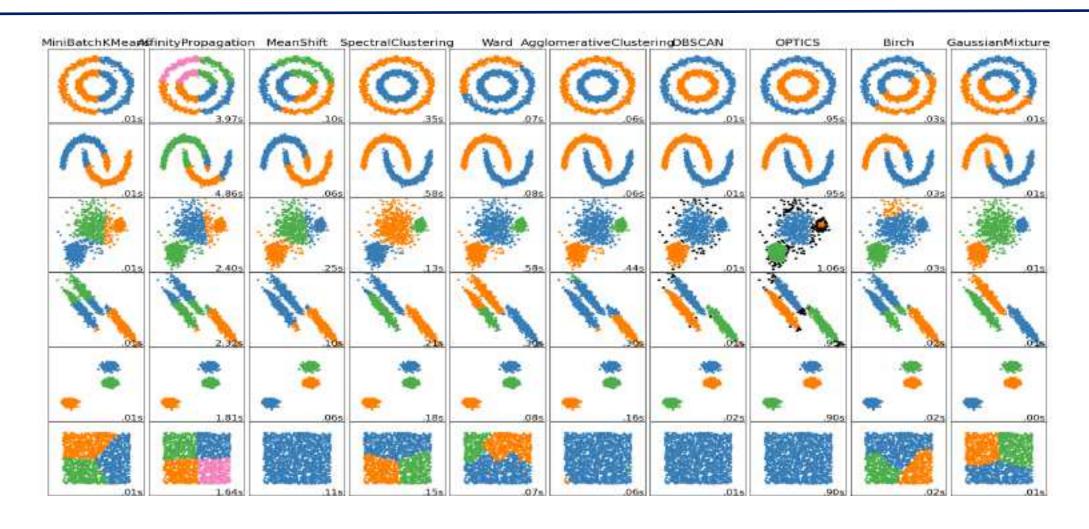
Recognizing hand-written digits







Clustering





Clustering

Method name	Parameters	Scalability	Usecase	Geometry (metric used)
K-Means	number of clusters	Very large n_samples, medium n_clusters with MiniBatch code	General-purpose, even cluster size, flat geometry, not too many clusters	Distances between points
Affinity propagation	damping, sample preference	Not scalable with n_samples	Many clusters, uneven cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Mean-shift	bandwidth	Not scalable with	Many clusters, uneven cluster size, non-flat geometry	Distances between points
Spectral clustering	number of clusters	Medium n_samples, small n_clusters	Few clusters, even cluster size, non-flat geometry	Graph distance (e.g. nearest-neighbor graph)
Ward hierarchical clustering	number of clusters or distance threshold	Large n_samples and n_clusters	Many clusters, possibly connectivity constraints	Distances between points
Agglomerative clustering	number of clusters or distance threshold, linkage type, distance	Large n_samples and n_clusters	Many clusters, possibly connectivity constraints, non Euclidean distances	Any pairwise distance
DBSCAN	neighborhood size	Very large n_samples, medium n_clusters	Non-flat geometry, uneven cluster sizes	Distances between nearest points
OPTICS	minimum cluster membership	Very large n_samples, large n_clusters	Non-flat geometry, uneven cluster sizes, variable cluster density	Distances between points
Gaussian mixtures	many	Not scalable	Flat geometry, good for density estimation	Mahalanobis distances to centers
Birch	branching factor, threshold, optional global clusterer.	Large n_clusters and n_samples	Large dataset, outlier removal, data reduction.	Euclidean distance between points

https://scikit-learn.org/stable/modules/clustering.html#clustering



Clustering

Cluster analysis or clustering is the task of grouping a set of objects in such a way that objects in the same group (called a cluster) are more similar (in some sense) to each other than to those in other groups (clusters). It is a main task of exploratory <u>data mining</u>, and a common technique for <u>statistical data analysis</u>, used in many fields, including <u>pattern recognition</u>, <u>image analysis</u>, <u>information retrieval</u>, <u>bioinformatics</u>, <u>data compression</u>, computer graphics and machine learning.

聚类分析指将物理或抽象对象的集合分组为由类似的对象组成的多个类的分析过程。它是一种重要的人类行为。聚类分析的目标就是在相似的基础上收集数据来分类。聚类源于很多领域,包括数学,计算机科学,统计学,生物学和经济学。在不同的应用领域,很多聚类技术都得到了发展,这些技术方法被用作描述数据,衡量不同数据源间的相似性,以及把数据源分类到不同的簇中。

聚类与分类的不同在于,聚类所要求划分的类是未知的。 聚类是将数据分类到不同的类或者簇这样的一个过程,所以同一个簇中的对象有很大的 相似性,而不同簇间的对象有很大的相异性。

https://scikit-learn.org/stable/modules/clustering.html#clustering



Any Question?





Preprocessing data

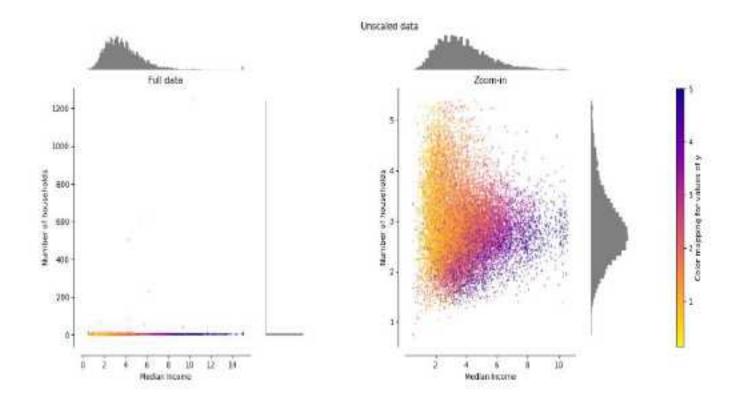
The sklearn.preprocessing package provides several common utility functions and transformer classes to change raw feature vectors into a representation that is more suitable for the downstream estimators.

In general, learning algorithms benefit from standardization of the data set. If some outliers are present in the set, robust scalers or transformers are more appropriate. The behaviors of the different scalers, transformers, and normalizers on a dataset containing marginal outliers is highlighted in Compare the effect of different scalers on data with outliers.

- Standardization, or mean removal and variance scaling
- 2. Non-linear transformation
- 3. Normalization
- 4. Encoding categorical features
- 5. <u>Discretization</u>
- 6. Imputation of missing values
- 7. Generating polynomial features
- 8. <u>Custom transformers</u>

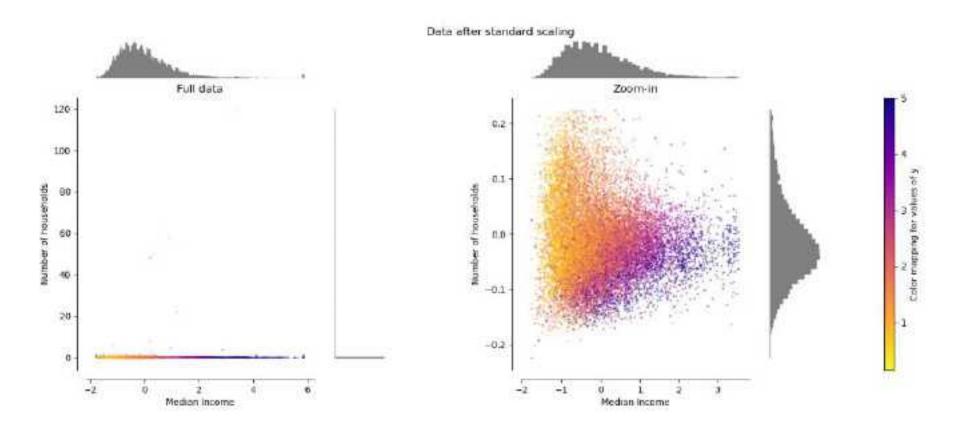
Compare the effect of different scalers Med on data with outliers

- 1. StandardScaler
- 2. MinMaxScaler
- 3. MaxAbsScaler
- 4. RobustScaler
- 5. PowerTransformer
- 6. <u>QuantileTransformer</u> (Gaussian output)
- 7. <u>QuantileTransformer (uniform output)</u>
- 8. <u>Normalizer</u>



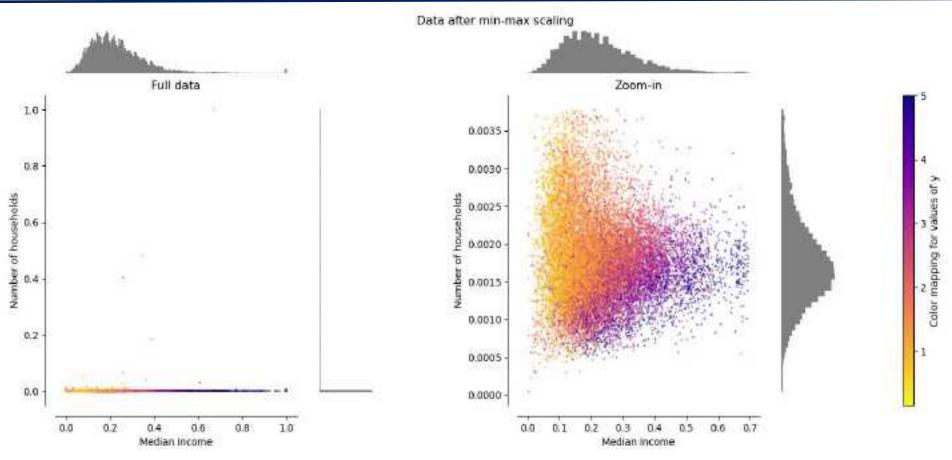
Original data (California housing dataset)

Compare the effect of different scalers Med on data with outliers



StandardScaler removes the mean and scales the data to unit variance. However, the outliers have an influence when computing the empirical mean and standard deviation which shrink the range of the feature values as shown in the left figure above

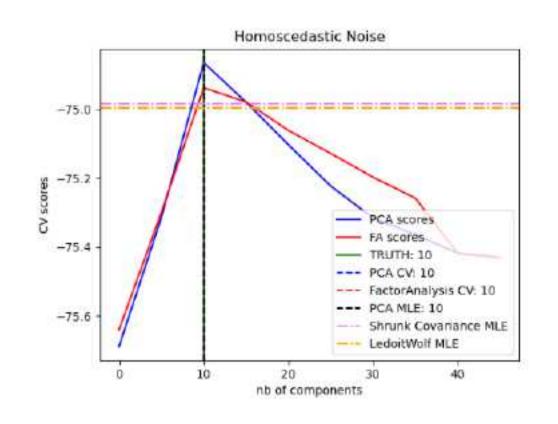
Compare the effect of different scalers Med on data with outliers



MinMaxScaler rescales the data set such that all feature values are in the range [0, 1] as shown in the right panel below. However, this scaling compress all inliers in the narrow range [0, 0.005] for the transformed number of households.

Compare the effect of different scalers Med on data with outliers

- 1. <u>Principal component analysis</u> (PCA)
- 2. <u>Truncated singular value</u> <u>decomposition and latent</u> <u>semantic analysis</u>
- 3. Dictionary Learning
- 4. <u>Factor Analysis</u>
- 5. <u>Independent component analysis</u> (ICA)
- 6. <u>Non-negative matrix</u> factorization (NMF or NNMF)
- 7. <u>Latent Dirichlet Allocation (LDA)</u>





Exercise: training linear models

https://github.com/ageron/handson-ml2/blob/master/04 training linear models.ipynb

Given a data set $\{y_i, x_{i1}, \dots, x_{ip}\}_{i=1}^n$ of n statistical units, a linear regression model assumes that the relationship between the dependent variable y and the p-vector of regressors x is linear. This relationship is modeled through a disturbance term or error variable ε — an unobserved random variable that adds "noise" to the linear relationship between the dependent variable and regressors. Thus the model takes the form

$$y_i = eta_0 + eta_1 x_{i1} + \dots + eta_p x_{ip} + arepsilon_i = \mathbf{x}_i^{\mathsf{T}} oldsymbol{eta} + arepsilon_i, \qquad i = 1, \dots, n,$$

where $^{\top}$ denotes the transpose, so that $\mathbf{x}_{i}^{\top}\boldsymbol{\beta}$ is the inner product between vectors \mathbf{x}_{i} and $\boldsymbol{\beta}$.

Often these n equations are stacked together and written in matrix notation as

$$egin{aligned} egin{aligned} L(D,ec{eta}) &= \|Xec{eta} - Y\|^2 \\ &= (Xec{eta} - Y)^T(Xec{eta} - Y) \\ &= Y^TY - Y^TXec{eta} - ec{eta}^TX^TY + ec{eta}^TX^TXec{eta} \end{aligned}$$



Exercise: MNIST Classification

Reference: https://github.com/ageron/handson-ml2/blob/master/03_classification.ipynb



Any Question?



PyTorch



目录

- 1 History and similar frameworks
- Architecture of Al model
- Common APIs
- 4 Some learning resources





- 2002--Torch
- 2011--Torch7



Merits: flexible dynamic

user-friendly

Demerits: based on Lua

Caffe

- 2013--Caffe
- 2017--Caffe2

Merits: fast(based on C++)

Demerits: not flexible

- 2017--PyTorch
- 2018--PyTorch v0.4






1 History and similar frameworks













History and similar frameworks

Evaluation


	PyTorch	TensorFlow 1	TensorFlow 2
性能	****	***	****
生态	***	****	**
工业界	***	****	**
学术界	****	***	**
上手难度	****	**	****
易用性	****	*	****
兼容性	****	*	*
发展前景	****	0	****



2 A

Archetecture of Al model

How to build a model?








Archetecture of AI model

How to build a model?





2

Archetecture of AI model

How to build a model?

A model may have:


- Convolutional layers
- Fully connetected layers
- Batch normolization layers
- Pooling layers
- Other operations: concat/split/add ...



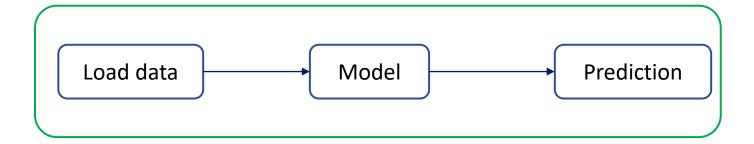


Archetecture of AI model

How to train a model?



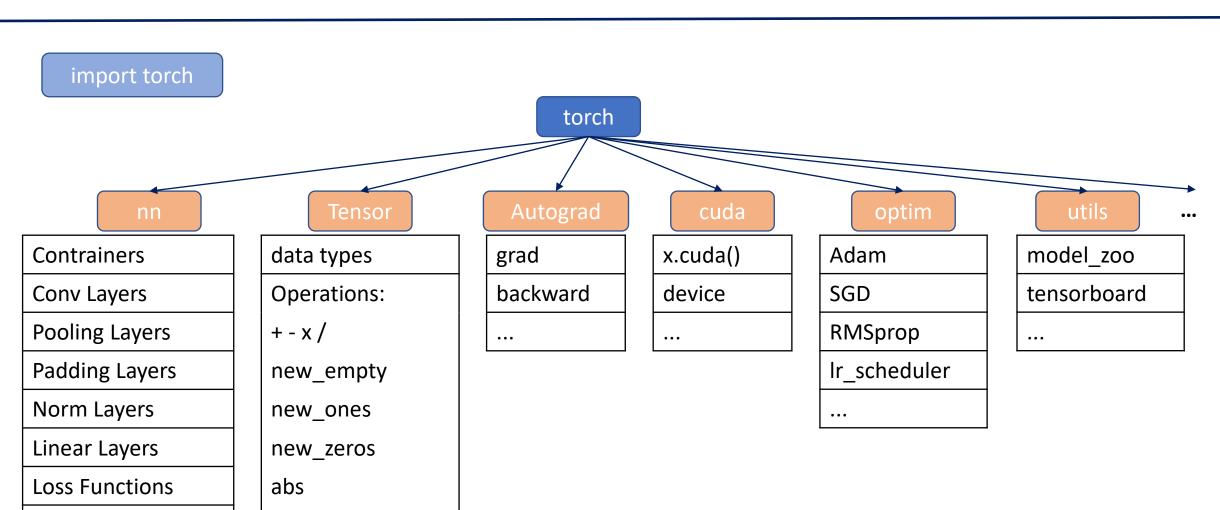
Predict:





asin/acos/atan

Vision Layers





torch.nn

Containers

Module	Base class for all neural network modules.
Sequential	A sequential container,
ModuleList	Holds submodules in a list.
ModuleDict	Holds submodules in a dictionary.
ParameterList	Holds parameters in a list.
ParameterDict	Holds parameters in a dictionary.



torch.nn

Convolution Layers

Curwlin	Applies a 1D convolution over an input signal composed of several input planes.
nd.Conv2d	Applies a 2D convolution over an input signal composed of several input planes.
Comau	Applies a 3D convolution over an input signal composed of soveral input planes.
me, CorreTs assupprise Life	Applies a 1D transposed convolution operator over an input image composed of several input planes.
, Спо у Узланиризм 2 d	Applies a 2D transposed convolution operator over an input image composed of several input planes.
an, ConyTzanapose3d	Applies a 3D transposed convolution operator over an input image composed of several input planes.
me, linfold	Extracts sliding local blocks from a batched input tensor.
No.Feli	Combines an array of sliding local blocks into a large containing tensor.



torch.nn

Pooling layers

or, RAPPORTER	Applies a 1D max pooling over an input signal composed several input prieries.	
en, Northerlad	Applies a 2D max pooling over an input signal composed of several input places.	
on floir-clad	Applies a 3D max pooling over an input signal composed of several input planes.	
er Matthewalld	Computes a partial inverse of humbrollid.	
on Maxillager 124	Computes a partial inverse of MarPani2H.	
en, Nathrynäääd	Computes a pertial inverse of Histoista.	
en, AugPool16	Applies a 1D average pooling over an input signal composed of several input planes.	
on. AvgPoot28	Applies a 2D average pooling over an input signal compose of several input planes.	

- top-sales	Applies a 3D everage pooling over an input signal correspond of several input planes.	
re-Tractions/Real-Model	Applies a 20 fractional may pooling over an input signal composed of several input planes.	
m_LPFootia	Applies a 1D power-overage pooling over an input signal composed of several input planes.	
en. DEPositio	Applies a 2D power-average pooling over an input signal composed of several input planes.	
to ManticobarPositio	Applies a 1D subgrove max pooling over an input signal composed of several input planer.	
rechlorischerheitze	Applies a 3D intertive must posting over an input signal composed of several input planes.	
in hautsomatustal	Applies a 3D adaptive max pooling over an input signal compared of several imput plants.	
re. MagriroshigFoolid	Applies a 1D adaptive average pooling over an input signal composed of sever-timpat planes.	
re-startive-logPoilin	Applies a 2D adaptive average pooling over an input signal composed of several input planer.	
nn. Adapt tor Ang Foot list	Applies a 3D adoption average pooling over an input signal compassed of anyeral imput planes.	





torch.nn

Padding Layers

or an annual section of the section	Pads the input tensor using the reflection of the input	
Am. RetlectionPadld	boundary.	
am ReflectionPad2d	Pads the input tensor using the reflection of the input boundary.	
- SepticationPaulid	Pads the input tensor using replication of the input boundary.	
on. ReplicationPod2d	Pads the input tensor using replication of the input	
	boundary.	
on.ReplicationPattid	Pads the input tensor using replication of the input boundary.	
m. ZazuPadžd	Pads the input tensor boundaries with zero.	
on ConstantPadLd	Pads the input tensor boundaries with a constant value.	
en.ConstantPad2d	Pads the input tensor boundaries with a constant value.	
un.ConstantPad3d	Pads the input tensor boundaries with a constant value.	




torch.Tensor

dtype	CPU tensor	GPU tensor
torch.float32 or torch.float	torch.FloatTensor	torch.cuda.FloatTensor
torch.float64 or torch.double	torch.DoubleTensor	torch, cudn.DoubleTensor
torch.float16 or torch.half	torch.HalfTensor	torch.cuda.HalfTensor
torch.bfloat16	torch.BFloat16Tensor	torch.cuda.BFloat16Tenso
	torch.float32 or torch.float64 or torch.double torch.float16 or torch.half	torch.float32 or torch.float64 or torch.float64 or torch.double torch.float16 or torch.half



import torchvision







4 Learning resources

Learning resources





3



4 Learning resources

Learning resources









Tensorflow-Introduction



1

TensorFlow is a <u>free</u> and <u>open-source</u> <u>software</u> <u>library</u> for <u>dataflow</u> and <u>differentiable</u> <u>programming</u> across a range of tasks. It is a symbolic math library, and is also used for <u>machine</u> <u>learning</u> applications such as <u>neural networks</u>. It is used for both research and production at <u>Google</u>.



Tensorflow-Development



2015年11月,Google正式发布了Tensorflow的白皮书并开源 TensorFlow 0.1 版本。

2017年02月, Tensorflow正式发布了1.0.0版本, 同时也标志 着稳定版的诞生。

2019年10月, TensorFlow在经历七个多月(2019年3月1日-2019年10月1日)的2.0 Alpha 版本的更新迭代后发布 2.0 正式版。



Tensorflow1.0 vs Tensorflow2.0



Functions, not sessions TF1 TF2 a = tf.constant(5) a = tf.constant(5) Symbolic Concrete b = tf.constant(3)b = tf.constant(3)c = a * b c = a * bprint(c) with tf.Session() as sess: print(sess.run(c))



Tensorflow1.0 vs Tensorflow2.0



TF 1.x

Difficult to debug

API confusion

It is difficult to get started, but it is still difficult to enter

A large number of researchers turn to PyTorch

TF 2.x


TF+Keras

Easy to use



Tensorflow2.0





https://www.youtube.com/watch?v=5ECD8J3dvDQ



Tensorflow2.0-Install



Install TensorFlow 2

TensorFlow is tested and supported on the following 64-bit

systems:

Python 3.5–3.8
Ubuntu 16.04 or later
Windows 7 or later (with C++ redistributable)
macOS 10.12.6 (Sierra) or later (no GPU support)
Raspbian 9.0 or later

```
# Requires the latest pip
$ pip install --upgrade pip

# Current stable release for CPU and GPU
$ pip install tensorflow

# Or try the preview build (unstable)
$ pip install tf-nightly
```



Tensorflow2.0-Tensor



When writing a TensorFlow program, the main object that is manipulated and passed around is the tf.Tensor.

A tf.Tensor has the following properties:

- a single data type (float32, int32, or string, for example)
- a shape

```
>>> # Compute some values using a Tensor

>>> c = tf.constant([[1.0, 2.0], [3.0, 4.0]])

>>> d = tf.constant([[1.0, 1.0], [0.0, 1.0]])

>>> e = tf.matmul(c, d)

>>> print(e)

tf.Tensor(

[[1. 3.]

[3. 7.]], shape=(2, 2), dtype=float32)
```



Tensorflow2.0-Tensors



Some useful examples:

```
# Strip leading and trailing 2 elements
foo = tf.constant([1,2,3,4,5,6])
print(foo[2:-2].eval()) # == [3,4]
foo = tf.constant([[1,2,3], [4,5,6], [7,8,9]])
print(foo[::2,::-1].eval()) # => [[8,2,1], [9,8,7]]
# Use scalar tensors as indices on both dimensions
print(foo[tf.constant(0), tf.constant(2)].eval()) # => 3
# Insert enother dimension
foo = tf.constant([[1,2,3], [4,5,6], [7,8,9]])
print(foo[tf.newaxis, :, :].eval()) = > [[[ 2.3], [4.5.6], [7.8.9]]]
print(foo[:, tf.newaxis, :]:eval()) = > [[[1.2.3]], [[4.5.6]], [[7.8.9]]]
print(foo[:, :, tf.newaxis].eval()) # => [[[1], [2], [3]], [[4], [5], [6]].
1171.[8].[9]]]
# Ellipses (8 equivalent operations)
foo = tf.constant([[1,2,3], [4,5,6], [7,8,9]])
print(foo[tf.newaxis, :, :].eval()) # = [[[1.2.3], [4.5.6], [7.8.9]]]
print(foo[tf.newaxis, ...].eval()) # > [[[1,2,3], [4,5,6], [7,8,9]]]
print(foo[tf.newaxis].eval()) # >> [[[1,2,3], [4,5,6], [7,8,0]]]
# Hasks
foo = tf.constant([[1,2,3], [4,5,6], [7,8,9]])
print(foo[foo > 2].eval()) # => [3, 4, 5, 6, 7, 8, 9]
```

Get items in a tensor



Tensorflow2.0-Tensors



Some useful examples:

```
x = tf.constant([5, 4, 6])
y = tf.constant([5, 2, 5])
tf.math.greater(x, y) ==> [False, True, True]

x = tf.constant([5, 4, 6])
y = tf.constant([5])
tf.math.greater(x, y) ==> [False, False, True]
```

Returns the truth value of (x > y) element-wise.

```
>>> a = tf.constant([1, 2, 3, 4, 5, 6], shape=[2, 3])
>>> a # 2-0 tensor
<tf.Tensor: shape=(2, 3), dtype=int32, numpy=
array([[1, 2, 3],
       [4, 5, 6]], dtype=int32)>
>>> b = tf.constant([7, 8, 9, 10, 11, 12], shape=[3, 2])
>>> b = 2-D tensor
<tf.Tensor: shape=(3, 2), dtype=int32, numpy=
array([[ 7, 8],
       [ 9, 10],
       [11, 12]], dtype=int32)>
>>> c = tf.matmul(a, b)
>>> c # # # #
<tf.Tensor: shape=(2, 2), dtype=int32, numpy=
array([[ 58, 64],
       [139, 154]], dtype=int32)>
```

2-D tensor matrix multiplication



Tensorflow2.0-Tensors



Some useful examples:


```
>>> c = tf.constant([True])
>>> x = tf.constant([False, True, True, False])
>>> tf.math.logical_xor(c, x)
<tf.Tensor: shape=(4,), dtype=bool, numpy=array([ True, False, False, True])>
```


```
>>> y = tf.constant([False, False, True, True])
>>> z = tf.constant([False, True, False, True])
>>> tf.math.logical_xor(y, z)
<tf.Tensor: shape=(4,), dtype=bool, numpy=array([False, True, True, False])>
```

Logical XOR function.



Tensorflow2.0-Computational Graph







Tensorflow2.0-Example



```
import tensorflow as tf
```

Load and prepare the MNIST dataset. Convert the samples from integers to floating-point numbers:

```
mnist = tf.keras.datasets.mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train, x_test = x_train / 255.8, x_test / 255.8
```

Build the <u>tf.keras.Sequential</u> model by stacking layers. Choose an optimizer and loss function for training:

```
model = tf.keras.models.Sequential([
   tf.keras.layers.Flatten(input_shape=(28, 28)),
   tf.keras.layers.Dense(128, activation='relu'),
   tf.keras.layers.Dropout(8:2),
   tf.keras.layers.Dense(18)
])
```



Tensorflow2.0-Example



For each example the model returns a vector of "logits" or "log-odds" scores, one for each class.

```
predictions = model(x_train[:1]).numpy()
predictions
```

The <u>tf.nn.softmax</u> function converts these logits to "probabilities" for each class:

```
tf.nn.softmax(predictions).numpy()
```

The <u>losses.SparseCategoricalCrossentropy</u> loss takes a vector of logits and a True index and returns a scalar loss for each example.



Tensorflow2.0-Example



This loss is equal to the negative log probability of the true class: It is zero if the model is sure of the correct class.

```
loss_fn(y_train[:1], predictions).numpy()
```

```
model.compile(optimizer='adam',
loss=loss_fn,
metrics=['accuracy'])
```

The Model.fit method adjusts the model parameters to minimize the loss:

```
model.fit(x_train, y_train, epochs=5)
```

The <u>Model.evaluate</u> method checks the models performance, usually on a "<u>Validation-set</u>" or "<u>Test-set</u>".

```
model.evaluate(x_test, y_test, verbose=2)
```





```
# A vision model.
                                                             Question input
                                                                           Image input
# Encode an image into a vector.
vision_model = Sequential()
vision_model.add(Conv2D(64, (3, 3),
                                                              Embedding
                           activation= relu ,
                                                                            Conv2D /
                                                                           MaxPooling20
                           input_shape=(224, 224, 3)))
                                                                             stack
vision_model.add(MaxPooling2D())
                                                                LSTM
vision_model.add(Flatten())
# Get a tensor with the output of your vision model
                                                                    Concatenate
image_input = Input(shape=(224, 224, 3))
encoded_image = vision_model(image_input)
                                                                   Dense / softmax
```





```
# A language model.
                                                               Question input
                                                                             Image input
# Encode the question into a vector.
question_input = Input(shape=(100,),
                          dtype='int32',
                                                                Embedding
                          name="Question")
                                                                              Conv2D /
                                                                             MaxPooling2D
                                                                               stack
embedded = Embedding(input_dim=10000,
                                                                 LSTM
                        output_dim=256,
                        input_length=100)(question_input)
                                                                      Concatenate
encoded_question = LSTM(256)(embedded_question)
                                                                     Dense / softmax
```





```
# Concatenate the encoded image and question
                                                                Question input
                                                                                 Image input
merged = layers.concatenate([encoded_image,
                                    encoded_question])
                                                                  Embedding
                                                                                 Conv2D /
                                                                                MaxPooling2D
                                                                                   stack
                                                                   LSTM
                                                                        Concatenate
                                                                       Dense / softmax
```






```
# Train a classifier on top.
                                                            Question input
                                                                           Image input
output = Dense(1000,
                  activation='softmax')(merged)
                                                             Embedding
                                                                            Conv2D7
                                                                          MaxPooling2D
# You can train w/ .fit, .train_on_batch,
                                                                             stack
                                                               LSTM.
# or with a GradientTape.
vqa_model = Model(inputs=[image_input,
                                 question_input],
                                                                   Concatenate
                                 outputs=output)
                                                                  Dense / softmax
```



Tensorflow2.0-Models








Tensorflow2.0-Models





Natural language question answering

Anone: gustions beset as the contact of a given

View code O

https://tensorflow.google.cn/js/models



Node Js Pitch Prediction


want least in

View code ()

That a converside model is already beneful pitch types.

Tensorflow2.0-Demos





Visualize Model Training

behavior and trialing Long to rule.

Explore demo > View code ()

See flow to visualize in browner browning and model

https://tensorflow.google.cn/js/demos





CS 103 -06

Mid-Term Review and AI Platform Introduction

Jimmy Liu 刘江

TA: 章晓庆, 张洋, 肖尊杰, 杨冰

2020-10-23