

Computer Vision

CS308

Feng Zheng

SUSTech CS Vision Intelligence and Perception

Week 3



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY




Content

- Brief Review
- Point Operators
- Linear Filtering
- More Neighborhood Operators

Brief Review



Pinhole Camera




f = focal length

c = center of the camera



Projection Matrix



$$\mathbf{x} = \mathbf{K} \begin{bmatrix} \mathbf{R} & \mathbf{t} \end{bmatrix} \mathbf{X}$$

\mathbf{x} : Image Coordinates: $(u, v, 1)$

\mathbf{K} : **Intrinsic Matrix** (3x3)

\mathbf{R} : Rotation (3x3)

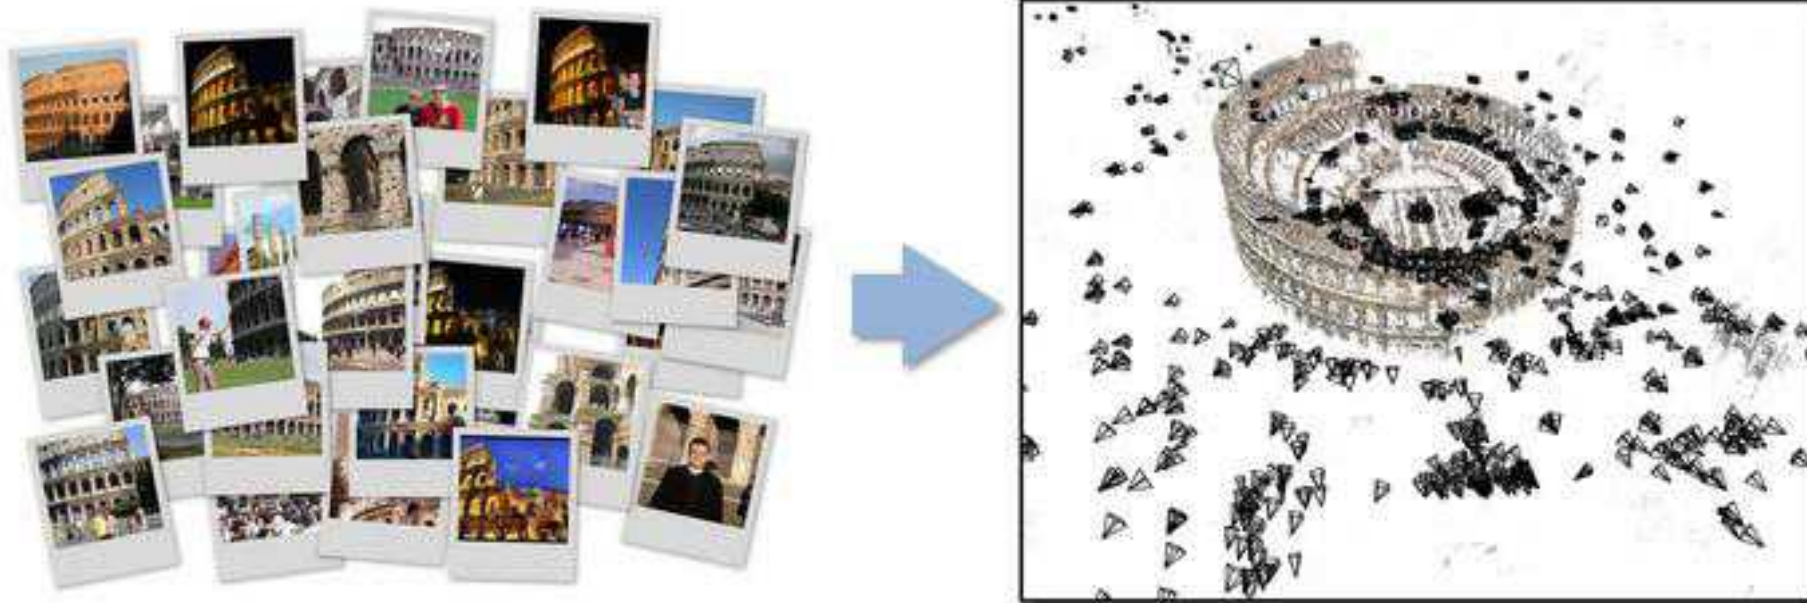
\mathbf{t} : Translation (3x1)

\mathbf{X} : World Coordinates: $(X, Y, Z, 1)$



An Example of Application

- Structure from Motion (SfM) for Unordered Image Collections





What is image processing?

- (a) original image
- (b) increased contrast
- (c) change in hue
- (d) quantized colors
- (e) blurred
- (f) rotated

0	0	0	1	1	0	0	0	0	2	2	0	0	0	0	1	1	5
0	0	0	0	0	0	0	0	2	2	0	0	0	0	0	0	1	0
0	0	0	0	0	1	1	0	0	0	0	0	0	0	0	0	0	0
1	0	0	0	0	0	0	0	34	19	0	0	4	97	34	0	0	0
0	0	0	0	1	1	0	0	88	152	56	42	145	194	29	0	0	0
0	0	0	2	2	0	0	0	58	208	237	231	251	206	19	0	1	2
0	0	2	2	1	0	0	0	47	229	238	230	220	196	47	0	2	2
0	1	1	0	1	1	0	0	71	236	195	200	218	204	49	0	1	0
1	1	0	0	0	0	0	0	134	231	234	216	198	136	0	0	0	0
1	1	0	0	0	0	0	0	46	108	66	58	99	107	21	0	0	0
0	0	0	0	0	0	7	164	120	0	18	36	36	6	0	0	0	1
0	0	1	0	4	126	231	183	62	14	20	44	20	0	0	1	1	5
0	1	1	0	84	223	224	240	190	14	11	80	36	0	0	0	0	0
1	1	0	2	164	224	217	235	186	16	54	187	37	0	0	0	0	0
1	1	0	28	211	221	218	234	201	46	166	204	12	0	0	0	0	0
0	0	0	57	219	218	224	238	182	50	211	152	0	0	1	1	0	5
0	0	0	121	217	226	230	229	126	45	196	106	0	0	0	2	3	3
0	0	46	209	217	225	238	189	50	94	230	161	49	0	1	3	1	0
0	0	91	235	218	222	243	164	71	200	241	250	196	7	0	0	0	0
1	0	72	233	231	223	244	157	134	243	243	255	185	3	0	0	0	0
0	0	36	206	244	232	244	183	188	249	244	255	178	1	0	0	0	0
0	0	14	93	109	185	251	137	137	221	181	168	120	5	0	1	1	0
0	0	7	22	1	45	130	74	61	77	42	9	0	0	0	1	1	5
0	0	0	0	3	24	23	7	0	2	4	0	0	0	1	1	0	0
1	0	0	0	0	1	1	0	0	0	0	0	1	1	0	0	0	0
1	1	0	0	0	0	1	1	0	0	0	0	1	1	0	0	0	0



(a)



(b)



(c)



(d)



(e)





(f)

Point (Pixel) Operators



Image Functions

- Image as a function in two-dimensional space





Point Operators

- A general image processing operator

- Continuous domain $g(\mathbf{x}) = h(f(\mathbf{x}))$ or $g(\mathbf{x}) = h(f_0(\mathbf{x}), \dots, f_n(\mathbf{x}))$
- Discrete images $g(i, j) = h(f(i, j))$

- Multiplication and **addition** $g(\mathbf{x}) = a f(\mathbf{x}) + b$
 - gain controls (points to a)
 - brightness (points to b)
 - spatially varying (points to the equation)

- Dyadic (**two-input**) operator $g(\mathbf{x}) = (1 - \alpha) f_0(\mathbf{x}) + \alpha f_1(\mathbf{x})$
- Gamma **correction** $g(\mathbf{x}) = [f(\mathbf{x})]^{1/\gamma}$
 - linear blend operator (points to the equation)

✓ Remove the non-linear mapping between input radiance and quantized pixel values


$$g(\mathbf{x}) = [f(\mathbf{x})]^{1/\gamma} \quad \gamma \approx 2.2$$

pixel locations (points to \mathbf{x})



Color Transforms

- (a-d) RGB
- (e-h) rgb
- (i-k) $L^*a^*b^*$
- (l-n) HSV
- **Brightening** a picture by adding a **constant** value to all three channels
 - Whether this achieves the desired effect of making the image look brighter?
 - Is there any undesirable **side-effects** or **artifacts**?



Some color ratio images multiplied by the middle gray value for better visualization



Compositing and Matting

- **Matting:** extracting the object from the original image
- **Compositing:** inserting one image into another image

$$C = (1 - \alpha)B + \alpha F$$



(a)



(b)




(c)



(d)

- (a) source image
(b) extracted foreground object
(c) alpha matte shown in grayscale
(d) new composite



B

(a)

\times

$(1 -$



α

(b)

$)$

$+$



αF

(c)

$=$




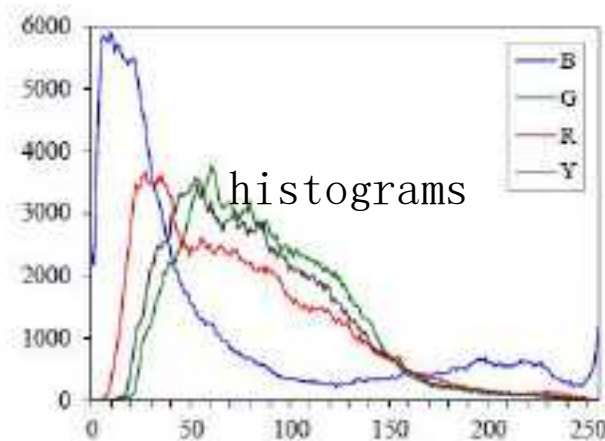
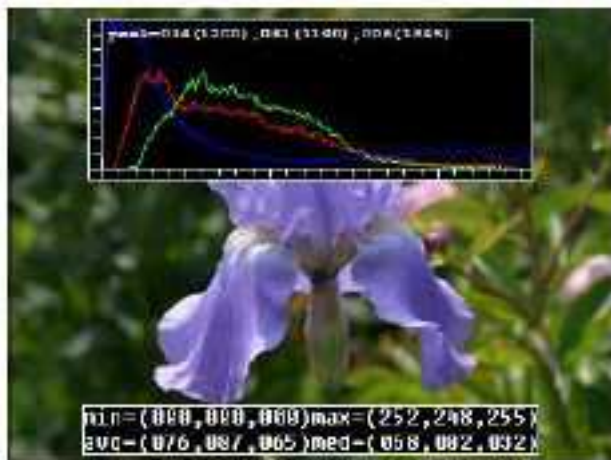
C

(d)



Histogram Equalization




- How can we automatically determine **best values** of pixels?
 - Look at the **darkest** and **brightest** values and map them to **pure** black and white
 - Find the **average** value, push it towards **middle** gray, and expand the range so that it more closely fills the **displayable** values
 - Histogram equalization: find an intensity mapping function $f(I)$ such that the **resulting histogram is flat**





Histogram Equalization

- A technique for **adjusting** image **intensities** to **enhance contrast**
 - An image would have a **linearized** cumulative distribution function (CDF)





Histogram Equalization

- An example




$$c(I) = \frac{1}{N} \sum_{i=0}^I \overset{\text{histogram}}{h(i)} = c(I-1) + \frac{1}{N} h(I) \quad I \text{ in } [0 \ 255]$$

- Why it is equalized?

$$f(I) = c(I)$$

$$f(I) = \alpha c(I) + (1 - \alpha)I$$

Maintain more of its original grayscale distribution while having a more appealing balance.





Histogram Equalization

- Calculation

Let f be a given image represented as a m_r by m_c matrix of integer pixel intensities ranging from 0 to $L - 1$. L is the number of possible intensity values, often 256. Let p denote the normalized histogram of f with a bin for each possible intensity. So

$$p_n = \frac{\text{number of pixels with intensity } n}{\text{total number of pixels}} \quad n = 0, 1, \dots, L - 1.$$

The histogram equalized image g will be defined by

$$g_{i,j} = \text{floor}\left((L - 1) \sum_{n=0}^{f_{i,j}} p_n\right), \quad (1)$$

where $\text{floor}()$ rounds down to the nearest integer. This is equivalent to transforming the pixel intensities, k , of f by the function

$$T(k) = \text{floor}\left((L - 1) \sum_{n=0}^k p_n\right).$$

	$f_{i,j}$	



	$g_{i,j}$	



Histogram Equalization

• Motivation

The motivation for this transformation comes from thinking of the intensities of f and g as continuous random variables X, Y on $[0, L-1]$ with Y defined by

$$Y = T(X) = (L-1) \int_0^X p_X(x) dx, \quad (2)$$

where p_X is the probability density function of f . T is the cumulative distributive function of X multiplied by $(L-1)$. Assume for simplicity that T is differentiable and invertible. It can then be shown that Y defined by $T(X)$ is uniformly distributed on $[0, L-1]$ namely that $p_Y(y) = \frac{1}{L-1}$.

变上、下限积分求导公式

$$\frac{d}{dx} \int_a^{u(x)} f(t) dt = f(u(x)) \cdot u'(x)$$

$$\frac{d}{dx} \int_{\phi(x)}^b f(t) dt = -f(\phi(x)) \cdot \phi'(x)$$

$$\frac{d}{dx} \int_{\phi(x)}^{u(x)} f(t) dt = f(u(x)) \cdot u'(x) - f(\phi(x)) \cdot \phi'(x)$$

$$[f^{-1}(x)]' = \frac{1}{f'[f^{-1}(x)]}$$

$$\begin{aligned} \int_0^y p_Y(z) dz &= \text{probability that } 0 \leq Y \leq y \\ &= \text{probability that } 0 \leq X \leq T^{-1}(y) \\ &= \int_0^{T^{-1}(y)} p_X(w) dw \end{aligned}$$

$$\frac{d}{dy} \left(\int_0^y p_Y(z) dz \right) = p_Y(y) = p_X(T^{-1}(y)) \frac{d}{dy} (T^{-1}(y)).$$

Note that $\frac{dT}{dy} T^{-1}(y) = \frac{dT}{dx} y = 1$, so


$$\frac{dT}{dx} \bigg|_{x=T^{-1}(y)} \frac{d}{dy} (T^{-1}(y)) = (L-1) p_X(T^{-1}(y)) \frac{d}{dy} (T^{-1}(y)) = 1,$$

which means $p_Y(y) = \frac{1}{L-1}$.



Locally **Adaptive** Histogram Equalization

- **Subdivide** the image into blocks and perform separate histogram equalization in each sub-block
- Re-compute the histogram for **every** block centered at each pixel
- Adaptive histogram equalization:
 - Compute **non-overlapped** block-based equalization functions
 - Smoothly **interpolate** the transfer functions



$$f_{s,t}(I) = (1-s)(1-t)f_{00}(I) + s(1-t)f_{10}(I) + (1-s)t f_{01}(I) + st f_{11}(I)$$

Distribute each input pixel into four adjacent lookup tables during the histogram accumulation phase

Linear Filtering



Linear Filtering

- (a) original image
- --(b) blurred
- --(c) sharpened
- --(d) smoothed with **edge-preserving** filter
- (e) binary image
- --(f) dilated
- --(g) distance transform
- --(h) connected components



(a)




(b)




(c)



(d)



(e)



(f)



(g)



(h)



Linear Filtering

- Linear filter: **correlation** operator
 - The entries in the weight kernel or mask are the filter coefficients

$$g(i, j) = \sum_{k, l} f(i + k, j + l) h(k, l) \quad g = f \otimes h$$

45	60	98	127	132	133	137	133
46	65	98	123	126	128	131	133
47	65	96	115	119	123	135	137
47	63	91	107	113	122	138	134
50	59	80	97	110	123	133	134
49	53	68	83	97	113	128	133
50	50	58	70	84	102	116	126
50	50	52	58	69	86	101	120

$f(x, y)$

*

0.1	0.1	0.1
0.1	0.2	0.1
0.1	0.1	0.1

$h(x, y)$

=

69	95	116	125	129	132
68	92	110	120	126	132
66	86	104	114	124	132
62	78	94	108	120	129
57	69	83	98	112	124
53	60	71	85	100	114

$g(x, y)$



Linear Filtering

- Linear filter: **convolution** operator

$$g(i, j) = \sum_{k, l} f(i - k, j - l) h(k, l) = \sum_{k, l} f(k, l) h(i - k, j - l)$$

$$g = f * h$$

- Additional nice properties, e.g., it is both **commutative** and **associative**
- The Fourier transform of **two convolved images** is the **product** of their individual Fourier transforms
- The continuous version of convolution

$$g(\mathbf{x}) = \int f(\mathbf{x} - \mathbf{u}) h(\mathbf{u}) d\mathbf{u}.$$



Linear Filtering

- Linear **shift-invariant** (LSI) operators

- Correlation operator
- Convolution operator

- The **superposition** principle $h \circ (f_0 + f_1) = h \circ f_0 + h \circ f_1$

- **Shift** invariance principle

$$g(i, j) = f(i + k, j + l) \Leftrightarrow (h \circ g)(i, j) = (h \circ f)(i + k, j + l)$$

- Matrix-vector multiply (remember cross product?)

$$g = Hf$$

$$\begin{bmatrix} 72 & 88 & 62 & 52 & 37 \end{bmatrix} * \begin{bmatrix} 1/4 & 1/2 & 1/4 \end{bmatrix} \Leftrightarrow \frac{1}{4} \begin{bmatrix} 2 & 1 & \cdot & \cdot & \cdot \\ 1 & 2 & 1 & \cdot & \cdot \\ \cdot & 1 & 2 & 1 & \cdot \\ \cdot & \cdot & 1 & 2 & 1 \\ \cdot & \cdot & \cdot & 1 & 2 \end{bmatrix} \begin{bmatrix} 72 \\ 88 \\ 62 \\ 52 \\ 37 \end{bmatrix}$$



Padding (border effects)

- How to keep the same size?
 - zero: set all pixels outside the source image to 0
 - constant (border color): set all pixels to a **specified** value;
 - clamp: **repeat** edge pixels indefinitely;
 - (cyclic) wrap (repeat or tile): loop "around" the image in a "toroidal" configuration;
 - mirror: **reflect** pixels across the image edge;
 - extend: **extend** the signal by subtracting the **mirrored** version of the signal from the edge pixel value.




Padding (border effects)

- The effects of padding an image
- Padding is an important step to design a NN
 - Padding
 - Size of filter
 - Stride



zero



wrap



clamp



mirror



blurred zero



normalized zero



blurred clamp



blurred mirror

The normalized zero image is the result of dividing (normalizing) the blurred zero padded RGBA image by its corresponding soft alpha value.



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10							

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20						

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30					

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30				

1	1	1
1	1	1
1	1	1

$\frac{1}{9}$

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30				

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30				
						?			
				50					

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Linear Filtering: An Example

$f[.,.]$

$h[.,.]$

$g[.,.] \frac{1}{9}$

1	1	1
1	1	1
1	1	1

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m,n] = \sum_{k,l} g[k,l] f[m+k,n+l]$$



Separable Filtering

- Problem
 - Not efficient: performing a convolution requires K^2 (multiply-add) operations **per pixel**
- Solution: separable filtering
 - First performing a one-dimensional **horizontal** convolution
 - Then being followed by a one-dimensional **vertical** convolution
- Why it works: outer product of the two kernels
 - $2K$ operations per pixel

$$K = vh^T$$

$$K = \sum_i \sigma_i u_i v_i^T$$



Separable Filtering

- Examples

$$\frac{1}{K^2}$$

1	1	...	1
1	1	...	1
\vdots	\vdots	1	\vdots
1	1	...	1

$$\frac{1}{16}$$

1	2	1
2	4	2
1	2	1

$$\frac{1}{256}$$

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

$$\frac{1}{8}$$

-1	0	1
-2	0	2
-1	0	1

$$\frac{1}{4}$$

1	-2	1
-2	4	-2
1	-2	1

$$\frac{1}{K}$$

1	1	...	1
---	---	-----	---

$$\frac{1}{4}$$

1	2	1
---	---	---

$$\frac{1}{16}$$

1	4	6	4	1
---	---	---	---	---

$$\frac{1}{2}$$

-1	0	1
----	---	---

$$\frac{1}{2}$$

1	-2	1
---	----	---



(a) box, $K = 5$



(b) bilinear



(c) "Gaussian"



(d) Sobel



(e) corner

First
derivative
image



Band-Pass

- Gaussian filter

$$G(x, y; \sigma) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

- **Second** derivative of a two-dimensional image


$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

- Discrete Laplacian filters

- Used to find areas of rapid change (edges) in images
- Finite-difference method or by the finite-element method

$$\begin{pmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{pmatrix}$$

(a) original image of Einstein;
(b) orientation map computed from the **second-order** oriented energy;
(c) Original image with oriented structures **enhanced**.



(a)



(b)



(c)



Band-Pass

- Problem


- Laplacian filter is sensitive to noise
- Common to smooth the image (e.g., using a Gaussian filter) before applying the Laplacian

- This two-step process

- Laplacian of Gaussian (LoG) filter

$$\nabla^2 G(x, y; \sigma) = \left(\frac{x^2 + y^2}{\sigma^4} - \frac{2}{\sigma^2} \right) G(x, y; \sigma)$$

0	0	3	2	2	2	3	0	0
0	2	3	5	5	5	3	2	0
3	3	5	3	0	3	5	3	3
2	5	3	-12	-23	-12	3	5	2
2	5	0	-23	-40	-23	0	5	2
2	5	3	-12	-23	-12	3	5	2
3	3	5	3	0	3	5	3	3
0	2	3	5	5	5	3	2	0
0	0	3	2	2	2	3	0	0





Steerable Filters

- Sobel operator

$$\begin{aligned} \mathbf{G}_x &= \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} * \mathbf{A} & \mathbf{G}_y &= \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} * \mathbf{A} \\ \mathbf{G}_x &= \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * ([-1 \ 0 \ +1] * \mathbf{A}) & \mathbf{G}_y &= \begin{bmatrix} -1 \\ 0 \\ +1 \end{bmatrix} * ([1 \ 2 \ 1] * \mathbf{A}) \end{aligned}$$

- Directional or oriented filter

- Smoothing with a Gaussian
- Taking a **directional derivative**

$$\hat{u} \cdot \nabla (G * f) = \nabla_{\hat{u}} (G * f) = (\nabla_{\hat{u}} G) * f \quad \hat{u} = (\cos \theta, \sin \theta)$$

- The smoothed directional derivative filter

$$G_{\hat{u}} = uG_x + vG_y = u \frac{\partial G}{\partial x} + v \frac{\partial G}{\partial y} \quad \leftarrow \text{a steerable filter}$$



Steerable Filters

- Steps

- First convolving with the pair of filters
- Steering the filter by multiplying this gradient with a unit vector

- Advantage

- A whole family of filters can be evaluated with very little cost


- A directional **second** derivative filter

$$G_{\hat{u}\hat{u}} = u^2 G_{xx} + 2uv G_{xy} + v^2 G_{yy}$$

- Steerable filters are often used

- Feature descriptors and edge detectors

Example of a
two-dimensional
Gabor filter





Summed Area Table: Integral Image

- The running sum of all the pixel values from the origin

$$s(i, j) = \sum_{k=0}^i \sum_{l=0}^j f(k, l) \quad s(i, j) = s(i-1, j) + s(i, j-1) - s(i-1, j-1) + f(i, j)$$

$$S(i_0 \dots i_1, j_0 \dots j_1) = s(i_1, j_1) - s(i_1, j_0 - 1) - s(i_0 - 1, j_1) + s(i_0 - 1, j_0 - 1)$$

3	2	7	2	3
1	5	1	3	4
5	1	3	5	1
4	3	2	1	6
2	4	1	4	8

(a) $S = 24$

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

(b) $s = 28$

3	5	12	14	17
4	11	19	24	31
9	17	28	38	46
13	24	37	48	62
15	30	44	59	81

(c) $S = 24$

More Neighborhood Operators (Non-linear)



Median Filtering

- Select the **median** from neighborhood
 - Filter away shot noise
 - It is **not** as efficient (rank) as Gaussian
- Alpha-trimmed **mean**
 - Average all of the pixels except for the alpha-fraction that are the smallest and the largest
- Weighted **median**
 - Each pixel is used a **number** of times depending on its **distance** from the center
 - Equivalent to minimizing the weighted objective function

1	2	1	2	4
2	1	3	5	8
1	3	7	6	9
3	4	8	6	7
4	5	7	8	9

(a) median = 4

1	2	1	2	4
2	1	3	5	8
1	3	7	6	9
3	4	8	6	7
4	5	7	8	9

(b) α -mean = 4.6

$$\sum_{k,l} w(k,l) |f(i+k, j+l) - g(i,j)|^p$$



Bilateral Filtering

- A bilateral filter is a **non-linear**, **edge-preserving**, and noise-reducing smoothing filter for images
- It replaces the intensity of each pixel with a **weighted average** of intensity values from nearby pixels.

$$g(i, j) = \frac{\sum_{k, l} f(k, l) w(i, j, k, l)}{\sum_{k, l} w(i, j, k, l)} \quad w(i, j, k, l) = \exp \left(-\frac{(i - k)^2 + (j - l)^2}{2\sigma_d^2} - \frac{\|f(i, j) - f(k, l)\|^2}{2\sigma_r^2} \right)$$
$$d(i, j, k, l) = \exp \left(-\frac{(i - k)^2 + (j - l)^2}{2\sigma_d^2} \right) \quad r(i, j, k, l) = \exp \left(-\frac{\|f(i, j) - f(k, l)\|^2}{2\sigma_r^2} \right)$$

Diagram illustrating the components of the bilateral filter weight function $w(i, j, k, l)$. The weight function is the product of a domain kernel $d(i, j, k, l)$ and a range kernel $r(i, j, k, l)$. Blue arrows point from the domain kernel term in the weight function to the domain kernel equation, and from the range kernel term to the range kernel equation.

product of a domain kernel

data-dependent range kernel



Bilateral Filtering

- This weight can be based on a Gaussian distribution. Crucially, the weights depend not only on **Euclidean distance** of pixels, but also on the **radiometric differences** (e.g., range differences, such as color intensity, depth distance, etc.). This preserves sharp edges.

	2	1	0	1	2
2	0.1	0.3	0.4	0.3	0.1
1	0.3	0.6	0.8	0.6	0.3
0	0.4	0.8	1.0	0.8	0.4
1	0.3	0.6	0.8	0.6	0.3
2	0.1	0.3	0.4	0.3	0.1

(c) domain filter

0.0	0.0	0.0	0.0	0.2
0.0	0.0	0.0	0.4	0.8
0.0	0.0	1.0	0.8	0.4
0.0	0.2	0.8	0.8	1.0
0.2	0.4	1.0	0.8	0.4

(d) range filter





An Example



(a)



(b)



(c)



(d)



(e)



(f)



(g)



(h)

(a) original image with **Gaussian** noise; (b) Gaussian filtered; (c) median filtered; (d) bilaterally filtered;
(e) original image with **shot** noise; (f) Gaussian filtered; (g) median filtered; (h) bilaterally filtered

Note that the bilateral filter fails to remove the shot noise because the **noisy pixels are too different from their neighbors.**



Morphology

- Binary images

- First convolve the binary image with a **binary structuring element**
 - ✓ It can be any shape, from a simple 3×3 box filter s

$$c = f \otimes s$$

- Then select a **binary** output value depending on the **thresholded** result of the convolution

$$\theta(f, t) = \begin{cases} 1 & \text{if } f \geq t, \\ 0 & \text{else,} \end{cases}$$

- **dilation:** $\text{dilate}(f, s) = \theta(c, 1);$

- **erosion:** $\text{erode}(f, s) = \theta(c, S);$ S number of pixels

- **majority:** $\text{maj}(f, s) = \theta(c, S/2);$

- **opening:** $\text{open}(f, s) = \text{dilate}(\text{erode}(f, s), s);$

- **closing:** $\text{close}(f, s) = \text{erode}(\text{dilate}(f, s), s)$



Morphology

- The structuring element for all examples is a 5*5 square



(a)



(b)



(c)



(d)



(e)



(f)

(a) original image; (b) dilation; (c) erosion; (d) majority; (e) opening; (f) closing



Distance Transforms (Binary Image)

- The distance transform

- Quickly pre-computing the distance to a **curve** or set of points

✓ **Manhattan** distance

$$d_1(k, l) = |k| + |l|$$

✓ **Euclidean** distance

$$d_2(k, l) = \sqrt{k^2 + l^2}$$

- How to calculate distances to the *nearest* background pixel

$$D(i, j) = \min_{k, l: b(k, l)=0} d(i - k, j - l)$$

- Forward and backward pass of a simple raster-scan algorithm

0	0	0	0	1	0	0
0	0	1	1	1	0	0
0	1	1	1	1	1	0
0	1	1	1	1	1	0
0	1	1	1	0	0	0
0	0	1	0	0	0	0
0	0	0	0	0	0	0

0	0	0	0	1	0	0
0	0	1	1	2	0	0
0	1	2	2	3	1	0
0	1	2	3			

0	0	0	0	1	0	0
0	0	1	1	2	0	0
0	1	2	2	3	1	0
0	1	2	2	1	1	0
0	1	2	1	0	0	0
0	0	1	0	0	0	0
0	0	0	0	0	0	0

0	0	0	0	1	0	0
0	0	1	1	1	0	0
0	1	2	2	2	1	0
0	1	2	2	1	1	0
0	1	2	1	0	0	0
0	0	1	0	0	0	0
0	0	0	0	0	0	0



Connected Components

- Connected components
 - Define regions of adjacent pixels that have the **same input** value
 - Consider pixels to be **adjacent** if they are immediate **N4 neighbors** and they have the same input value
- Connected components can be used to
 - Finding individual letters
 - Finding objects
 - Compute their area statistics
 - ✓ The area (number of pixels)
 - ✓ The perimeter (number of boundary pixels)
 - ✓ The centroid (average x and y values)
 - ✓ The second moments






Connected Components

- First step

- Split the image into **horizontal runs** of adjacent pixels
- Then color the runs with **unique labels**
- Re-use the labels of vertically adjacent runs whenever possible

- Second step

- Adjacent runs of different colors are then **merged**



Conclusions



Conclusions

- Point operators
 - Include **brightness** and **contrast** adjustments as well as **color** correction and transformations
- Neighborhood operators: filter images (linear)
 - Add soft **blur**, **sharpen** details, accentuate **edges**, or remove **noise**
- More neighborhood operators (non-linear)
 - Edge preserving **median**, **bilateral** filters,
 - **Morphological** operators that operate on binary images,
 - Semi-global operators that compute **distance** transforms and find **connected** components in binary images



Thanks



zhengf@sustc.edu.cn