

Computer Vision

CS308

Feng Zheng

SUSTech CS Vision Intelligence and Perception

Week 6



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



Content

- Brief Review
- Fitting Techniques
 - Least Squares
 - Total Least Squares
- Random Sample Consensus (RANSAC)
- Hough Voting
- Image Alignment

Brief Review



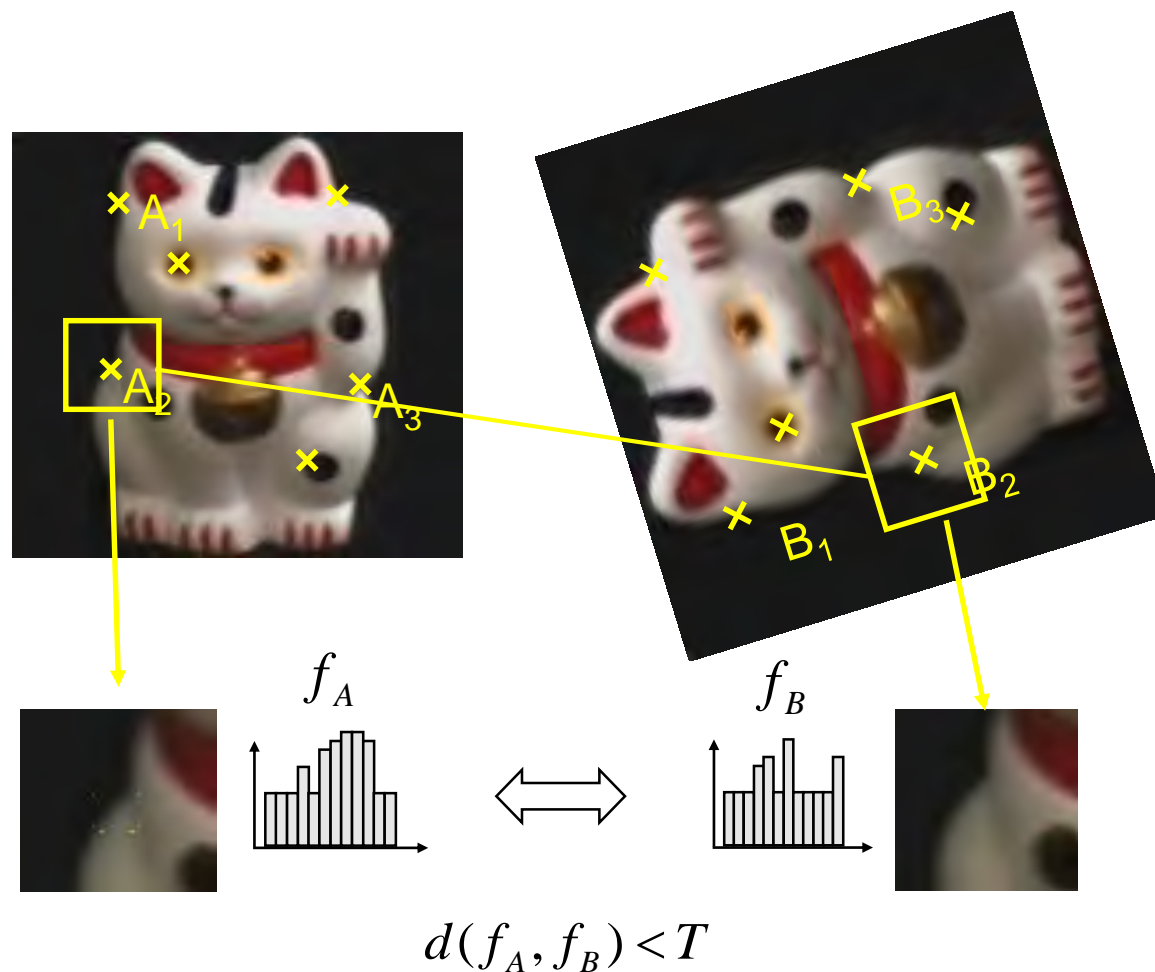
Overview of Keypoint Matching

- Steps

- **Find** a set of distinctive keypoints
- Define a **region** around each keypoint
- Compute a local **descriptor** from the region
- **Match** local descriptors

- Goals

- Detect points that are **repeatable** and **distinctive**



Fitting Techniques



How Do We Build Panorama?

- We need to match (align) images





Matching with Features

- Steps

- Detect feature points in both images






Matching with Features

- Steps

- Detect feature points in both images
- Find corresponding pairs





Matching with Features

- Steps


- Detect feature points in both images
 - Find corresponding pairs
 - Use these pairs to align images
- } Previous Lecture





Fitting: Building a Model for a Set of Features


- Choose a **parametric model** to represent **a set of features**



Simple model: **lines**



Simple model: **circles**



Complicated model: **car**



Fitting: Issues

- Case study: Line detection
 - **Noise** in the measured feature locations
 - **Extraneous** data: clutter (outliers), multiple lines
 - **Missing** data: occlusions





Fitting: Issues


- If **we know which points belong to the line**, how do we find the "optimal" line parameters?
 - Least squares
- What if there are **outliers**?
 - Robust fitting, RANSAC
- What if there are **many lines**?
 - Voting methods: RANSAC, Hough transform
- What if we're **not even sure** it's a line?
 - Model selection



Line Fitting: Ordinary Least Squares

- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = mx_i + b$
- Find (m, b) to minimize

$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$




We know which points belong to the line



Line Fitting: Ordinary Least Squares

- Data: $(x_1, y_1), \dots, (x_n, y_n)$
- Line equation: $y_i = mx_i + b$
- Find (m, b) to minimize



$$E = \sum_{i=1}^n (y_i - mx_i - b)^2$$

$$\begin{aligned} E &= \sum_{i=1}^n \left(y_i - \begin{bmatrix} x_i & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right)^2 = \left\| \begin{bmatrix} y_1 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} x_1 & 1 \\ \vdots & \vdots \\ x_n & 1 \end{bmatrix} \begin{bmatrix} m \\ b \end{bmatrix} \right\|^2 = \|Y - XB\|^2 \\ &= (Y - XB)^T (Y - XB) = Y^T Y - 2(XB)^T Y + (XB)^T (XB) \end{aligned}$$



Line Fitting: Ordinary Least Squares

- Normal equations: least squares solution to $XB=Y$

$$\frac{dE}{dB} = 2X^T XB - 2X^T Y = 0$$

$$X^T XB = X^T Y$$

- Problem with "vertical" least squares
 - Not rotation-invariant
 - Fails completely for vertical lines

$$X^T X$$




Total Least Squares

- Distance between point (x_i, y_i) and line $ax+by=d$ ($a^2+b^2=1$):

$$|ax_i + by_i - d|$$

$$E = \sum_{i=1}^n (ax_i + by_i - d)^2$$





Total Least Squares

$$\frac{\partial E}{\partial d} = \sum_{i=1}^n -2(ax_i + by_i - d) = 0 \quad \Rightarrow \quad d = \frac{a}{n} \sum_{i=1}^n x_i + \frac{b}{n} \sum_{i=1}^n y_i = a\bar{x} + b\bar{y}$$

$$E = \sum_{i=1}^n (a(x_i - \bar{x}) + b(y_i - \bar{y}))^2 = \left\| \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix} \right\|^2 = (UN)^T (UN)$$

\uparrow
 U

\uparrow
 N

$$\frac{dE}{dN} = 2(U^T U)N = 0$$

- Solution to $(U^T U)N = 0$, subject to $\|N\|^2 = 1$: eigenvector of $U^T U$ associated with the smallest eigenvalue (least squares solution to *homogeneous linear system* $UN = 0$)




Total Least Squares

- Second moment matrix

$$U = \begin{bmatrix} x_1 - \bar{x} & y_1 - \bar{y} \\ \vdots & \vdots \\ x_n - \bar{x} & y_n - \bar{y} \end{bmatrix}$$



$$U^T U = \begin{bmatrix} \sum_{i=1}^n (x_i - \bar{x})^2 & \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) \\ \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y}) & \sum_{i=1}^n (y_i - \bar{y})^2 \end{bmatrix}$$





OLS vs. TLS


- The difference between standard OLS regression and "orthogonal" TLS regression





Total Least Squares


- Robustness to **noise**:
least squares fit to the
red points





Total Least Squares


- Robustness to noise: Least squares fit with an **outlier**
- Problem: squared error **heavily** penalizes outliers





Robust Estimators


- General approach---minimize: $\sum_i \rho(r_i(x_i, \theta); \sigma)$
 $r_i(x_i, \theta)$ – **residual** of i th point
w.r.t. model parameters θ
 ρ – robust function with
scale parameter σ
- The robust function ρ
behaves like squared
distance for **small values** of
the residual u but saturates
for **larger values** of u





Choosing the Scale: Just Right


- The effect of the **outlier** is minimized, when choosing a **just right** scale





Choosing the Scale: Too **Small**


- The error value is almost the same for every point and the fit is very **poor**





Choosing the Scale: Too Large


- Behaves much the same as least squares





Curve Fitting

- **Find Polynomial:** $y = f(x) = ax^3 + bx^2 + cx + d$
 - That best fits the given points (x_i, y_i)



- **Minimize:** $\frac{1}{N} \sum_i [y_i - (ax_i^3 + bx_i^2 + cx_i + d)]^2$
- **Using:** $\frac{\partial E}{\partial a} = 0$, $\frac{\partial E}{\partial b} = 0$, $\frac{\partial E}{\partial c} = 0$, $\frac{\partial E}{\partial d} = 0$

- **Note:** $f(x)$ is **LINEAR** in the parameters (a, b, c, d)

Random Sample Consensus



RANSAC

- Robust fitting (TLS) can deal with **a few** outliers - what if we have very **many**?
- Random sample consensus (RANSAC): Very general framework for model fitting in the **presence of outliers**
- Outline
 - Choose a **small subset** of points uniformly at random
 - Fit **a model** to that subset
 - Find all remaining points that are "**close**" to the **model** and reject the rest as outliers
 - Do this **many** times and choose the **best** model



RANSAC for Line Fitting

- Algorithm
- Repeat N times:
 - Draw s points uniformly at random
 - Fit line to these s points (TLS)
 - Find inliers to this line among the remaining points (i.e., points whose distance from the line is less than t)
 - If there are d or more inliers, accept the line and refit using all inliers
- End
- Four parameters: s , t , d and N



Choosing the Parameters

- Initial number of points s
 - Minimum number needed to fit the model
✓ 2 points
- Distance threshold t
 - (1) Choose t so probability for inlier is p (e.g. 0.95)
 - (2) Zero-mean Gaussian noise with standard deviation σ : $t^2 = 3.84\sigma^2$



Choosing the Parameters

- Number of times N
 - Choose N so that, with probability p , at least one random sample is free from outliers (e.g. $p=0.99$)

Desired success rate after N times: p

Outlier ratio (Unknown): e

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$

N	proportion of outliers e						
s	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177




Choosing the Parameters

- Consensus set size d (number of inliers)
 - Should match expected **inlier ratio**

$$\left(1 - (1 - e)^s\right)^N = 1 - p$$

$$N = \log(1 - p) / \log(1 - (1 - e)^s)$$





Adaptively determining the number of samples

- Inlier ratio e is often unknown a priori, so pick worst case, e.g. 50%, and adapt if more inliers are found, e.g. 80% would yield $e=0.2$
- Adaptive procedure:
 - $N=\infty$, $sample_count=0$
 - While $N > sample_count$
 - ✓ Choose a sample (fitting) and count the number of inliers
 - ✓ Set $e = 1 - (\text{number of inliers})/(\text{total number of points})$
 - ✓ Recompute N from e :


$$N = \log(1-p) / \log(1-(1-e)^s)$$

- ✓ Increment the $sample_count$ by 1




RANSAC

- An example



A data set with many outliers for which a line has to be fitted.



Fitted line with RANSAC; outliers have no influence on the result.



RANSAC pros and cons

- Pros

- Simple and general
- Applicable to many different problems
- Often works well in practice

- Cons

- **Lots** of parameters to tune
- **Can't always** get a good initialization of the model based on the minimum number of samples
- Sometimes **too many iterations** are required
- Can fail for extremely **low inlier ratios**
- We can often do better than **brute-force sampling**

Hough transform



Voting Schemes

- Principal of voting
 - Let **each** feature (voter) vote for **all** the models that are compatible with it
 - Hopefully the **noise** features (voter) will **not vote consistently** for **any** single model (nominator)
 - **Missing data doesn't matter** as long as there are **enough** features remaining to agree on a good model



Hough Transform

- An early type of voting scheme
- General outline:
 - Discretize parameter space into bins
 - For **each** feature point in the image, put a **vote** in every bin in the parameter space that **could have generated this point**
 - Find bins that have the **most** votes






Image space




Hough parameter space



Parameter Space Representation


- A **line** in the image corresponds to a **point** in Hough space





Parameter Space Representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?





Parameter Space Representation

- What does a point (x_0, y_0) in the image space map to in the Hough space?
 - Answer: the solutions of $b = -x_0m + y_0$
 - This is a **line** in Hough space





Image space




Hough parameter space



Parameter Space Representation



- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?





Parameter Space Representation

- Where is the line that contains both (x_0, y_0) and (x_1, y_1) ?
 - It is the **intersection** of the lines $b = -x_0m + y_0$ and $b = -x_1m + y_1$






Parameter Space Representation

- Problems with the (m, b) space:
 - Unbounded parameter domain
 - Vertical lines require infinite m
- Alternative: polar representation

$$x \cos \theta + y \sin \theta = \rho$$




- Each point will add a sinusoid in the (θ, ρ) parameter space



Algorithm Outline

- Initialize accumulator H to all zeros
- For **each** edge point (x,y) in the image
 - For $\theta = 0$ to 180
 - $\rho = x \cos \theta + y \sin \theta$
 - $H(\theta, \rho) = H(\theta, \rho) + 1$
 - end
- end
- Find the value(s) of (θ, ρ) where $H(\theta, \rho)$ is a **local maximum**
 - The detected line in the image is given by
$$\rho = x \cos \theta + y \sin \theta$$

H: accumulator array (votes)





Basic Illustration

- A line




Image space



Votes

Horizontal axis is θ
Vertical is ρ .



Basic Illustration

- A line with noise





Image space



Votes



Basic Illustration

- Scattered points





Image space



Votes



Mechanics of the Hough transform


- Difficulties
 - How big should the **cells** be? (too big, and we merge quite different lines; too small, and noise causes lines to be missed)
- How many lines?
 - Count the **peaks** in the Hough array
 - Treat **adjacent** peaks as a single peak
- Which points belong to each line?
 - Search for points close to the line
 - Solve again for line and iterate




Real World Example




Original



Edge Detection



Parameter Space



Found Lines



Finding Circles by Hough Transform



- Equation of Circle:

$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- If radius is known:

- 2D Hough Space

- Accumulator Array: $A(a, b)$





Finding Circles by Hough Transform

- Equation of Circle:


$$(x_i - a)^2 + (y_i - b)^2 = r^2$$

- If radius is not known:

- 3D Hough space!

- Use Accumulator array: $A(a, b, r)$

- What is the surface in the Hough space?





Finding Circles by Hough Transform

- Hough transform for circles






Image space



Hough parameter space




Finding Coins

Original



Edges (note noise)






Finding Coins

- Note that because the quarters and penny are **different sizes**, a different Hough transform (with separate accumulators) was used for each circle size

Penny


Quarters





Generalized Hough Transform



- We want to find a fixed shape (known) defined by its boundary points and a reference point





Generalized Hough Transform

- We want to find **a fixed shape (known)** defined by its boundary points and a reference point
- For every boundary point p , we can compute the **displacement** vector $r = a - p$ as a function of **gradient** orientation θ






Generalized Hough Transform

- Construct a model for a shape:
 - Construct a table indexed by θ storing displacement vectors r as function of **gradient direction**
- Detect using the model
 - For each **edge point** p with gradient orientation θ :
 - ✓ Retrieve **all** r indexed with θ
 - ✓ For **each** $r(\theta)$, put a vote in the Hough space at $p + r(\theta)$
 - **Peak** in this Hough space is **reference** point with most supporting edges
- Assumption: translation is the only transformation here, i.e., orientation and scale are fixed



Example: a Known and Fixed Shape


- **Model** shape
 - Gradient orientation
 - No rotation





Example: Building a Table

- Displacement vectors for model points




r

θ



Example: Detection

- Range of voting locations for **test** point




r

θ		



Example: Detection


- Range of voting locations for test point





Example: Detection

- **Votes** for points with $\theta = \uparrow$




r

θ		



Example: Building a Table

- Displacement vectors for model points




r

θ		



Example: Detection

- Range of **voting** locations for test point



θ


r		



Example: Detection

- Votes for points with

$$\theta = \swarrow$$




r

θ




Application in Recognition

- Instead of indexing **displacements** by gradient orientation, index by "**visual codeword**"



training image

What is the codeword?



visual codeword with
displacement vectors



Application in Recognition

- Instead of indexing displacements by gradient orientation, index by "visual codeword"




test image

Image Alignment



Image Alignment


- Two broad approaches:
 - Direct (pixel-based) alignment
 - ✓ Search for alignment where most pixels agree
 - Feature-based alignment
 - ✓ Search for alignment where *extracted features* agree
 - ✓ Can be verified using pixel-based alignment





Alignment as Fitting

- Previously: fitting a model to features in **one** image





Find model M that minimizes

$$\sum_i \text{residual}(x_i, M)$$



Alignment as Fitting

- Alignment: fitting a model to a transformation between pairs of features (*matches*) in **two** images




Find transformation T
that minimizes

$$\sum_i \text{residual}(T(x_i), x'_i)$$




2D Transformation Models


- Similarity
(translation, scale, rotation)



- Affine




- Projective
(homography)





Affine Transformations


- Simple fitting procedure (linear least squares)
- Approximates **viewpoint changes** for roughly planar objects and roughly orthographic cameras
- Can be used to initialize fitting for more complex models





Affine Transformations

- Assume we know the **correspondences (???)**, how do we get the **transformation**?



$$\begin{bmatrix} x'_i \\ y'_i \end{bmatrix} = \begin{bmatrix} m_1 & m_2 \\ m_3 & m_4 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \end{bmatrix} + \begin{bmatrix} t_1 \\ t_2 \end{bmatrix}$$



$$\begin{bmatrix} \dots & \dots & \dots & \dots & \dots & \dots \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & \dots & \dots & \dots & \dots & \dots \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$




Affine Transformations

- Linear system with **six** unknowns
- Each match gives us **two linearly independent equations**: need at least **three** to solve for the transformation parameters

$$\begin{bmatrix} \dots & & & & & \\ x_i & y_i & 0 & 0 & 1 & 0 \\ 0 & 0 & x_i & y_i & 0 & 1 \\ \dots & & & & & \end{bmatrix} \begin{bmatrix} m_1 \\ m_2 \\ m_3 \\ m_4 \\ t_1 \\ t_2 \end{bmatrix} = \begin{bmatrix} \dots \\ x'_i \\ y'_i \\ \dots \end{bmatrix}$$




Feature-Based Alignment Outline



- Extract features




Feature-Based Alignment Outline



- Extract features
- Compute *putative matches*




Feature-Based Alignment Outline



- Extract features
- Compute *putative matches*
- Loop:
 - *Hypothesize* transformation T



Feature-Based Alignment Outline



- Extract features
- Compute *putative* matches
- Loop:
 - *Hypothesize* transformation T
 - *Verify* transformation (search for other matches consistent with T)



Feature-Based Alignment Outline



- Extract features
- Compute *putative matches*
- Loop:
 - Hypothesize transformation T
 - Verify transformation (search for other matches consistent with T)

Conclusions



Conclusion

- Fitting techniques
 - Least Squares
 - Total Least Squares
- RANSAC
- Hough Voting
- Alignment as a fitting problem



Thanks



zhengf@sustc.edu.cn