

---

# **MACHINE LEARNING**

## **CHAPTER 0: INTRODUCTION**

---

# Contact Information

---

**Instructor:** Qi Hao

**E-mail:** hao.q@sustc.edu.cn

**Office:** South Tower Room 415

**Office Hours:** M 2:00-4:00pm

Available other times by appointment or the open door policy

**Phone:** 186-6495-7027

**QQ:** 820786590 2022机器学习

**Web:** <http://hqlab.isus.tech/teaching/CS405>

**BB:** Machine Learning Fall 2022

**OJ:** <http://oj.isus.tech/>

---

# Class Schedule

---

■ **Lectures:** T 8:00 am – 9:50 am Teaching Building I room 506

■ **Study:** F 7:00 pm – 8:50 pm Teaching Building I Room 506

■ **Grading policy:**

**Final Exam (in-class): 20%**

**Assignments (8~12 times): 20%**

**Lab Projects : 20%**

**Bonus Credits: <=5%**

**Midterm Exam (take-home): 10%**

**Quizzes (<=10 times): 10%**

**Final Projects (4 per group): 20%**

90~93: A-

80~82: B-

70~72: C-

60~62: D-

94~97: A

83~86: B

73~76: C

63~66: D

98~100: A+

87~89: B+

77~79: C+

67~69: D+

# Textbook and Lecture Notes

---

## **Textbooks:**

- [1] Pattern Recognition and Machine Learning, by Christopher M. Bishop, 2006 Springer
- [2] Machine Learning in Action, by Peter Harrington, 2012, Manning

## **Other books:**

- [1] 机器学习, 周志华
- [2] Dive in Deep Learning, by Aston Zhang, Zachary C. Lipton, Mu Li, and Alexander J. Smola
- [3] Reinforcement Learning: An Introduction, by Richard S. Sutton
- [4] The Elements of Statistical Learning, by Trevor Hstie, Rober Tibshirani, Jerome Friedman

## **Paper reading:**

- [1] Ghahramani Z. Probabilistic machine learning and artificial intelligence, Nature, 2015
- [2] Lecun Y, Bengio Y, Hinton G. Deep learning, Nature, 2015
- [3] Littman M L. Reinforcement learning improves behavior from evaluative feedback, Nature, 2015

## **Lecture notes:**

<http://hqlab.isus.tech/teaching/CS405>

---

# Other Resources

---

**Assignment platform:** `bb.sustech.edu.cn`

**Textbook resource:** <https://www.microsoft.com/en-us/research/people/cmbishop/#prml-book>

**Textbook Matlab codes :** <http://prml.github.io/>

**Matlab toolboxes :** Machine Learning, Neural Networks

---

# Teaching Objectives

---

- Fundamental knowledge about machine learning and pattern recognition, from Bayesian approaches to deep learning frameworks through lectures, quizzes and exercises
  - Machine learning system development methods in Python based platforms (numpy, scikit-learn, pytorch) through labs and projects
  - Model-based and data-driven machine learning system design and integration skills through the final project, literature surveys and reports
-

# Lecture Schedule

---

Section 0	Course Introduction	
Section 1	Preliminary	(HW1)
Section 2	Probability Distributions	(HW2)
Section 3	Linear Regression and Classification	(HW3)
Section 4*	Dimension Reduction and Feature Selection	(HW4)
Section 5	Neural Networks	(HW5)
Section 6	Sparse Kernel Machine	(HW6)
Section 7	Clustering and EM learning	(HW7)
<i>-- Midterm Exam --</i>		
Section 8*	Ensemble Learning	(HW8)
Section 9	Hidden Markov Models	(HW9)
Section 10*	Bayesian Networks	(HW10)
Section 11	Markov Decision Process	(HW11)
Section 12*	Reinforcement Learning	(HW12)
<i>-- Final Exam --</i>		

---

# Lab Schedule

---

Section 0     Lab Introduction

Section 1     Preliminary

Section 2     Bayes

Section 3     Regression and Classification

*--Final Project Proposal--*

Section 4     Decision Tree

Section 5     Random Forest (Ensemble Learning)

Section 6     KNN and Support Vector Machine

Section 7     K-Mean and EM Clustering

Section 8     Neural Network (I)

Section 9     Neural Network (II)

Section 10    Neural Network (III)

Section 11    Reinforcement Learning

*--Final Project Report--*

---



# Final Project Examples

---

- [1] Reinforcement learning based planning using a self-driving car simulator
  - [2] Segmentation of 2D/3D measurements for self-driving applications
  - [3] Detection and recognition of traffic signs for self-driving applications
  - [4] Detection and tracking of 2D/3D objects for self-driving applications
  - [5] Federated learning for model fusion of networked vehicle applications
  - [6] GNN for self-driving data augmentation
-

# Bonus Credits

---

- AI companies
- Survey Papers
- Attendance
- Bonus Credits



# Plagiarism

---

**From Spring 2018**, the plagiarism policy applied by the Computer Science and Engineering department is the following:

- \* If an assignment is found to be plagiarized, the first time the score of the assignment will be 0.
- The second time the score of the course will be 0.

**As it may be difficult** when two assignments are identical or nearly identical who actually wrote it, the policy will apply to BOTH students, unless one confesses having copied without the knowledge of the other.

---

# What is OK, and what isn't OK

---

## It's OK

- to work on an assignment with a friend, and think together about the program structure, share ideas and even the global logic. At the time of actually writing the code, you should write it alone.
- to use in an assignment a piece of code found on the web, as long as you indicate in a comment where it was found and don't claim it as your own work.
- to help friends debug their programs (you'll probably learn a lot yourself by doing so).
- to show your code to friends to explain the logic, as long as the friends write their code on their own later.

## It's NOT OK

- **to take the code of a friend, make a few cosmetic changes (comments, some variable names) and pass it as your own work.**
-

# Make a Promise to Keep

---

**Sign**

the “Student Commitment for Assignments”

**Keep**

the promise during the whole semester!

---

# Learning Objectives

---

1. What is the history of machine learning?
  2. What are the most important functionalities of machine learning?
  3. What are the major technical challenges for developing machine learning systems?
  4. What are the most useful tools for developing machine learning systems?
  5. What are the most popular software and hardware platforms for developing machine learning systems?
  6. What are the most promising applications for machine learning?
-

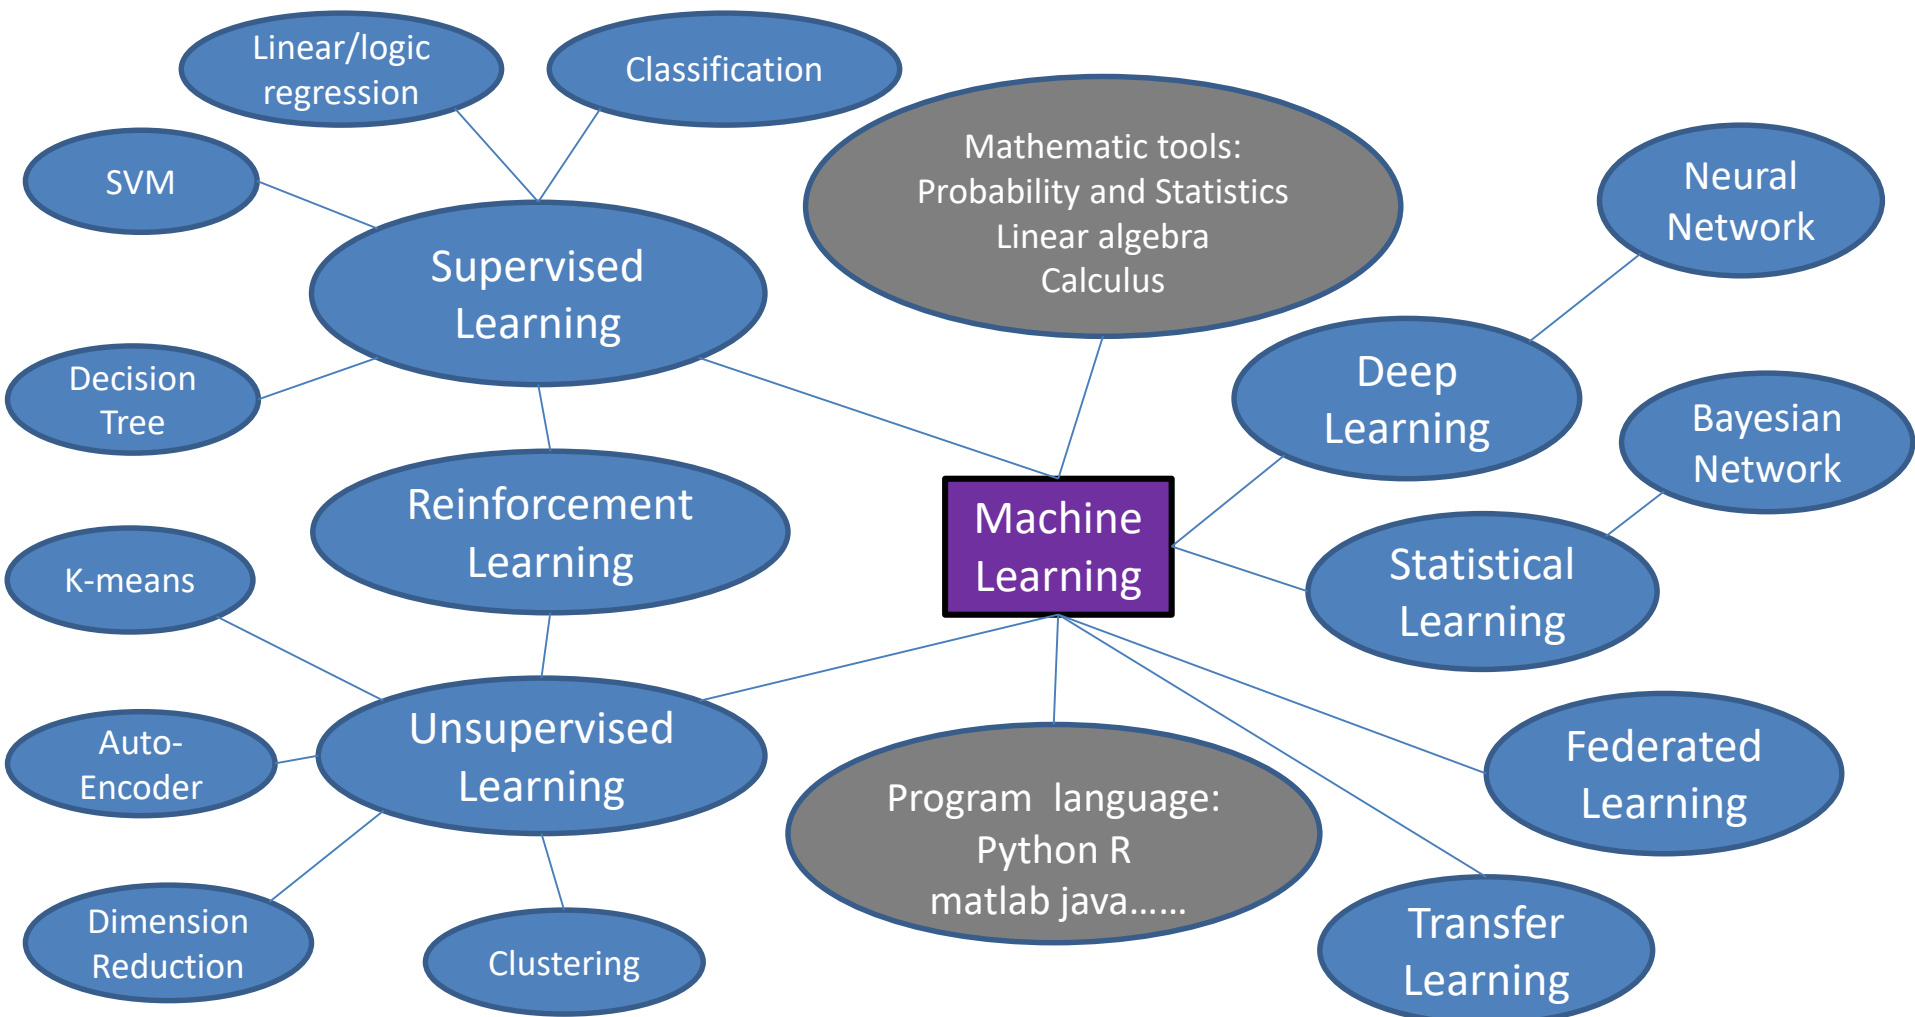
# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Framework

---





# Outlines

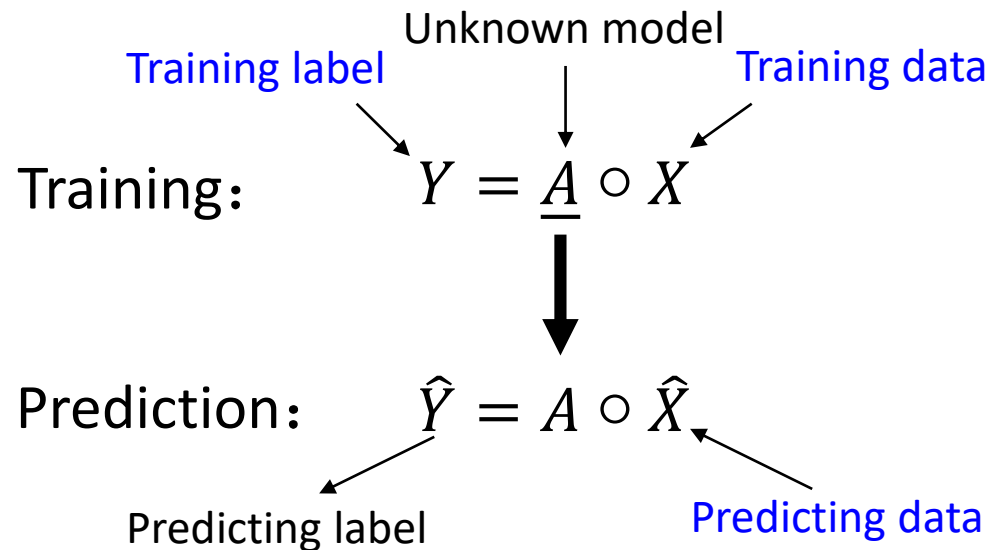
---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Problem Statement

---

- **Problem:** Predict the label  $\hat{Y}$  and data  $\hat{X}$  with training set  $(X, Y)$  ?



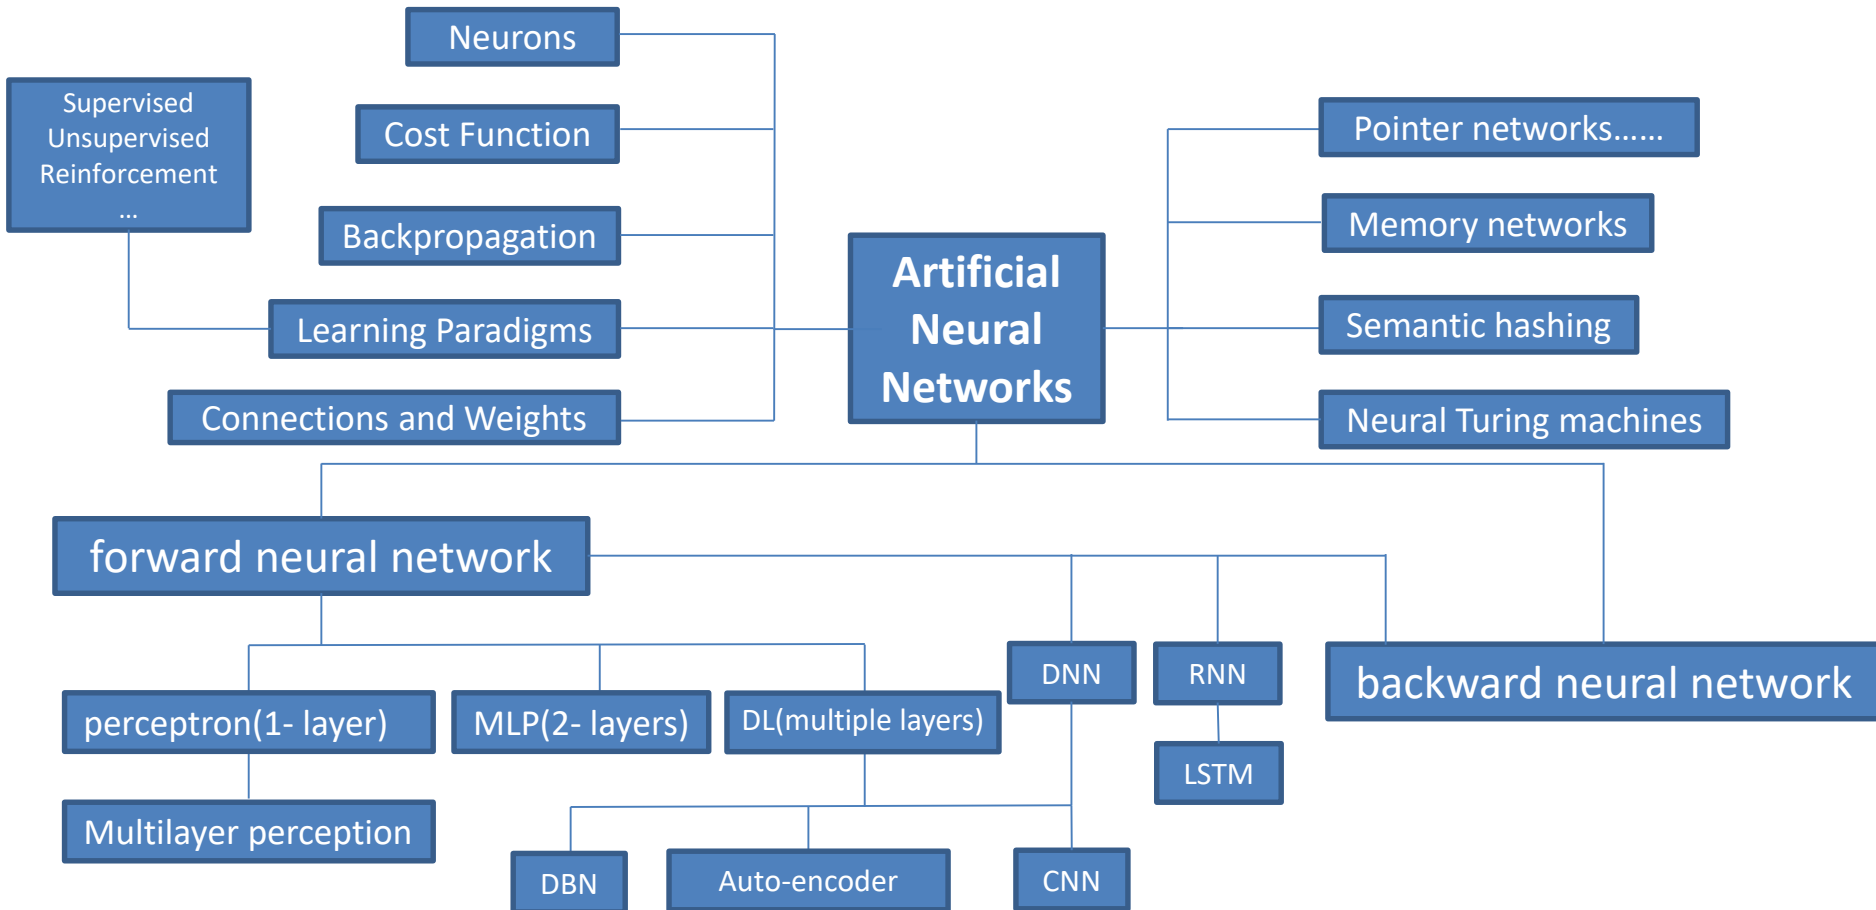
$\left\{ \begin{array}{l} Y \text{ and } X \text{ is known: supervised learning} \\ Y \text{ or } X \text{ is unknown: unsupervised learning} \end{array} \right.$   $\left\{ \begin{array}{l} Y, \hat{Y} \text{ are continuous: Regression} \\ Y, \hat{Y} \text{ are discrete: classification} \end{array} \right.$

$Y \text{ is known and } \text{Dim}(Y) > \text{Dim}(X) :$  dimensionality reduction

---

# Neural Network Models

---



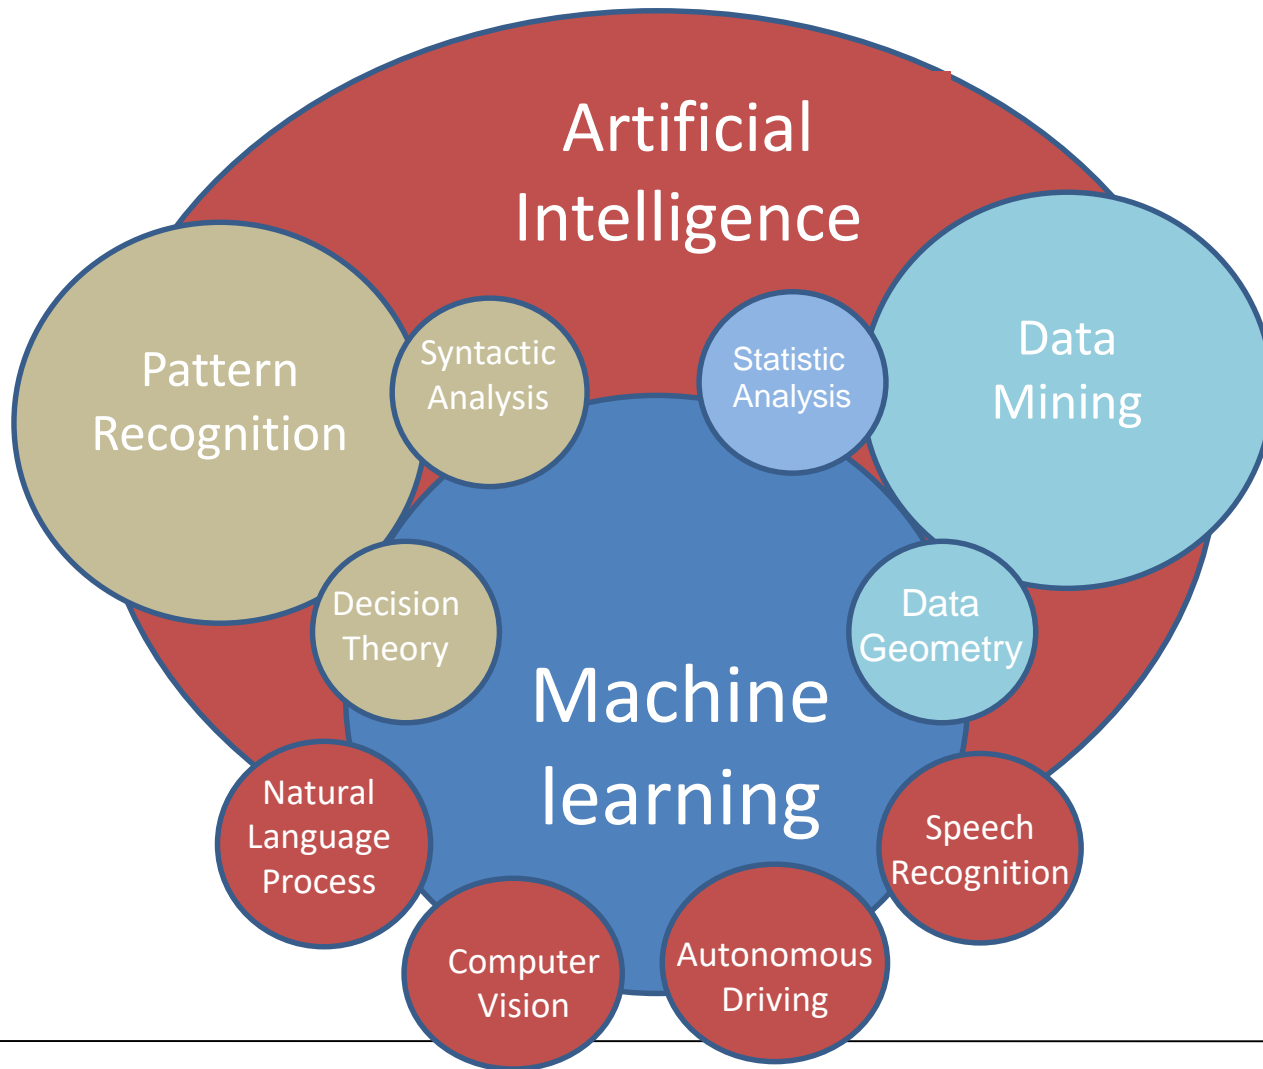
# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# The Whole Picture

---

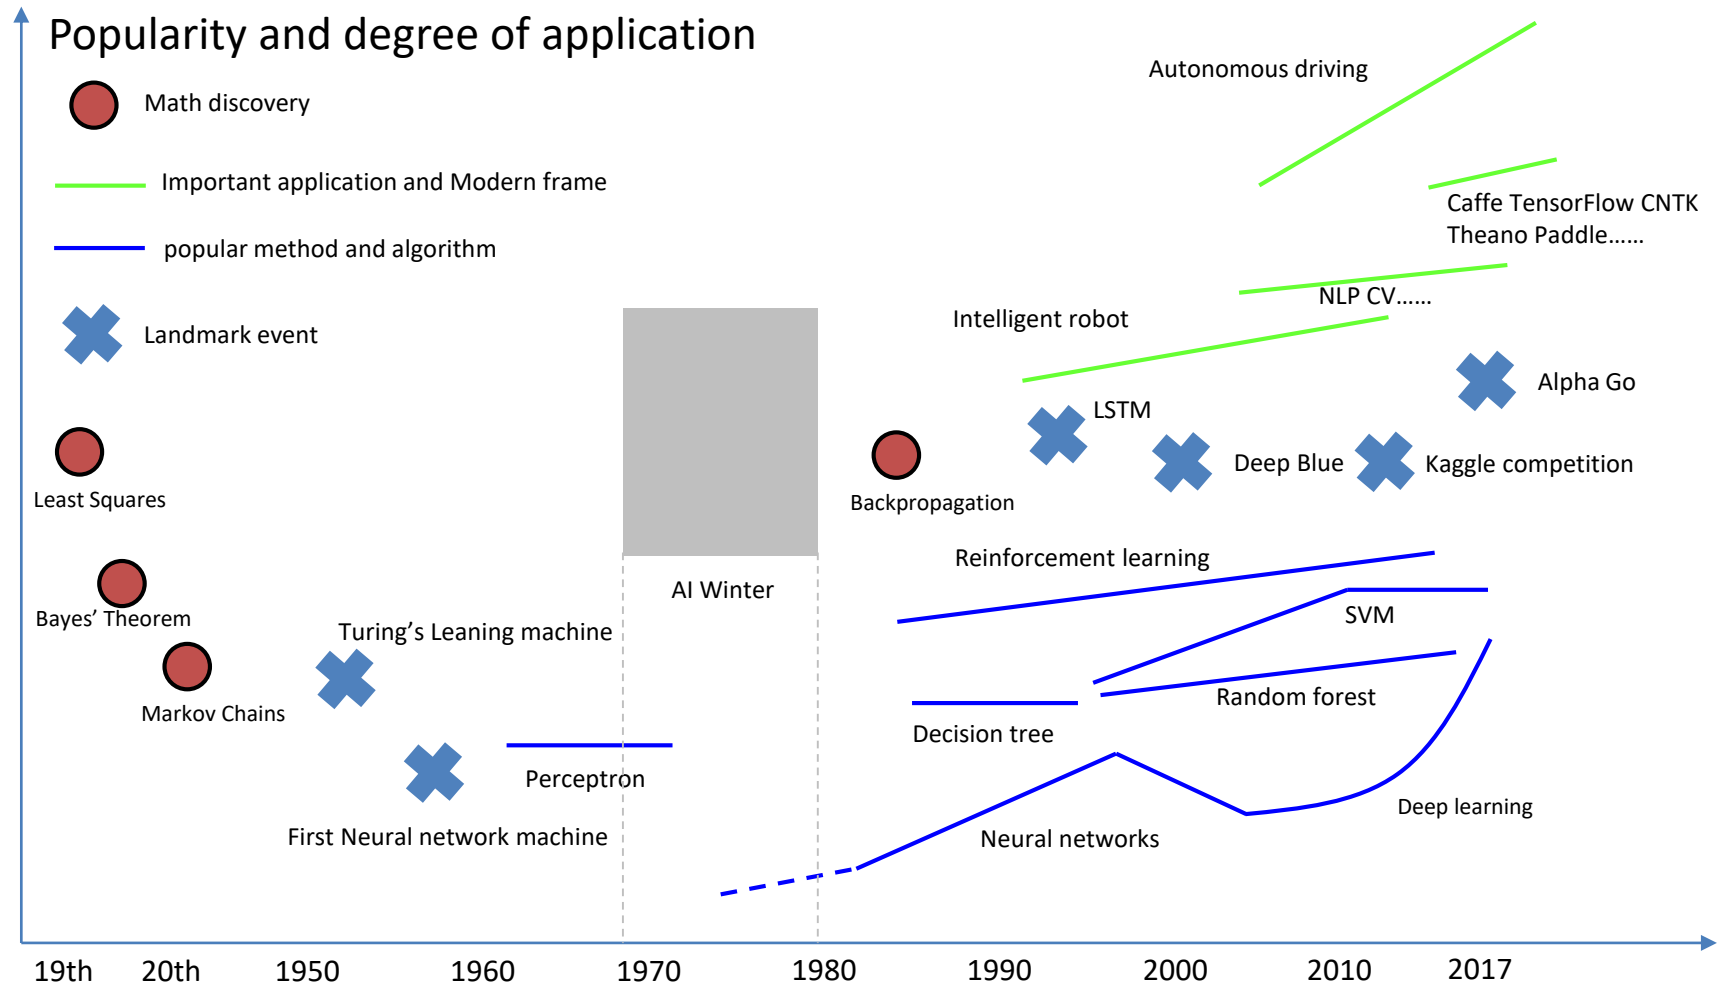


# Outlines

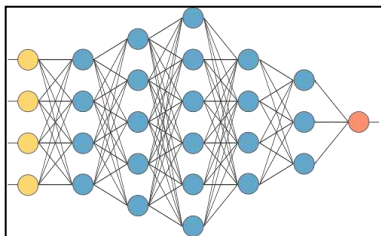
---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

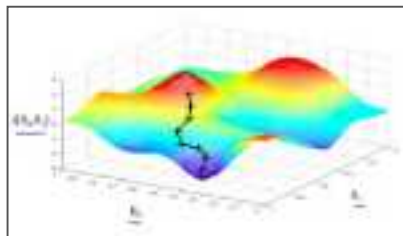
# History



# Deep ML vs Conventional ML



model



optimization



data



platform

**CNN**

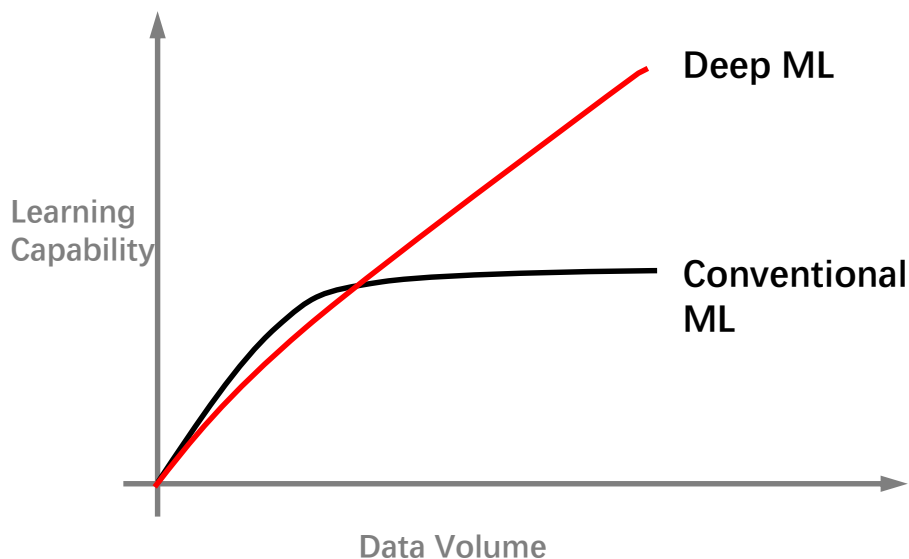
**RNN**

**LSTM**

**Auto Encoder**

**GAN**

**RL & FL**



**GPU & TPU**

**Cloud Computing**

**Data Visualization**

**GPU Service**

**TensorFlow**

**Caffe & Pytorch**



# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Machine learning

---

- **Machine Learning**—minimization of some loss function for generalizing data sets with models.
  - **Datasets** —annotated, indexed, organized
  - **Models** —tree, distance, probabilistic, graph, bio-inspired
  - **Optimization** —algorithms can minimize the loss.
-

# Datasets

---

- Collection
  - Storage
  - Annotation
  - Indexing
  - Organization
  - Access
-

# Simulators

---

- Data visualization
- Generate training data
- Algorithm evaluation



# Benchmark Metrics

---

- System functionalities
  - System scalability
  - System robustness
  - System efficiency
-

# Models

---

- Tree Models
  - Distance-based Models
  - Probabilistic Models
  - Neural Network Models
  - Graph-based Models
-

# Models

---

- Boosting
  - Mixed Models
  - Ensemble Learning
-

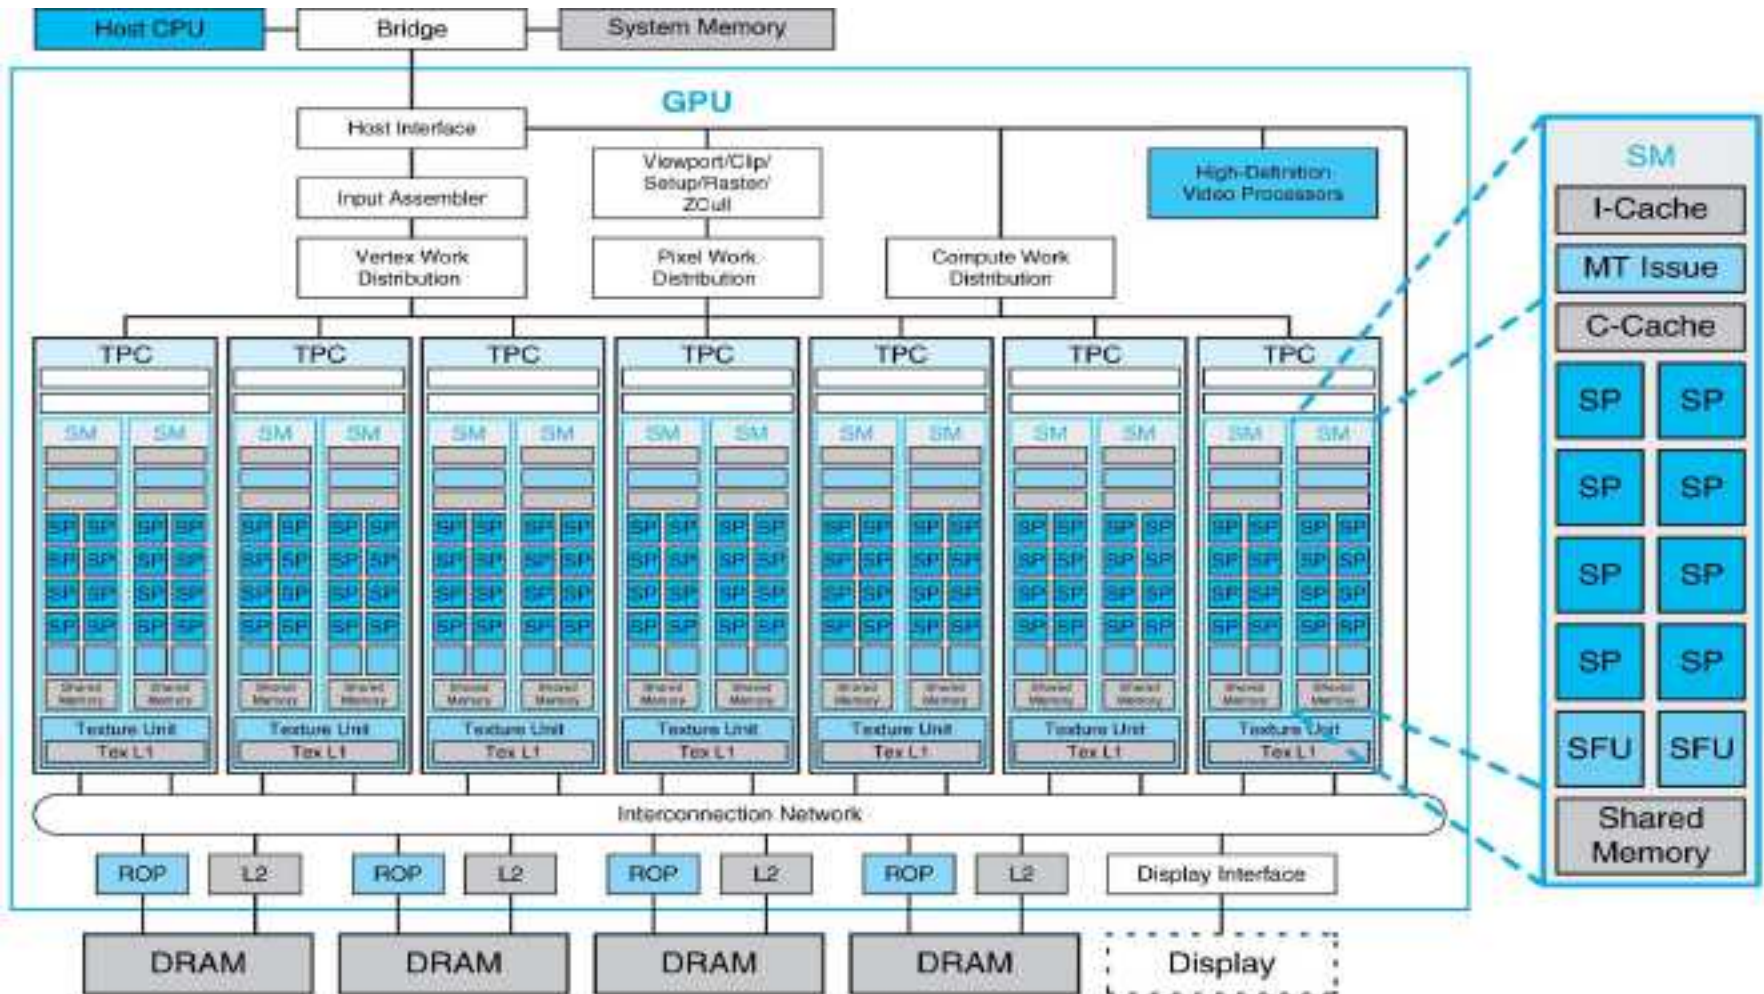
# Hardware Platform (GPU Server)

---





# Hardware Platform (GPU)



# 构筑业界最强AI算力平台，极简易用、极致性能

## 行业应用

智慧城市、制造、能源、交通、金融、运营商、教育等更多行业应用

## 应用使能



### MindX 昇腾应用使能



## AI框架



最佳匹配昇腾AI处理器算力的全场景AI计算框架

TensorFlow/PyTorch等第三方框架

可基于第三方框架开发的模型进行二次开发、训练和推理

## 异构计算架构



统一异构计算架构，释放昇腾硬件澎湃算力

## 系列硬件



端

边

云

Atlas



全流程开发工具链

MindStudio



管理运维工具

FusionDirector/SmartKit



昇腾社区

hiascend.com

# Atlas系列硬件打造人工智能算力平台基石

## Atlas训练系列硬件



Atlas 300T训练卡  
单卡算力**业界领先**

**320** TFLOPS FP16



Atlas 800  
训练服务器



Atlas 900 PoD



Atlas 900 AI集群

## Atlas推理系列硬件



Atlas 300I 推理卡



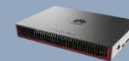
Atlas 800  
推理服务器



Atlas 200  
AI加速模块



Atlas 500  
智能小站



Atlas 200DK



Atlas 500 Pro  
智能边缘服务器

# 异构计算架构CANN，软硬协同充分释放澎湃算力



## 使能全场景

- 支持**10+**端边云设备形态
- 支持**14+**操作系统
- 支持多种AI框架



## 使能极简开发

- 统一API, AscendCL
- 自定义算子开发方式  
TBE-DSL/TBE-TIK/AI CPU



## 使能极致性能

- 亲和昇腾的图编译技术
- **1200+**高性能算子



# [M]<sup>5</sup> 开源AI框架MindSpore，构建端边云全场景生态



## 简单的开发体验

帮助开发者实现网络自动划分，只需单行到达就能实现并行训练，降低门槛，简化开发流程。



## 灵活的调试模式

具备引擎过程静态执行和动态调试能力，开发者通过变更一行代码即可切换模式，快速在线定位问题。



## 充分发挥硬件潜能

精准匹配异构处理器，最大程度发挥硬件能力，帮助开发者缩短训练时间，提升推理性能。



## 全场景快速部署

支持云、边缘和手机上的快速部署，实现更好的资源利用和数据隐私保护，让开发者专注于AI应用的创造。



## 全自动并行

静态图自动混合并行训练 **性能提升40%**

动态图优化性能 **超越业界60%**



## 全场景协同

云端分布式推理

边缘AI加速

超轻量IoT设备推理



## 全流程极简

第三方框架转换工具  
业务快速迁移

## 开发者生态

**51万+** **2300+**

下载量

社区贡献者

# 开放AI应用使能套件MindX，加速人工智能应用创新

## MindX：昇腾应用使能

MindX DL  
深度学习使能

MindX Edge  
智能边缘使能

**MindX SDK**  
行业应用开发套件

ModelZoo  
**250+**预训练模型

## MindX SDK：沉淀行业知识，使能行业应用 极简开发



2人月  
传统应用开发方式



2人天  
基于SDK开发方式

已支撑 **20+** 场景化解决方案高效开发

华为松山湖产线  
PCB板质检 散热器缺失检测

友达光电  
切片AOI检测

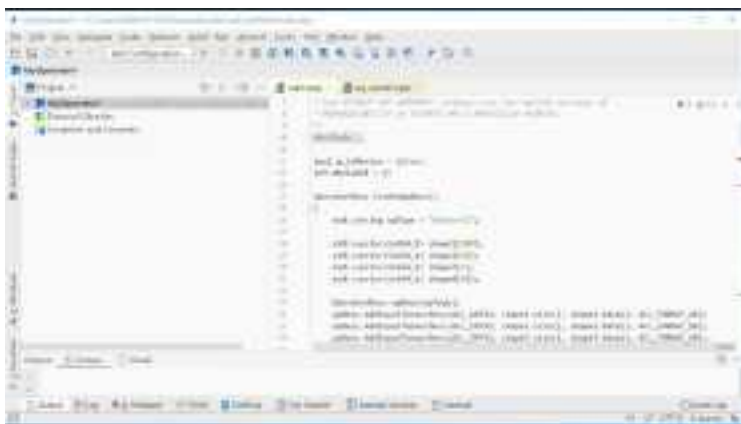
南瑞继远  
变电站 变电站  
火警检测 人员着装检测





# 一站式开发环境MindStudio，打造高效、便捷的全流程开发工具链

MindStudio是一套基于IntelliJ框架的开发工具平台。提供了应用开发、调试、模型转换功能，同时还提供了网络移植、**优化和分析等功能**，为用户开发应用程序带来了极大的便利。



- 训练脚本转换
- 模型转换
- 模型对比
- Profiling性能分析
- System Profiling工具
- AI Core Error分析工具

## 应用开发

- 基于MindX SDK开发应用
- 基于新工程开发应用

## 应用工程调试

## 模型开发

- 查询模型
- 模型训练
- 模型可视化

## 算子开发

- 查询算子
- 开发流程
- 算子分析
- 工程创建
- TBE算子开发 (TensorFlow)
- TBE算子开发 (PyTorch)
- TBE算子开发 (MindSpore)
- AI CPU算子开发 (TensorFlow)



# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-



# Machine learning and Optimization

---

- **Machine Learning**—minimization of some loss function for generalizing data sets with models.
  - **Datasets** —annotated, indexed, organized
  - **Models** —tree, distance, probabilistic, graph, bio-inspired
  - **Optimization** —algorithms can minimize the loss.
-

# What is optimization?

---

- Finding (one or more) minimizer of a function subject to constraints

$$\arg \min_x f_0(x)$$

$$s.t. f_i(x) \leq 0, i = \{1, \dots, k\}$$

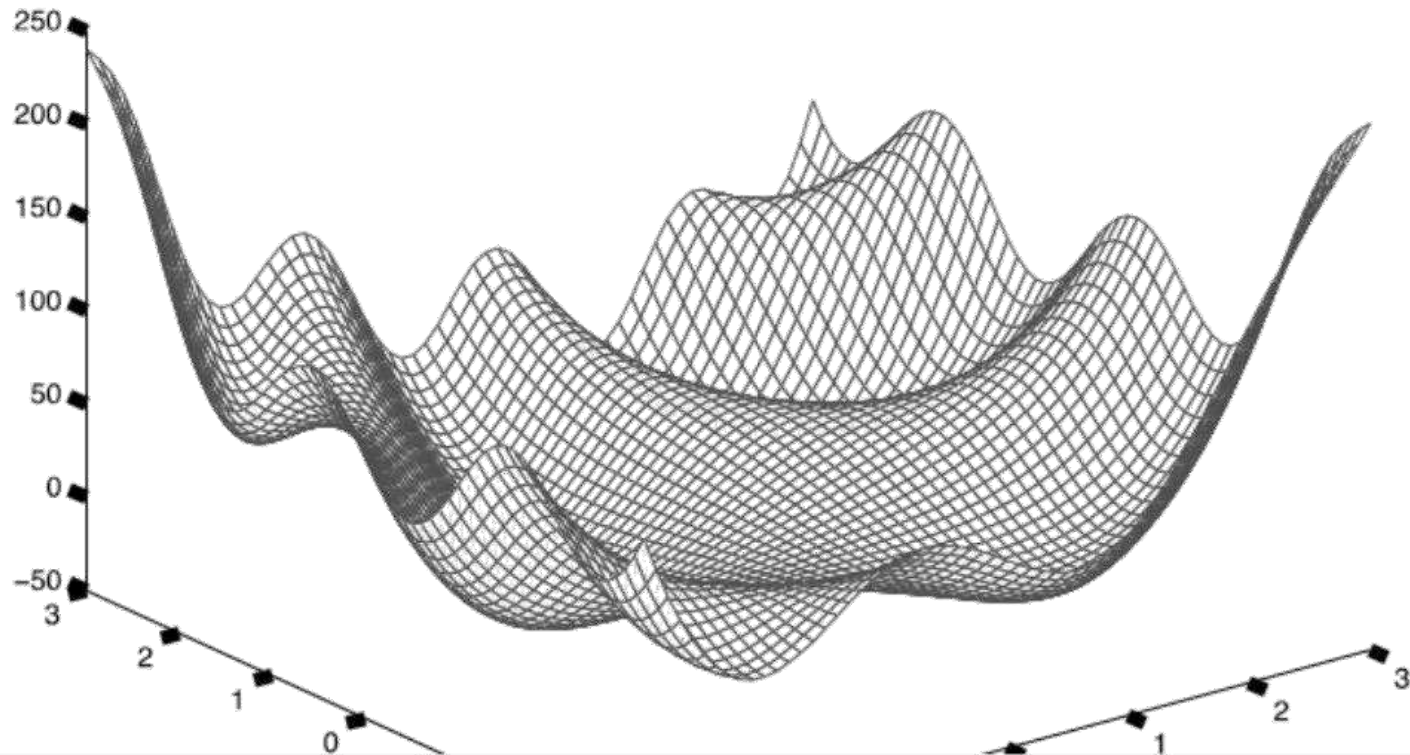
$$s.t. h_j(x) = 0, j = \{1, \dots, l\}$$

- Most of the machine learning problems are, in the end, optimization problems
-

# General Problem

---

■ Minimize  $f(x)$



# Linear Optimization

---

$$\boxed{Y = AX + w} \quad w \sim \mathcal{N}(0, R)$$

$$X^* = \min_X (Y - AX)^T R^{-1} (Y - AX)$$

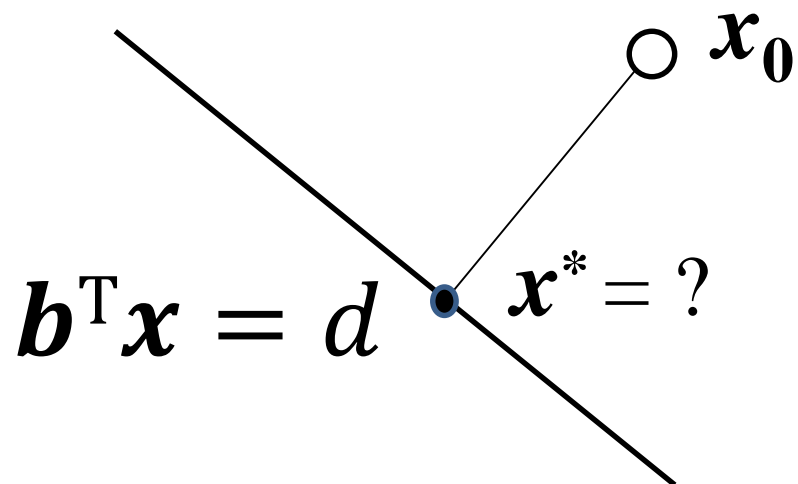
$$\frac{\partial}{\partial X^T} (Y - AX)^T R^{-1} (Y - AX) = 0$$

$$\Rightarrow X^* = (A^T R^{-1} A)^{-1} A^T R^{-1} Y$$

---

# Linear Optimization

---



$$x^* = x_0 - \frac{(b^T x_0 - d)b}{b^T b}$$

$$x^* = \min_x (x - x_0)^T (x - x_0)$$

$$\text{s. t. } b^T x - d = 0$$

---

# Nonlinear Optimization

---

## ■ Convex Optimization

- Unconstrained optimization
- Constrained optimization
- SVMs and Bayesian models

## ■ Non-convex Optimization

- Heuristic algorithms
  - Random search
-

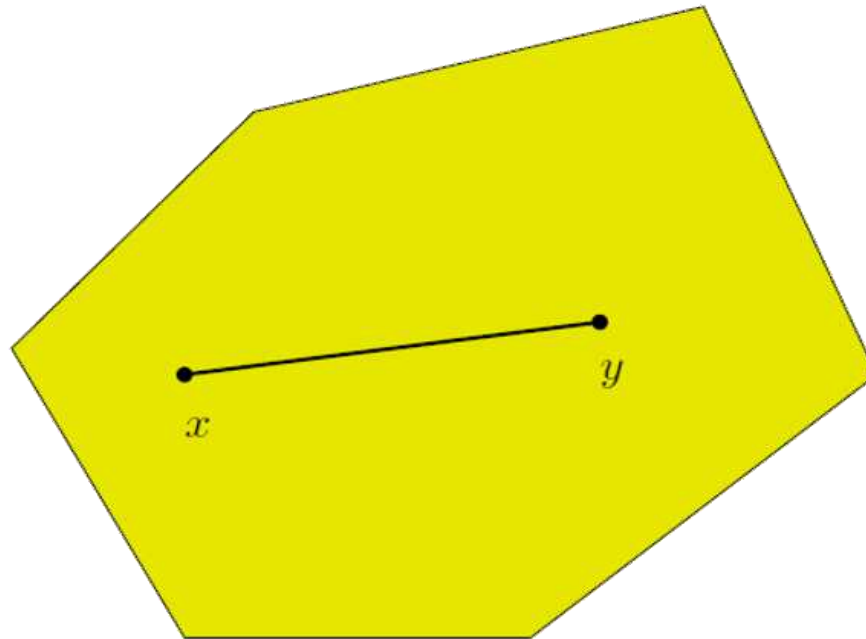
# What is Convex?

---

## ■ Convex sets

Def: A set  $C \subseteq \mathbb{R}^n$  is convex if for  $x, y \in C$ ;  $a \in [0, 1]$

$$ax + (1 - a)y \in C$$

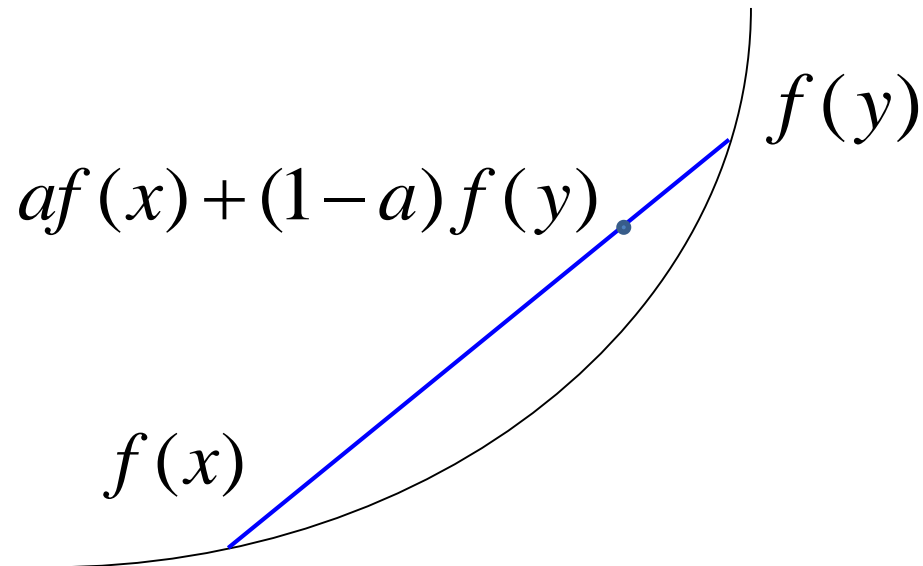


# What is Convex?

---

## ■ Convex functions

$$f(ax + (1-a)y) \leq af(x) + (1-a)f(y)$$

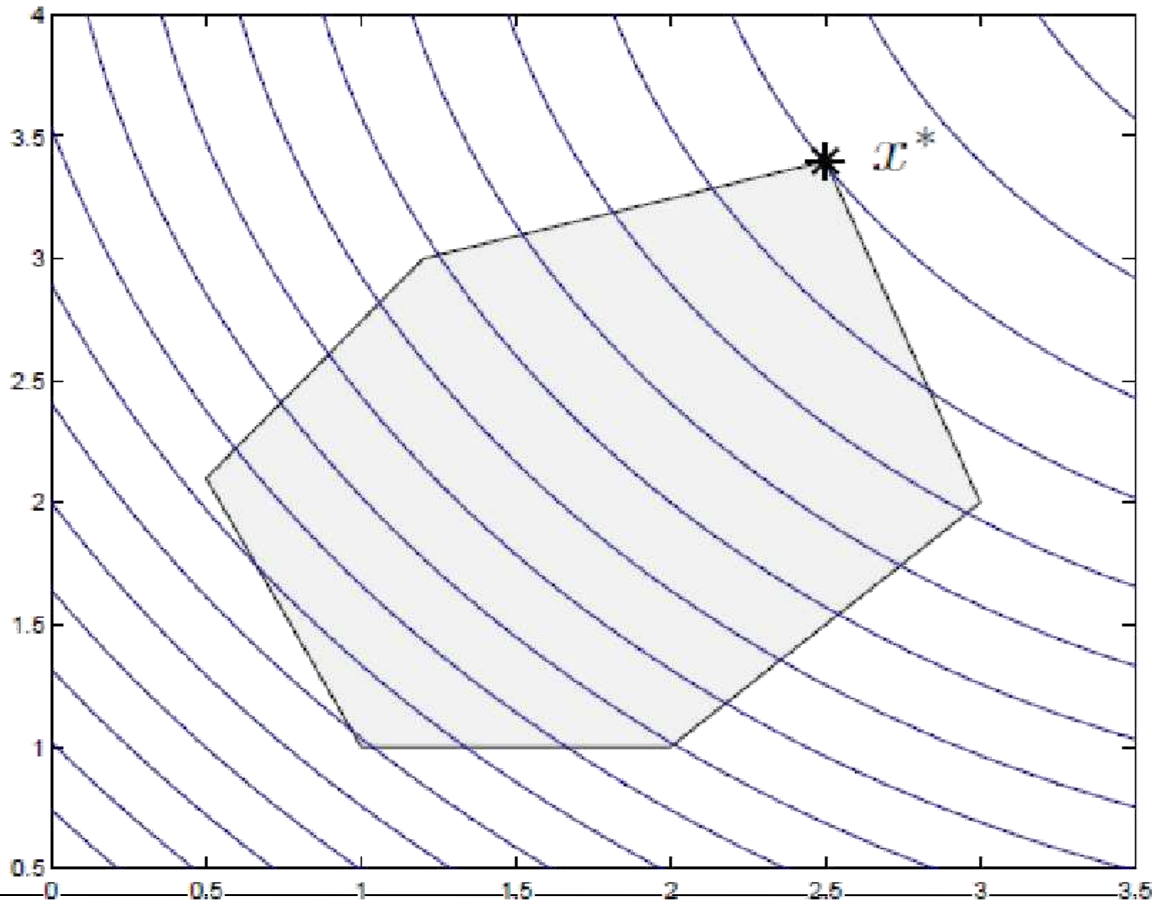




# Convex Optimization

---

- Local minimizer = Global minimizer



# Convex Optimization

---

## ■ Unconstrained optimization

- Gradient descent
- Gauss-Newton's method
- Batch learning
- Stochastic Gradient Descent

## ■ Constrained optimization

- Lagrange methods
  - Bayesian methods
-

# Convex optimization

---

## ■ Unconstrained optimization

- Gradient descent
- Gauss-Newton's method
- Batch learning
- Stochastic Gradient Descent

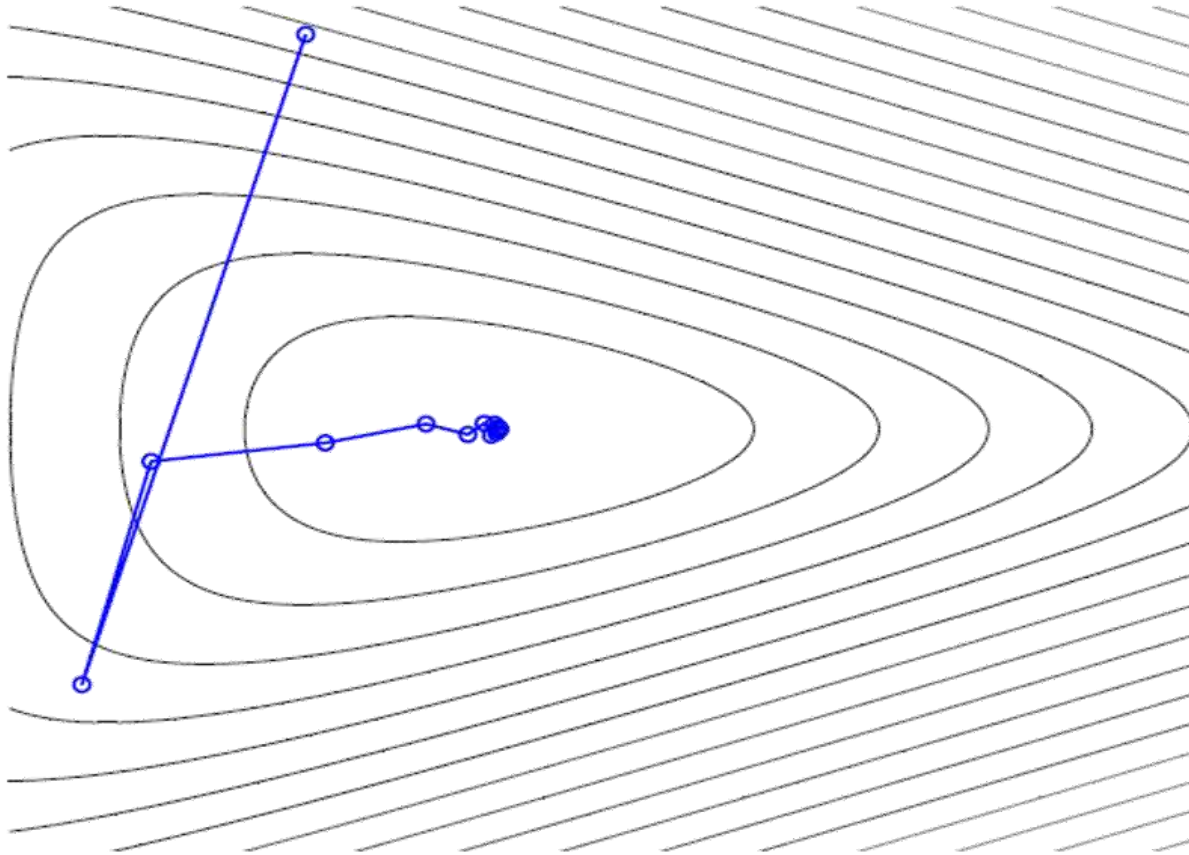
## ■ Constrained optimization

- Lagrange methods
  - Bayesian methods
-

# Gradient Descent

---

$$f(x_{t+1}) = f(x_t) - \eta \nabla f(x_t)^T (x - x_t)$$



# Gauss-Newton's Method

---

- Idea: use a second-order approximation to function

$$f(x + \Delta x) \approx f(x) + \nabla f(x)^T \Delta x + \frac{1}{2} \Delta x^T \nabla^2 f(x) \Delta x$$

- Choose  $\Delta x$  to minimize above:

$$\Delta x = -[\nabla^2 f(x)]^{-1} \nabla f(x)$$

- This is descent direction:

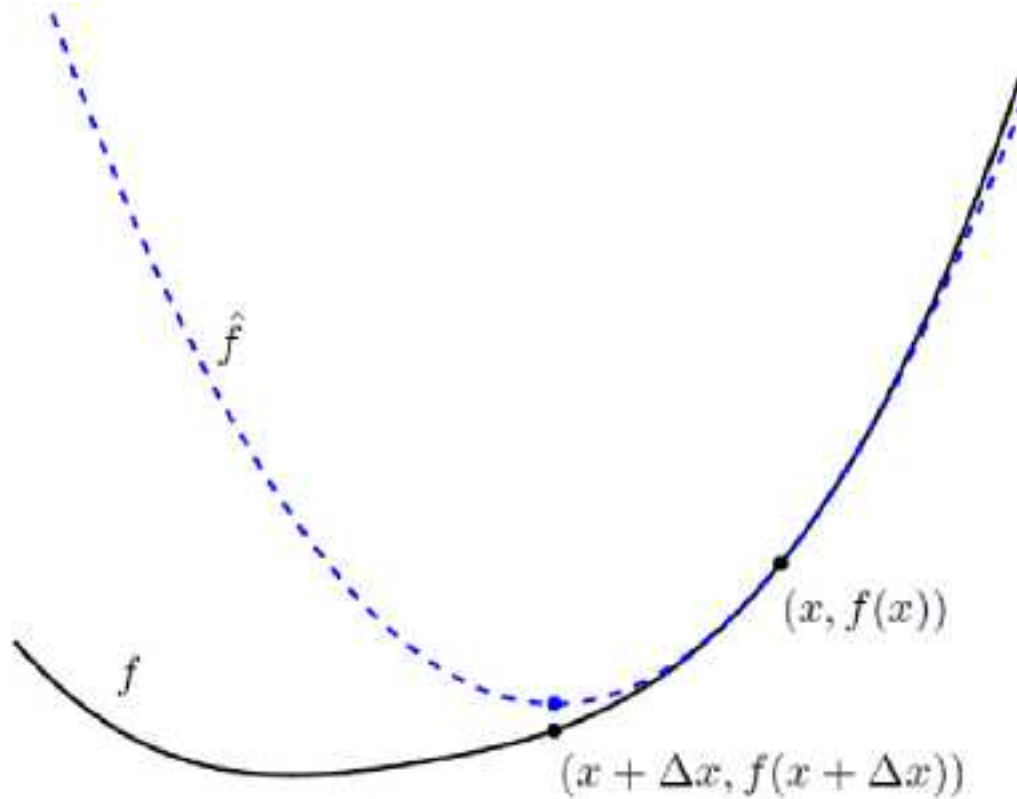
$$\nabla f(x)^T \Delta x = -\nabla f(x)^T [\nabla^2 f(x)]^{-1} \nabla f(x) < 0$$

---

# Gauss-Newton's Method

---

$\hat{f}$  is 2-order approximation,  $f$  is true function.



# Batch Gradient Descent

---

- Minimize empirical loss, assuming it's convex and unconstrained

- Gradient descent on the empirical loss

- At each step:

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \left( \frac{1}{n} \sum_{i=1}^n \frac{\partial L(w, x_i, y_i)}{\partial w} \right)$$

- Note: at each step, gradient is the average of the gradient for all samples ( $i=1, \dots, n$ )

- Very slow when  $n$  is very large

---

# Stochastic Gradient Descent

---

- Alternative: compute gradient from just one (or a few samples)
- Known as stochastic gradient descent:

- At each step,

$$w^{(k+1)} \leftarrow w^{(k)} - \eta_t \frac{\partial L(w, x_i, y_i)}{\partial w}$$

(choose one sample  $i$  and compute gradient for that sample only)

---



# Convex Optimization

---

## ■ Unconstrained optimization

- Gradient descent
- Gauss-Newton's method
- Batch learning
- Stochastic Gradient Descent

## ■ Constrained optimization

- Lagrange methods
  - Bayesian methods
-

# Lagrange Methods

---

- Start with an optimization problem:

$$\arg \min_x f_0(x)$$

$$s.t. f_i(x) \leq 0, i = \{1, \dots, k\}$$

$$s.t. h_j(x) = 0, j = \{1, \dots, l\}$$

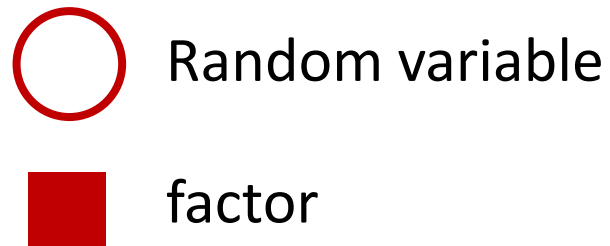
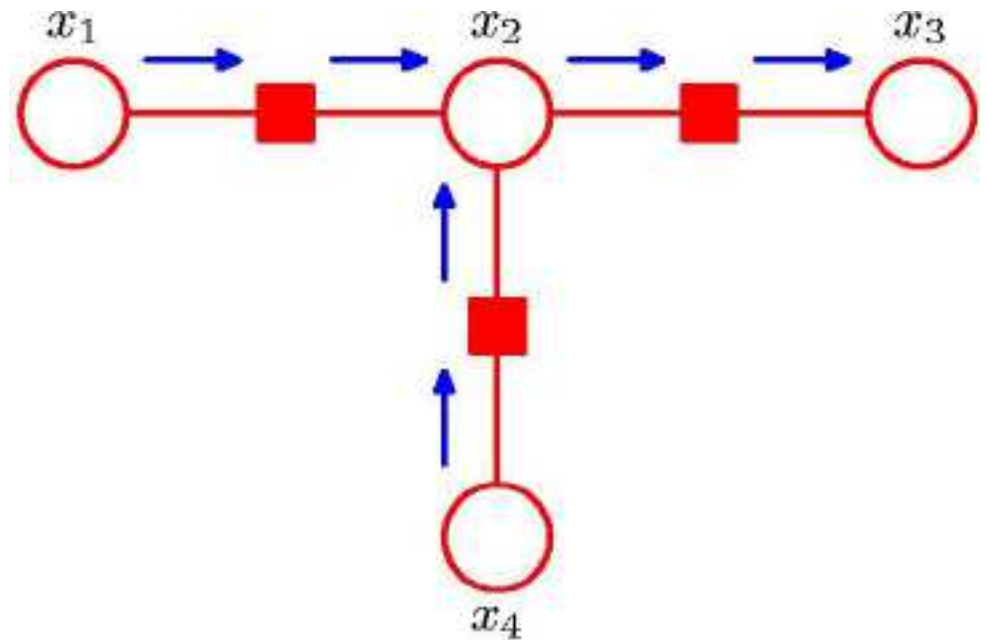
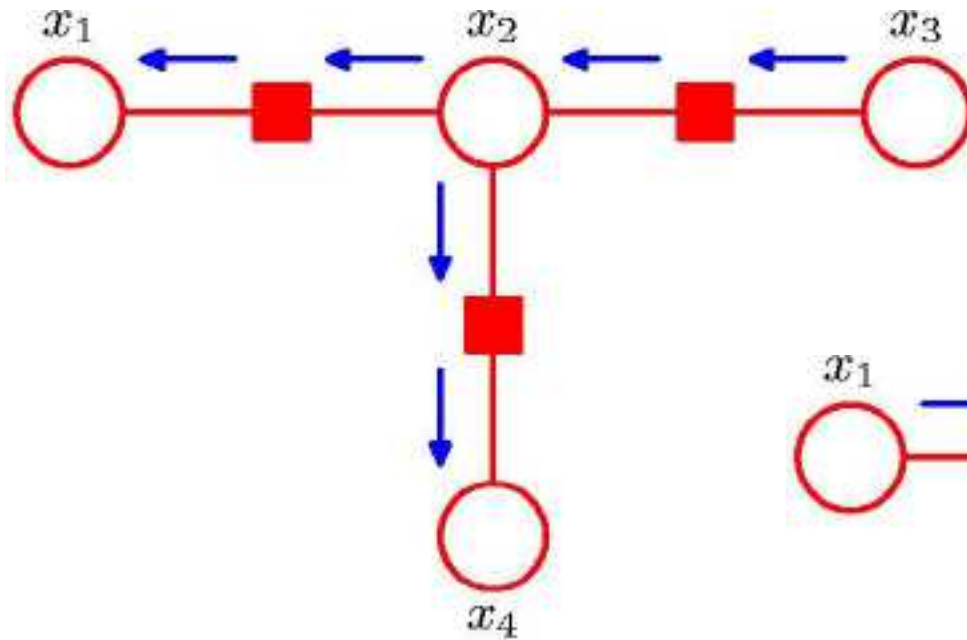
- Is equivalent to min-max optimization:

$$\arg \min_x \left[ \max_{\lambda \geq 0, \gamma > 0} \left( f_0(x) + \sum_{i=1}^k \lambda_i f_i(x) + \sum_{j=1}^l \gamma_j h_j(x) \right) \right].$$

---

# Bayesian Methods

---



# Convex Optimization for Machine Learning

---

## ■ Gradient Based Methods

- Neural networks

## ■ Lagrange Methods

- Support vector machines

## ■ Bayesian Methods:

- Expectation-Maximization methods (mixture models)
  - Variational methods (approximate models)
  - Graph optimization (belief propagation models)
-

# Non-convex Optimization

---

- Convex Optimization

- Unconstrained optimization

- Constrained optimization

- Non-convex Optimization

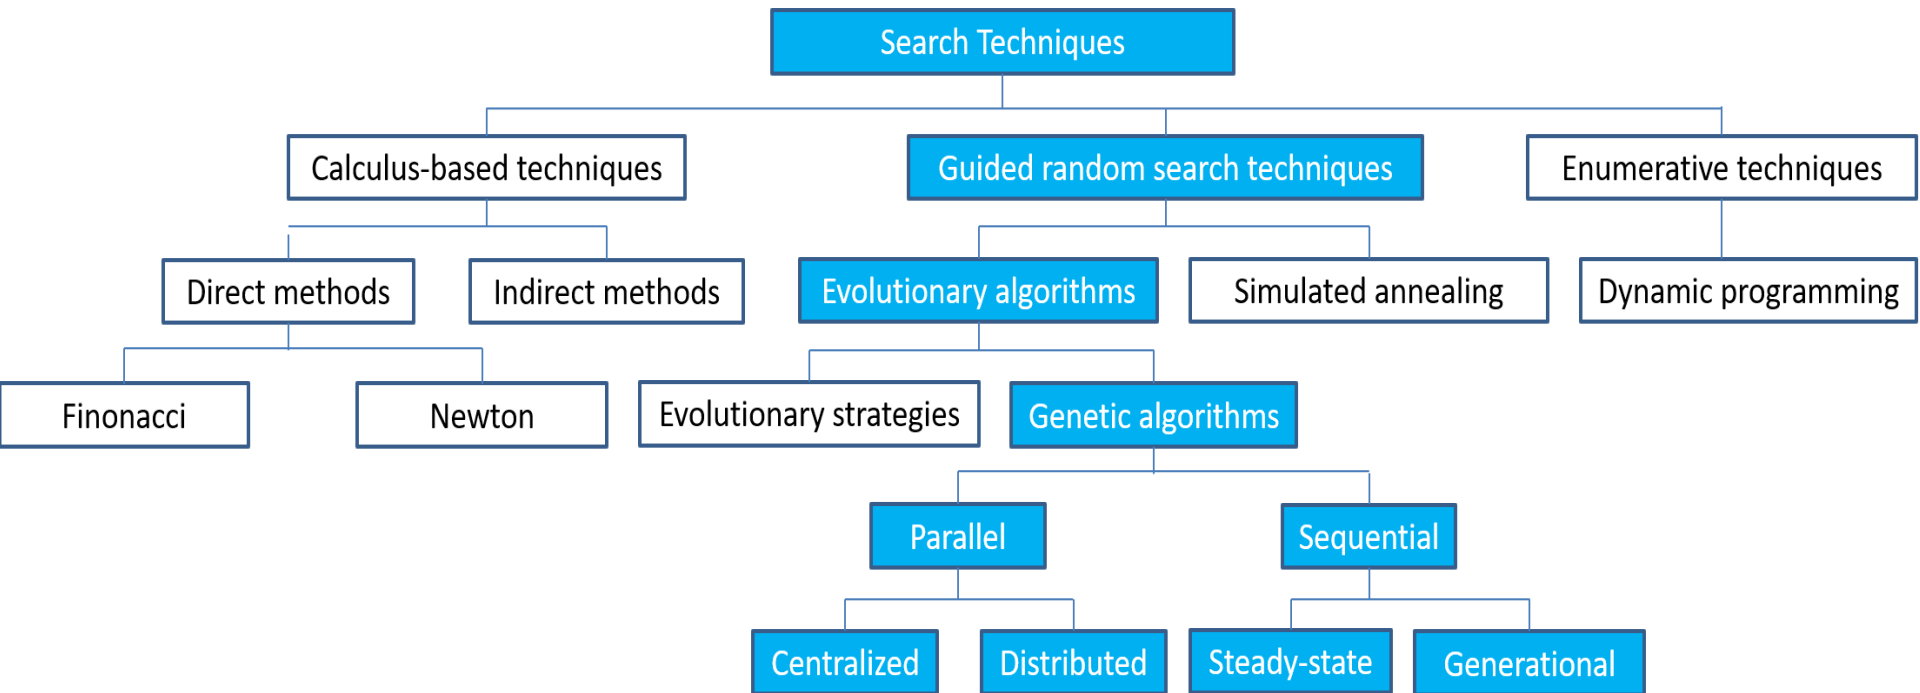
- Heuristic algorithms

- Random search

---

# Heuristic and Random Search

---



# Outlines

---

- Framework
  - Problem Statement
  - Related Areas
  - History
  - Datasets and Learning Models
  - Optimization Methods
  - Algorithms
  - Examples
-

# Algorithms

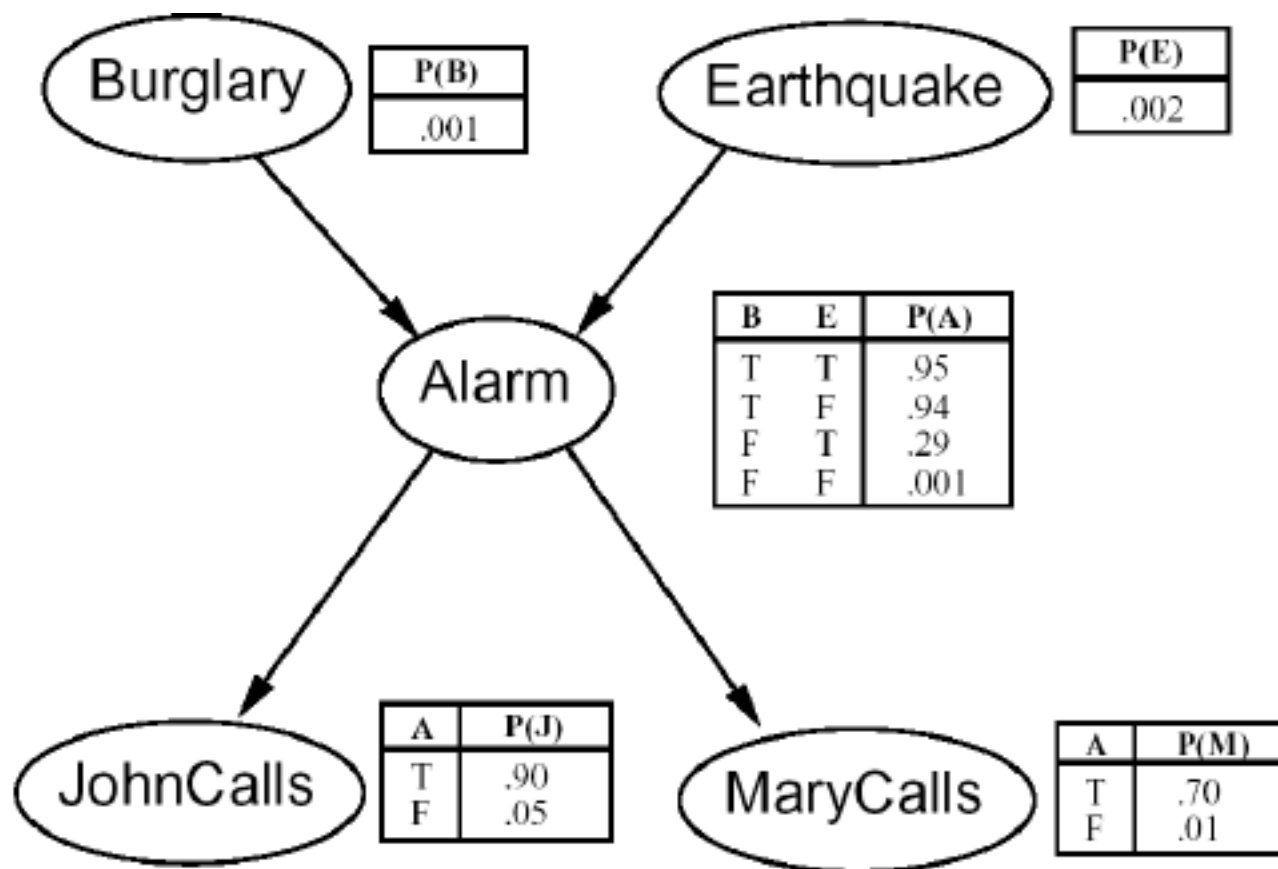
---

- Bayes
  - KNN and K-means
  - Decision tree
  - Support Vector Machine
  - Boosting and Ensemble Learning
  - Linear Statistical Learning (PCA, ICA, NMF)
  - Nonlinear Statistical Learning (Manifold learning)
  - Deep Neural Networks
  - Generative Adversarial Networks
  - Bayesian Networks
  - Reinforcement Learning
  - Federated Learning
-



# Bayes

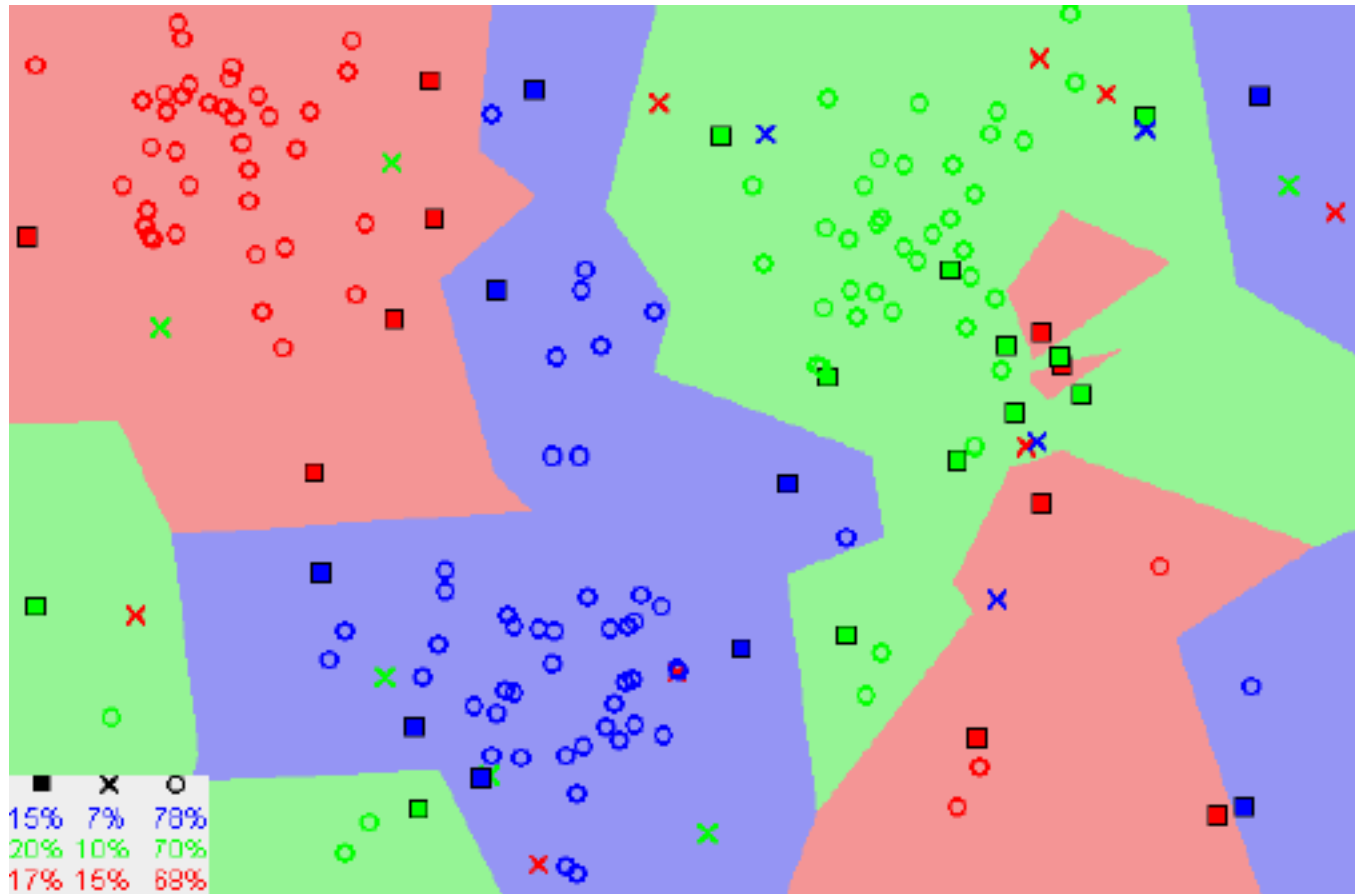
---



# K-Nearest Neighbors

---

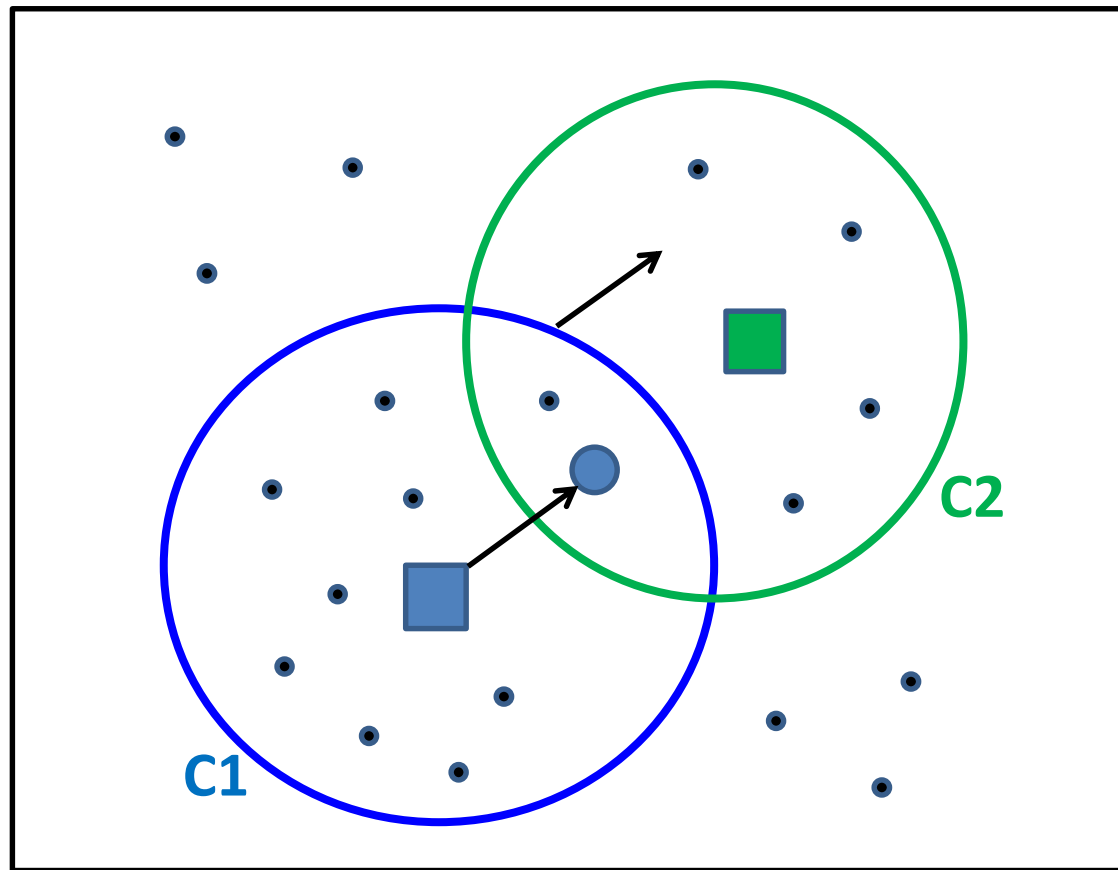
- Use training data for classification



# K-Means

---

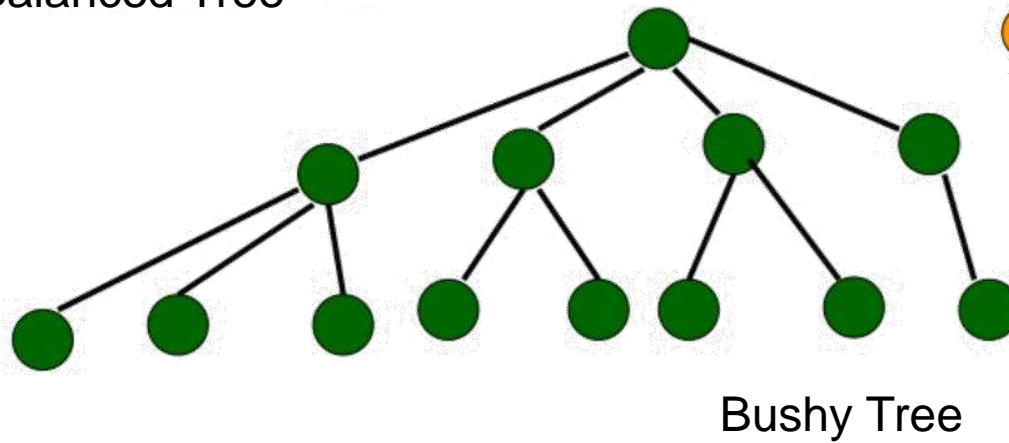
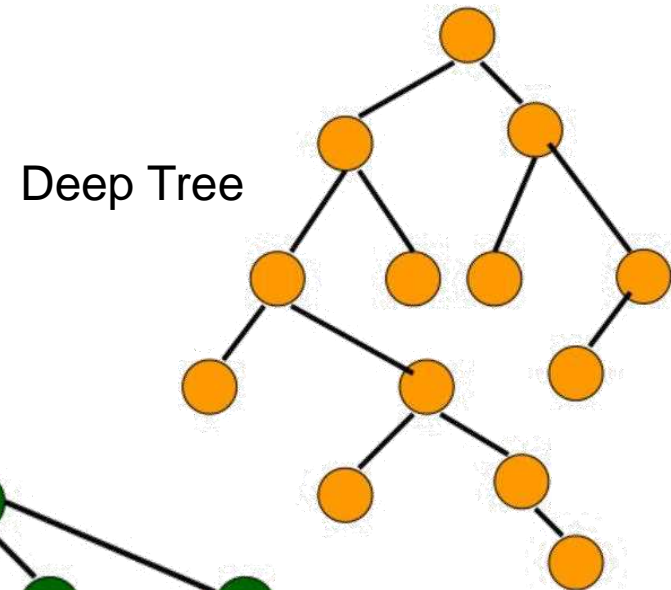
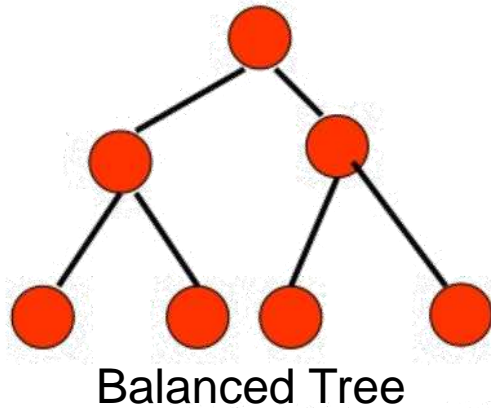
- Mean-shift for clustering



# Decision Tree

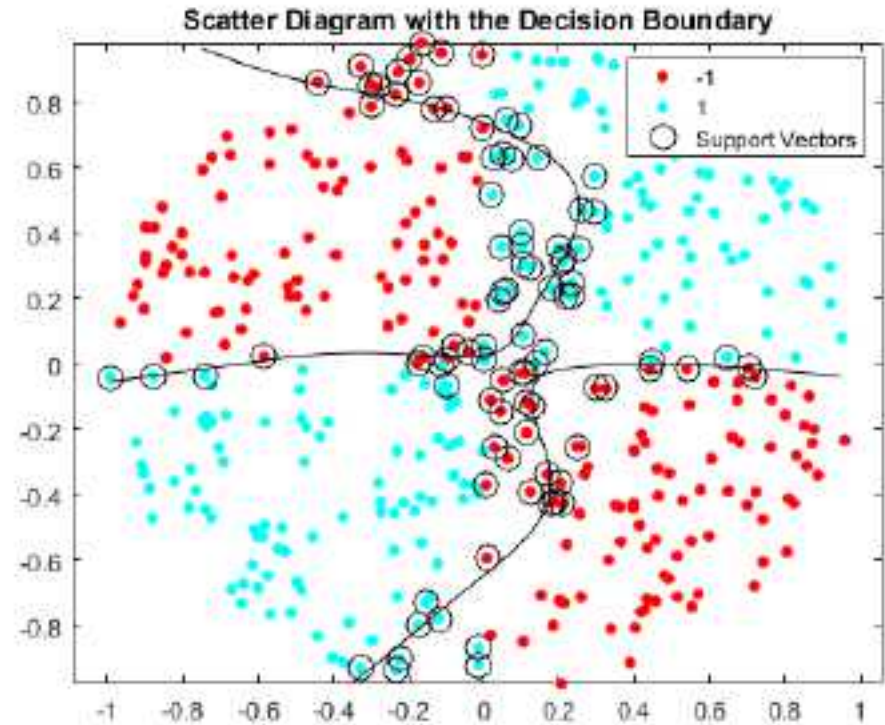
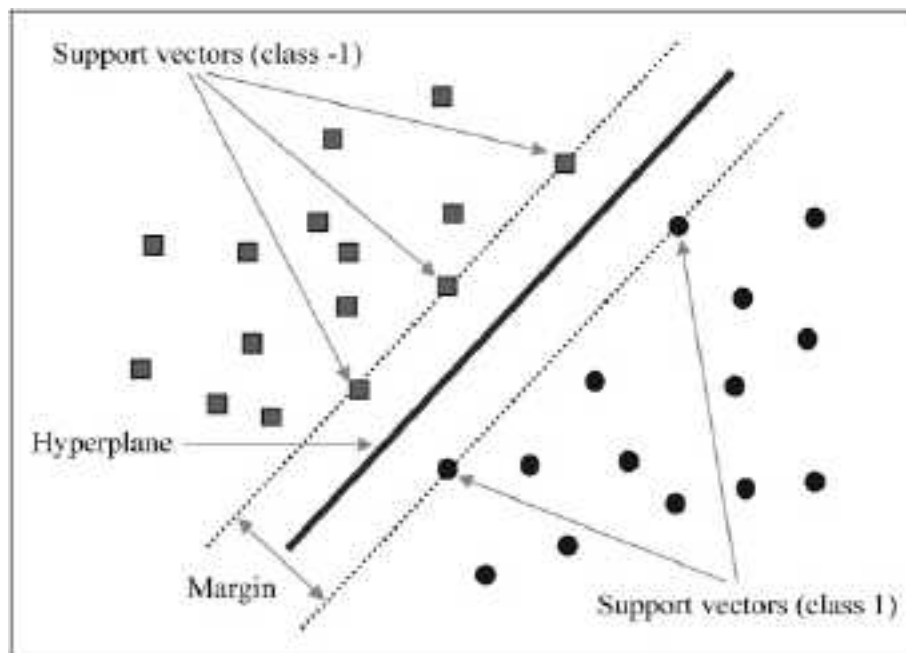
---

## ■ Types of Decision Tree:



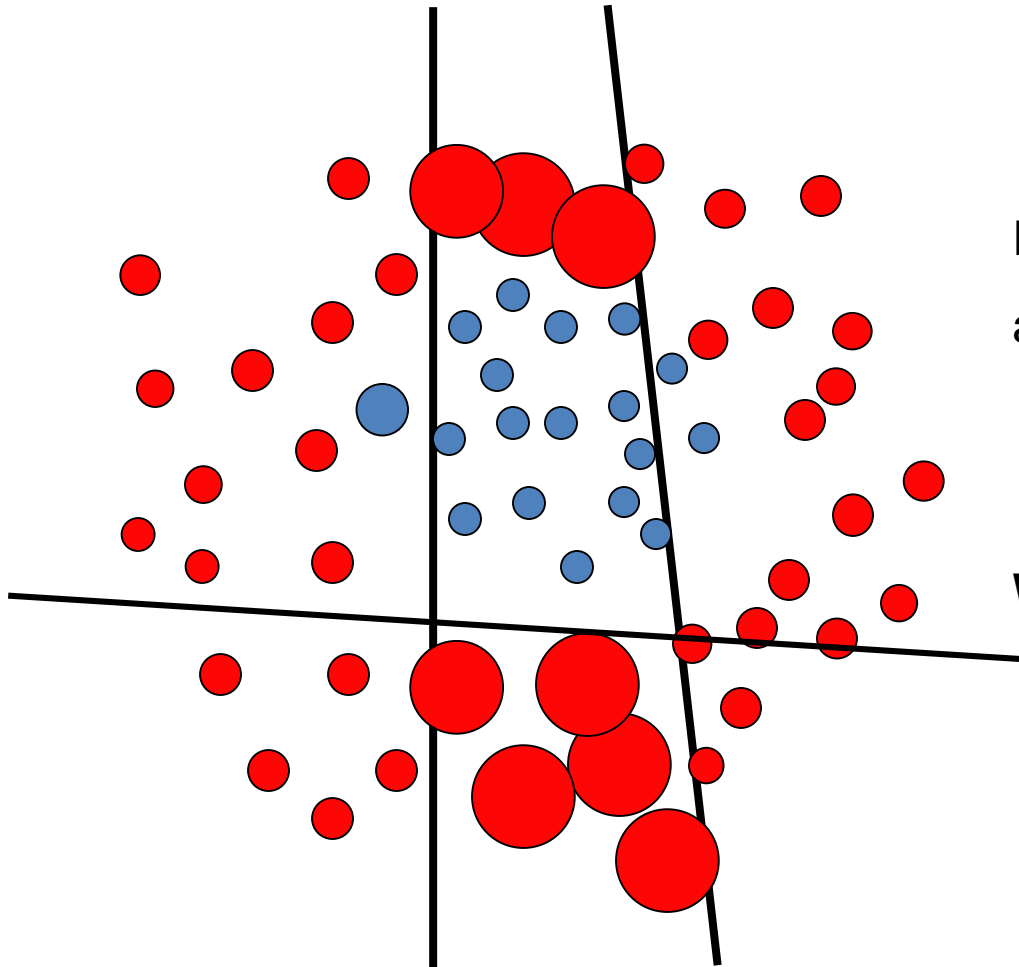
# SVM

- eg. Linear SVM: 
$$\arg \min_w \sum_{i=1}^n \|w\|^2 + C \sum_{i=1}^n \xi_i$$
$$\text{s.t. } 1 - y_i x_i^T w \leq \xi_i$$
$$\xi_i \geq 0$$



# Boosting

---



Each data point has  
a class label:

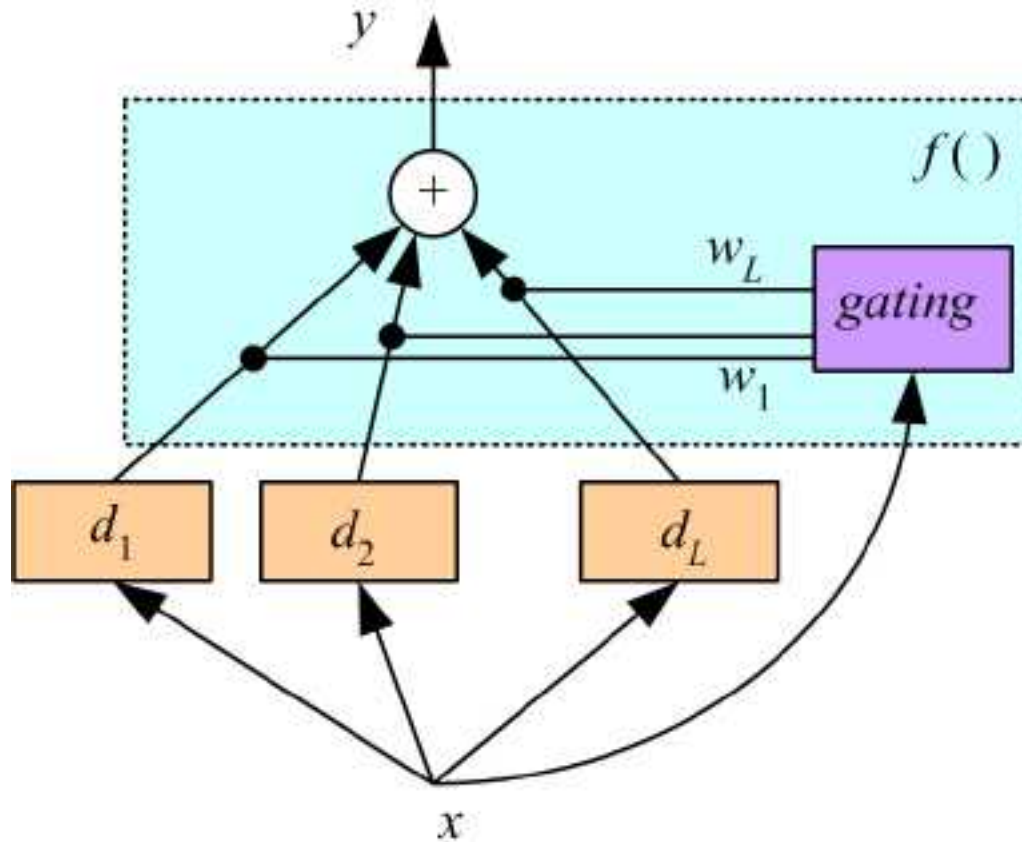
$$y_t = \begin{cases} +1 & (\text{red circle}) \\ -1 & (\text{blue circle}) \end{cases}$$

**We update the weights:**

$$w_t \leftarrow w_t \exp\{-y_t H_t\}$$

# Ensemble Learning

---

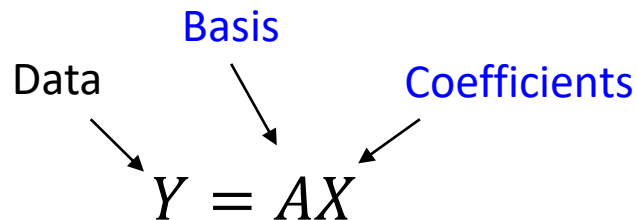


$$y = \sum_{j=1}^L w_j d_j$$

# Linear Statistical Learning

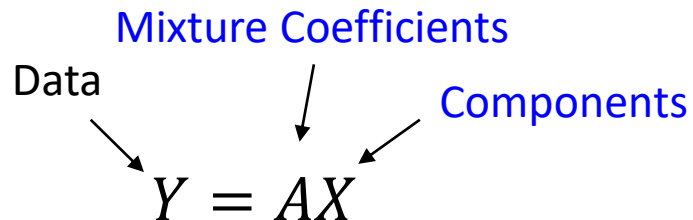
---

## ■ PCA



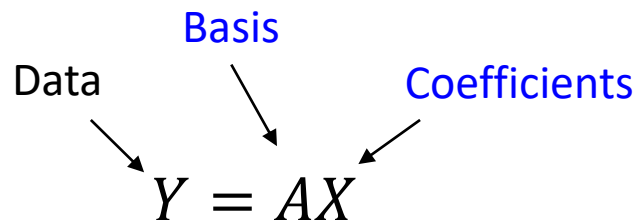
$$A_i \perp A_j$$

## ■ ICA



$$\min I(X)$$

## ■ NMF



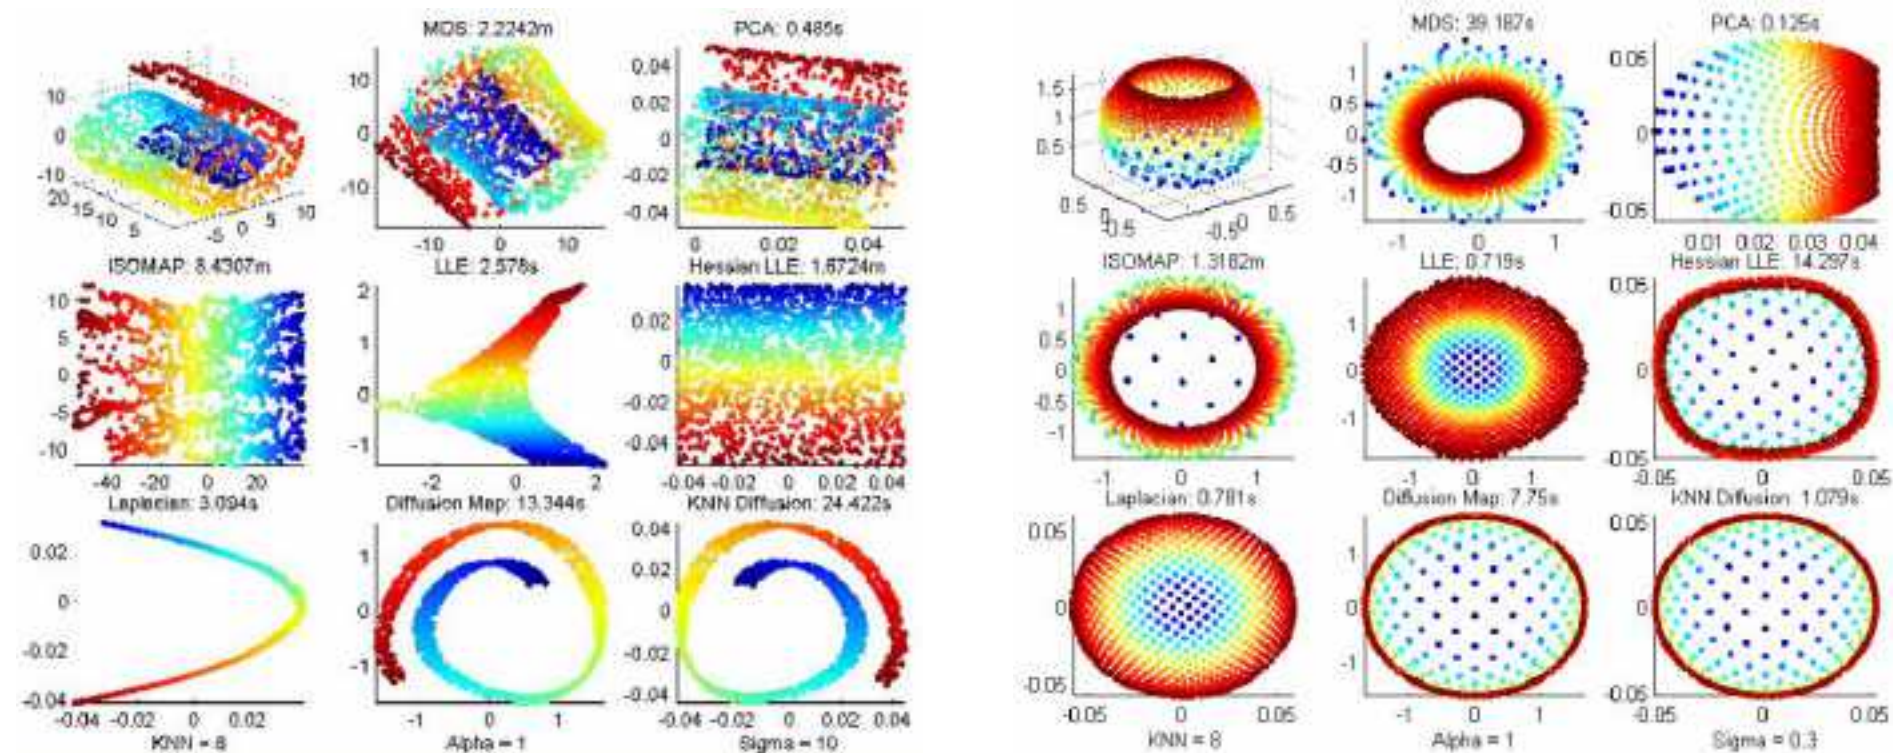
$$A, X > 0$$

---



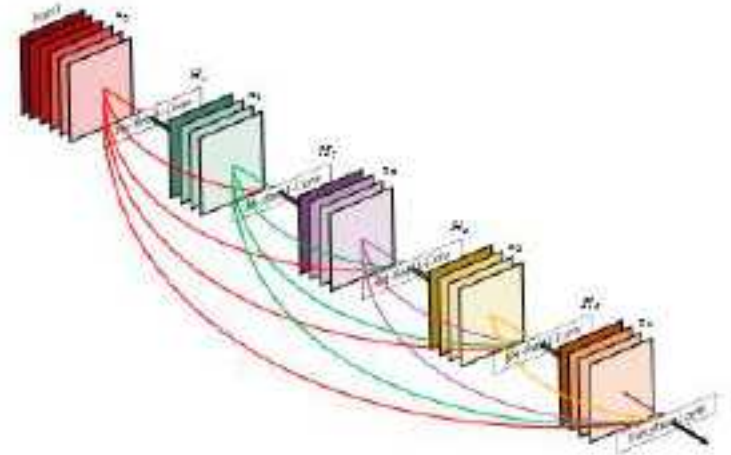
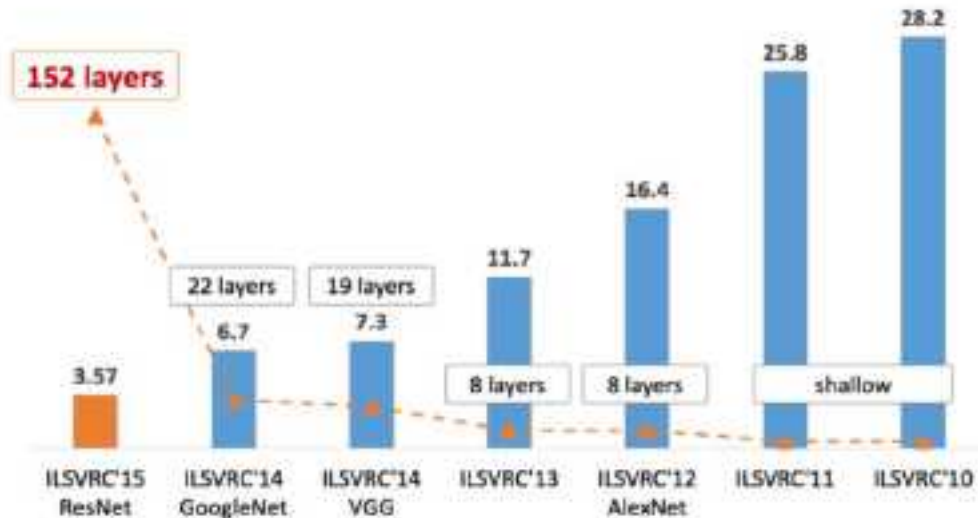
# Nonlinear Statistical Learning

## ■ Manifold learning



# Deep Neural Networks

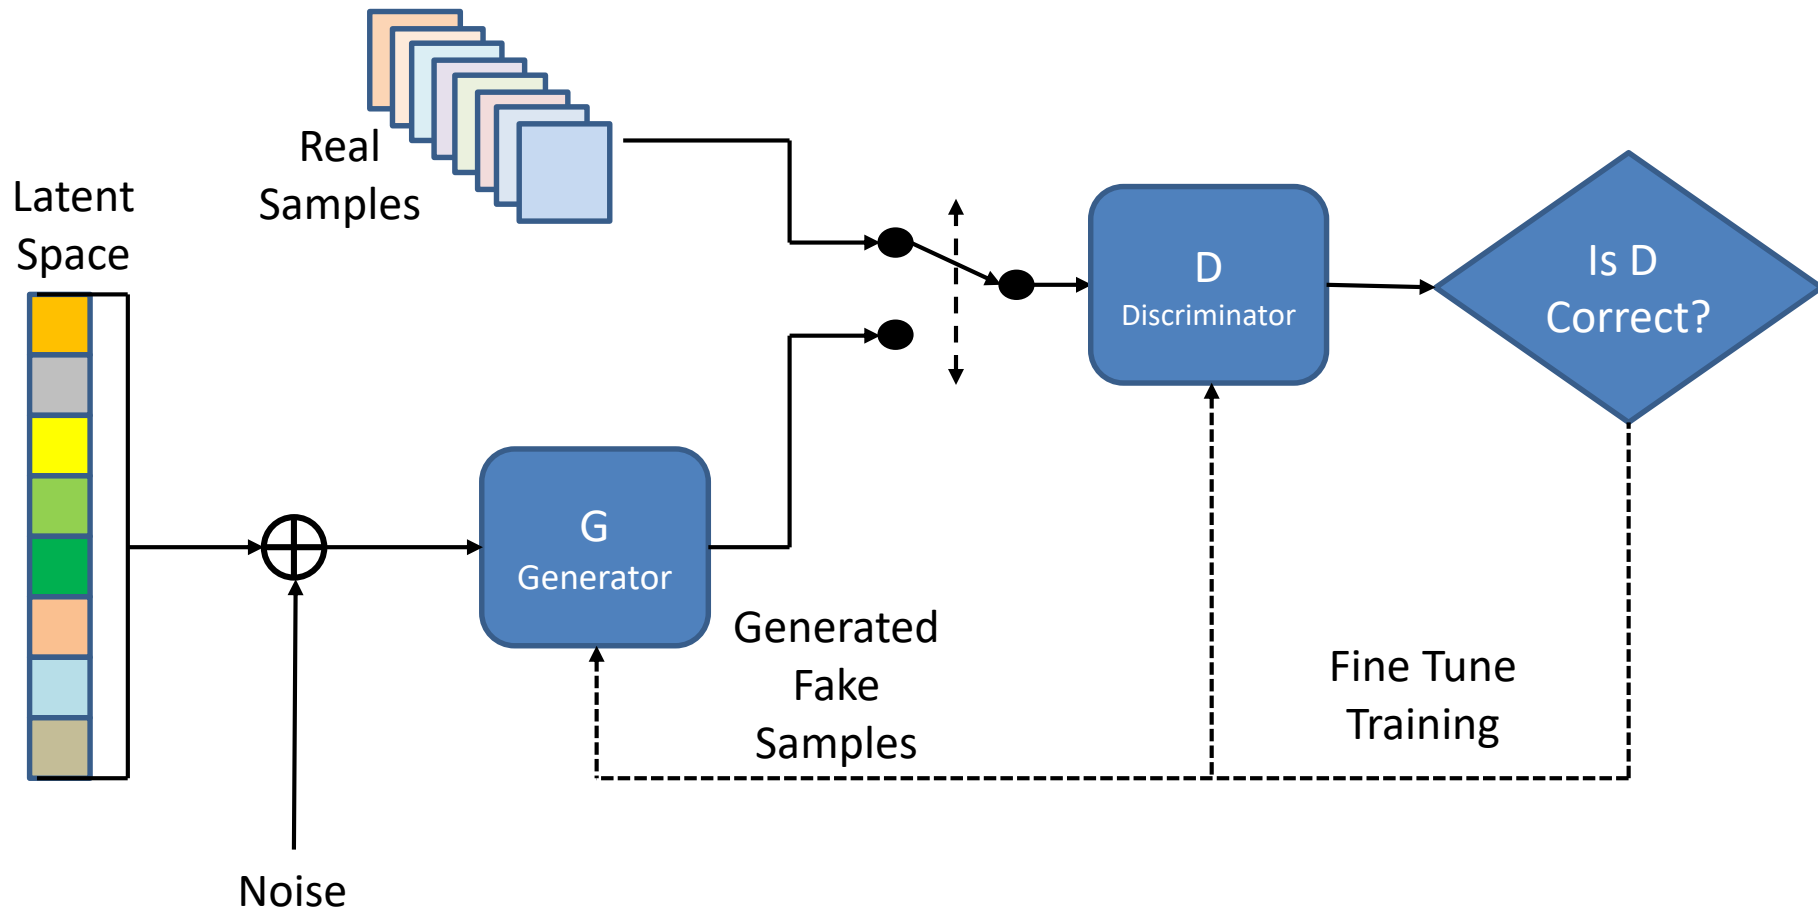
---



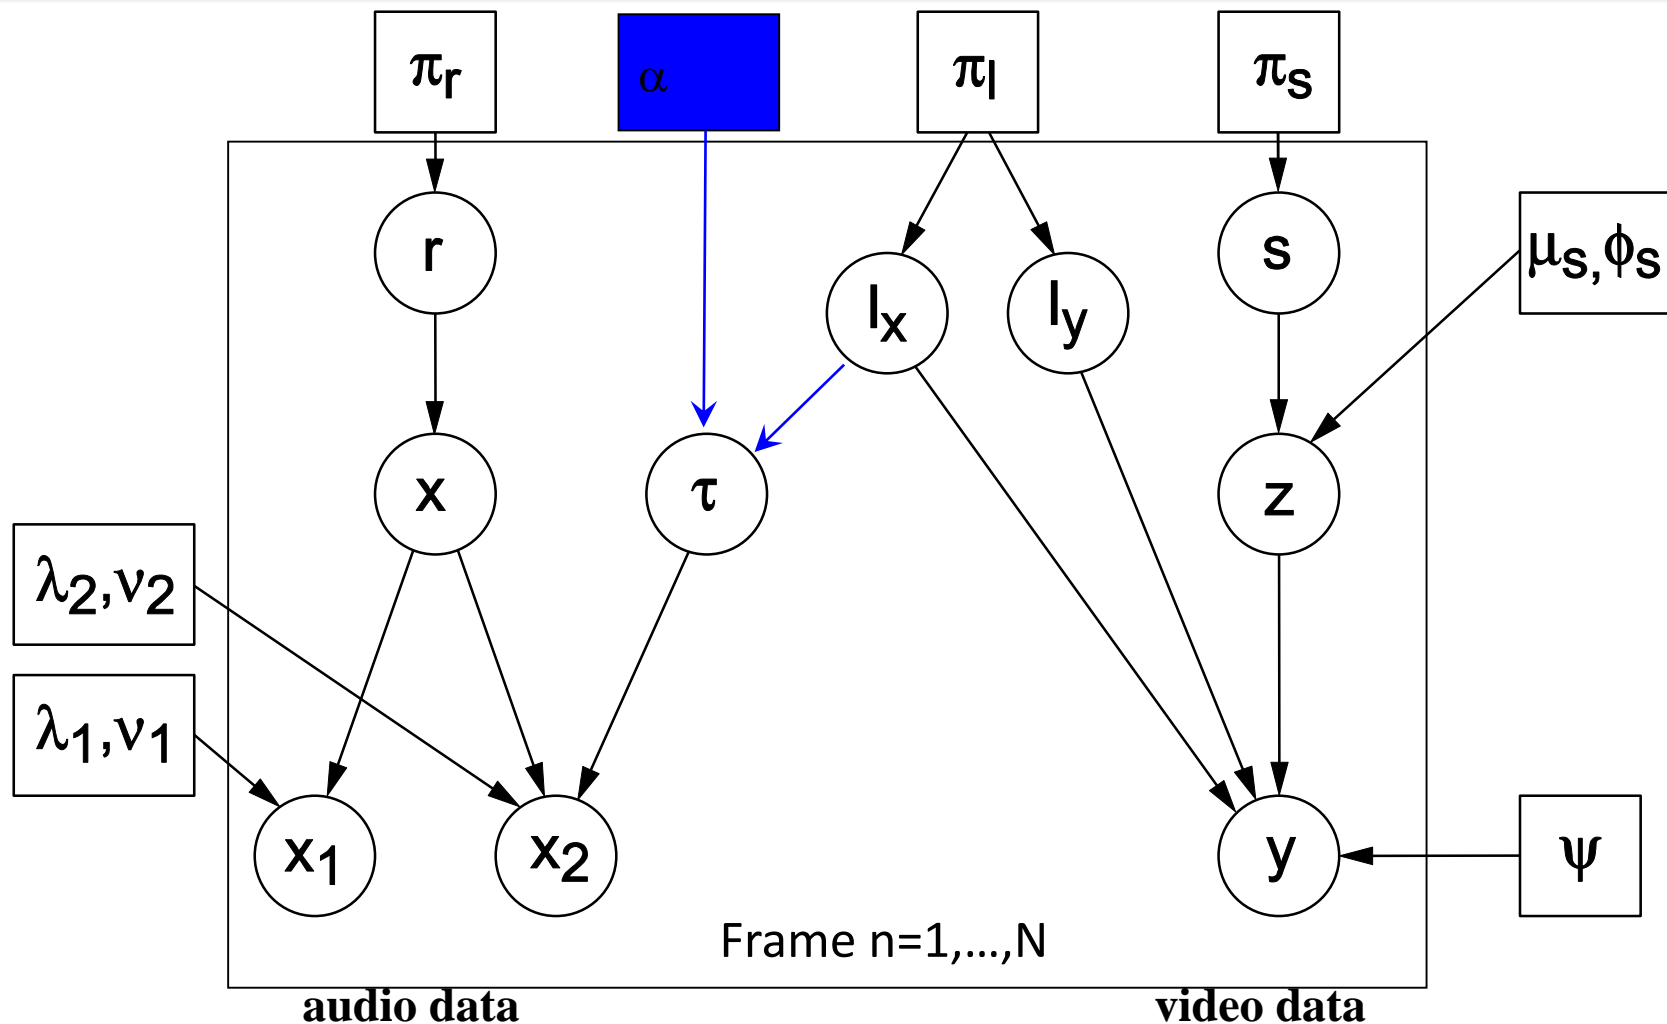
- Task: recognition
- Dataset: ILSVRC
- Huang G, Liu Z, Weinberger K Q, et al. Densely connected convolutional networks[J]. arXiv preprint arXiv:1608.06993, 2016.

# Generative Adversarial Networks

---



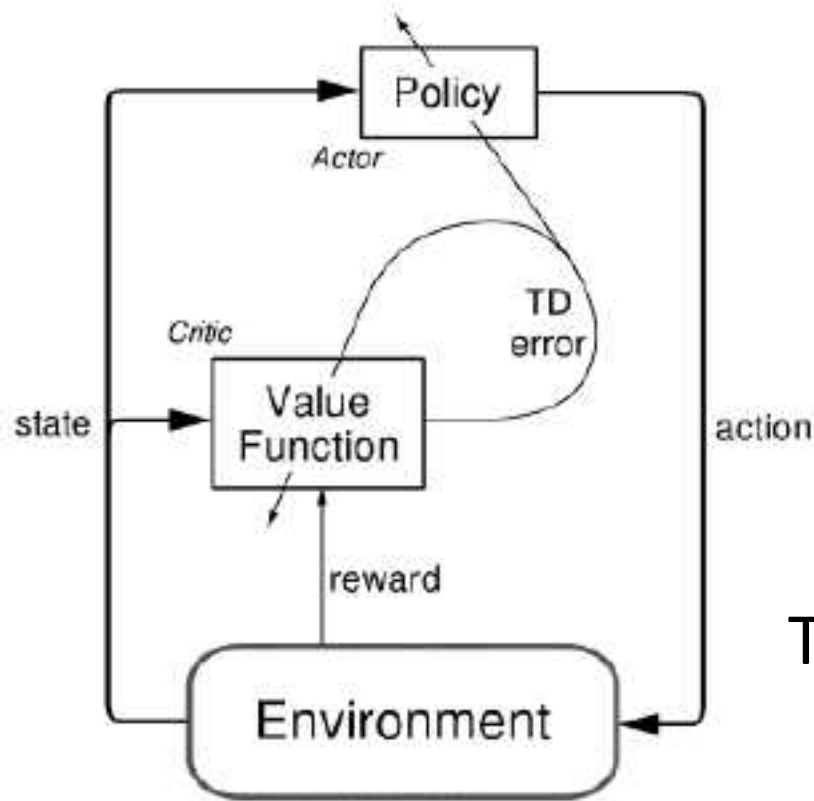
# Bayesian Networks



# Reinforcement Learning

---

- State, action, and Reward



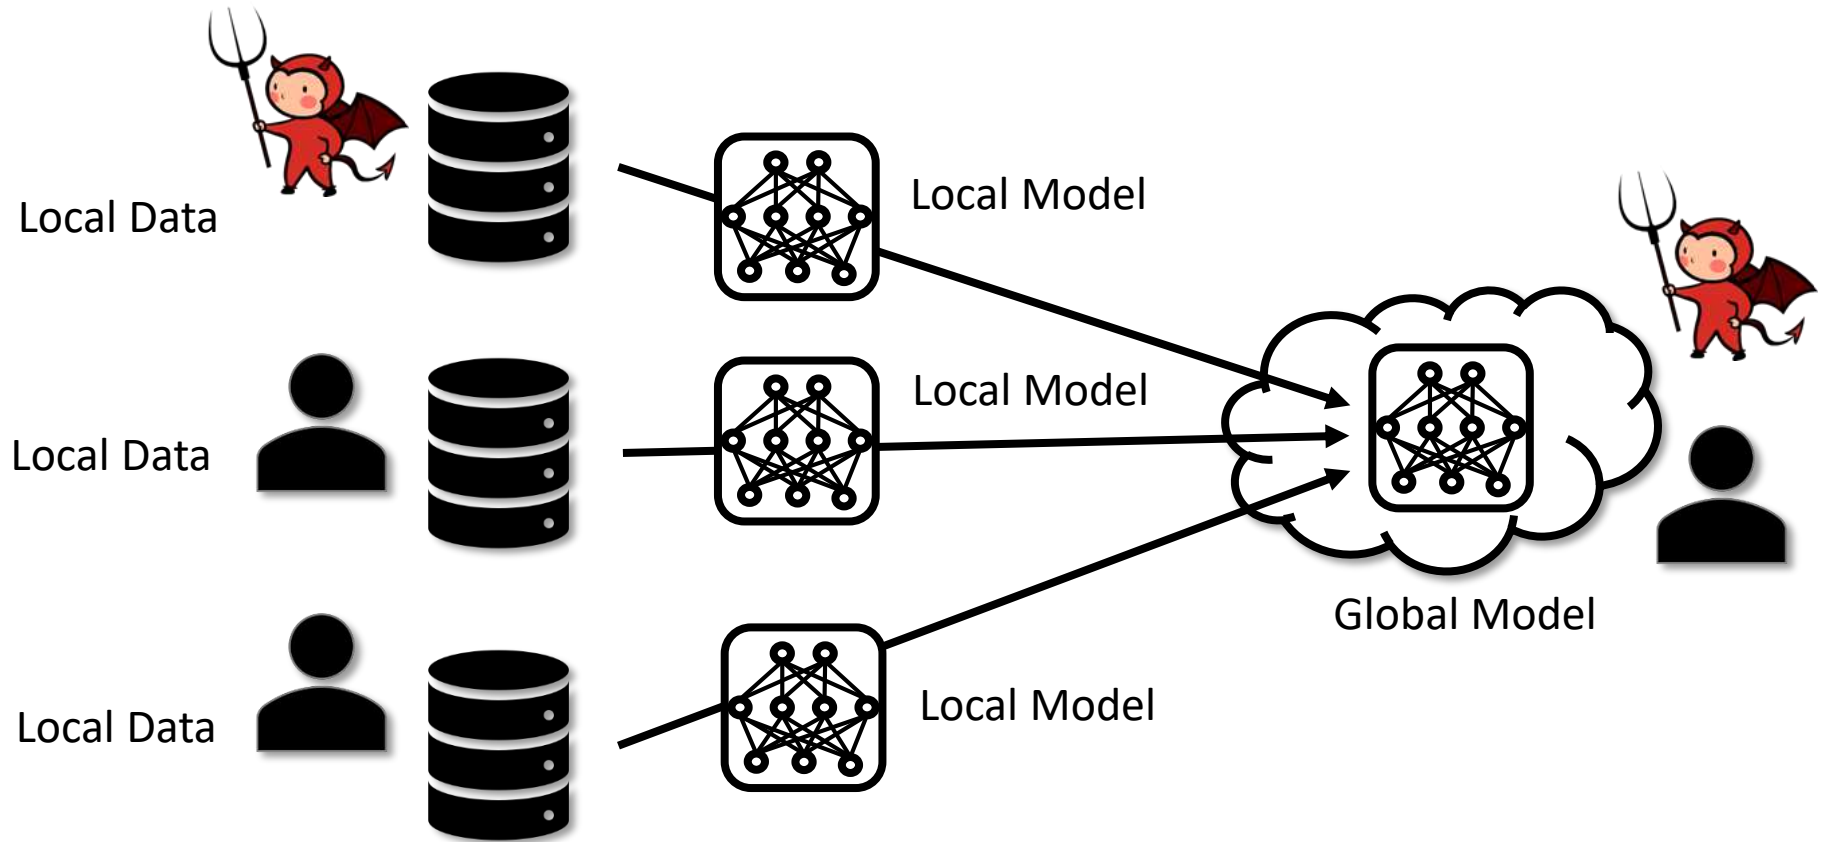
Update: Policy Function  
Value Function

TD Error: Temporal Difference  
between Real Reward  
and Estimated Reward

# Federated Learning

---

- Collaborative Learning



# More Course Links

---

## **Stanford Machine Learning:**

<https://see.stanford.edu/Course/CS229/47>

**MIT Machine Learning:** <https://ocw.mit.edu/courses/electrical-engineering-and-computer-science/6-867-machine-learning-fall-2006/index.htm>

**Stanford CNN for Vision:** <http://cs231n.stanford.edu>

**Stanford Deep Learning:** <http://cs230.stanford.edu/syllabus.html>

**MIT Deep Learning:** <http://introtodeeplearning.com/>

---