

CS302 Assignment 2

何泽安 12011323

2023.3.9

1. The usage of each parameters in the following command:

```
qemu-system-riscv64 \  
-machine virt \  
-nographic \  
-bios default \  
-device loader,file=bin/ucore.bin,addr=0x80200000
```

- `qemu-system-riscv64` is the executable file under `path/to/qemu/installation/riscv64-softmmu`, as we have added this path into the environment variable `PATH`, we can directly run the riscv64 VM by calling its name.
- `-machine virt` is the command that can select the emulated machine, and the param `virt` is to use the RESC-V VirtIO board.

```
heza12011323@VM-8-14-ubuntu:~/qemu-5.0.0/riscv64-softmmu$ qemu-system-riscv64 -h | grep machine  
-machine [type=]name[,prop[=value][,...]]  
    selects emulated machine ('-machine help' for list)  
    specify machine UUID  
    configure or create an on-board (or machine default) NIC and  
--preconfig pause QEMU before machine is initialized (experimental)  
            xenpv machine type).  
heza12011323@VM-8-14-ubuntu:~/qemu-5.0.0/riscv64-softmmu$ qemu-system-riscv64 -machine help  
Supported machines are:  
none          empty machine  
sifive_e      RISC-V Board compatible with SiFive E SDK  
sifive_u      RISC-V Board compatible with SiFive U SDK  
spike         RISC-V Spike Board (default)  
spike_v1.10   RISC-V Spike Board (Privileged ISA v1.10)  
spike_v1.9.1  RISC-V Spike Board (Privileged ISA v1.9.1)  
virt          RISC-V VirtIO board
```

- `-nographic` specify the VM to disable graphical output and redirect serial I/Os to the console.
- `-bios default` sets the filename (default) for the BIOS
- `-device loader,file=bin/ucore.bin,addr=0x80200000` is adding a device (using the *Generic Loader*), specifying the file `bin/ucore.bin` as the minimal OS kernel binary file (ELF), and specifying that the ELF should be loaded to the base address `0x80200000`, which will be executed after start up.

2. Explain the meaning of codes in `kernel.ld`: (as marked in the comment)

```
OUTPUT_ARCH(riscv) /* specify the output/target's CPU architecture is RISC-V */  
ENTRY(kern_entry) /* specify the entry point of the target is kern_entry */  
  
BASE_ADDRESS = 0xFFFFFFFFC0200000; /* define the base address of location counter  
*/  
  
SECTIONS /* start describing the memory layout of the output */
```

```

{
/* Load the kernel at this address: "." means the current address
 * this actually specify the starting address of the
 * `.text` section of the output file
 */
. = BASE_ADDRESS;

/* define the sequence of items in the `.text` section of output */
.text : {
    /* the `.text` of output will be filled with the below elements in sequence
     * first find the kern_entry, then is the `.text`, `.stub` section,
     * then fill all block entries under the `.text` section, finally
     * is the gun linkonce for inline assembler text section
     */
    *(.text.kern_entry .text .stub .text.* .gnu.linkonce.t.*)
}

/* using the PROVIDE keyword can define a symbol, when we need to reference it
 * but is not defined. this can prevent the confliction with function name in C
 * this will act as the `extern` variables mentioned below
 */
PROVIDE(etext = .); /* Define the 'etext' symbol to `.` (location counter
here)*/

/* define the sequence of items in the `.rodata` (read only) section of output
*/
.rodata : {
    /* the `.rodata` of output will be filled with the below elements in
sequence
     * first is all files' `.rodata` section, then are all the blocks under
     * the `.rodata` section, then is the gun linkonce for inline assembler
     * rodata section
     */
    *(.rodata .rodata.* .gnu.linkonce.r.*)
}

/* Adjust the address for the data segment to the next page,
 * as the page size of RISCV is 4kb = 0x1000byte
 */
. = ALIGN(0x1000);

/* define the sequence of items in the `.data` segment of output */
.data : {
    *(.data) /* first fills with all files' `.data` section */
    *(.data.*) /* then all blocks entry under `.data` section */
}

/* define the sequence of items in the `.sdata` (static data) segment of output
*/

```

```

.sdata : {
    *(.sdata) /* first fills with all files' `.sdata` section */
    *(.sdata.*) /* then all blocks entry under `.sdata` section */
}

PROVIDE(edata = .); /* Define the 'edata' symbol to `.` (location counter)*/

/* define the sequence of items in the `.bss` (block starting symbol) segment
of output */
.bss : {
    *(.bss) /* first is all files' `.bss` section */
    *(.bss.*) /* then all blocks entry under `.bss` section */
    *(.sbss*) /* finally, all files' `.sbss` (static bss) section */
}

PROVIDE(end = .); /* Define the 'end' symbol to `.` (location counter)*/

/* specify which sections to be discarded, aka, these sections in the input
files
    * will not be included in the output file
    */
/DISCARD/ : {
    *(.eh_frame .note.GNU-stack) /* discard `.eh_frame` and `.note.GNU-stack`
sections */
}
}

```

3. Explain the memset operation in kern_init:

```

extern char edata[], end[]; // two pointers from kernel.ld
memset(edata, 0, end - edata);

```

As we defined the `edata` and `end` symbol in `kernel.ld` (explained above), that between `edata` and `end` is the `bss` section. From the function prototype that `memset(void* dest, int ch, size_t count)`, this instruction is to fill all the memory bytes between the address `edata` and `end` point to, aka. the `bss` segment (that should be initially set to zero, which satisfies the requirements of program and OS).

4. Explain how `cputs` prints a string via sci:

```

// kern/libs/stdio.c
int cputs(const char *str) {
    int cnt = 0;
    char c;
    while ((c = *str++) != '\0') {
        cputch(c, &cnt); //<----
    }
    cputch('\n', &cnt);
}

```

```

        return cnt;
    }

    static void cputch(int c, int *cnt) {
        cons_putc(c);    //<----
        (*cnt) ++;
    }

    // kern/driver/console.c
    void cons_putc(int c) {
        bool intr_flag;
        local_intr_save(intr_flag);
        {
            sbi_console_putchar((unsigned char)c);    //<----
        }
        local_intr_restore(intr_flag);
    }

    // libs/sbi.h

#define SBI_CALL(which, arg0, arg1, arg2) ({           \
    register uintptr_t a0 asm ("a0") = (uintptr_t)(arg0); \
    register uintptr_t a1 asm ("a1") = (uintptr_t)(arg1); \
    register uintptr_t a2 asm ("a2") = (uintptr_t)(arg2); \
    register uintptr_t a7 asm ("a7") = (uintptr_t)(which); \
    asm volatile ("ecall"                               \
        : "+r" (a0)                                     \
        : "r" (a1), "r" (a2), "r" (a7)                 \
        : "memory");                                   \
    a0;                                                  \
    })

#define SBI_CALL_1(which, arg0) SBI_CALL(which, arg0, 0, 0)
#define SBI_CONSOLE_PUTCHAR 1

static inline void sbi_console_putchar(int ch)
{
    SBI_CALL_1(SBI_CONSOLE_PUTCHAR, ch);    //<----
}

```

The full call stack is listed above, where a C-string is first passed into the `cputs` function, which processes the string char-by-char (upcast to int). `cputch` then handle each char (in the C-string, and an additional `\n`): since our target is to display the char in the console, it passes the char (upcasted to int) into the `cons_putc` to do the display job, and it also manages a counter that helps `cputs` keep tracking of the number of chars it put to the console.

Diving down to the `cons_putc` function, one must notice the `local_intr_save` and `local_intr_restore`, these operations first disables interrupt delivery on the current processor after saving the current interrupt state into `intr_flag` and later restores that state which was stored into `intr_flag` by `local_intr_save` and ensures that during sbi call the system won't be trapped. Finally, the execution of `sbi_console_putchar` utilize the macro `SBI_CALL_1` which further dive into the `SBI_CALL` that consists of inline assembly and uses `ecall` to make a request to the supporting execution environment and `register` to pass data.

5. Implement the `shutdown` function.

As the `cputs` function is a wrapper of `cons_putc` defined in `console.h`, comparing the functionality that the `shutdown` function fits the `console` more than `stdio`, I choose to define the function and implement it in `console.h/c`

```
// kern/driver/console.h
void shutdown(void);

// kern/driver/console.c
void shutdown(void) {
    sbi_shutdown();
}
```

Similar to `cons_putc` that makes a sbi call, our `shutdown` also use a sbi call, that was already defined in `sbi.h`, but it's implementation was just added by me:

```
// libs/sbi.h
void sbi_shutdown(void);

// libs/sbi.c
uint64_t SBI_SHUTDOWN = 8; // this const was already here

void sbi_shutdown() {
    sbi_call(SBI_SHUTDOWN, 0, 0, 0);
}
```

Finally, we make the function call in `init.c`

```
// kern/init/init.c

int kern_init(void) {
    extern char edata[], end[];
    memset(edata, 0, end - edata);

    const char *message = "os is loading ...\n";
    cputs(message);

    // -----start-----
    cputs("The system will close.\n");
}
```

```
shutdown();  
// -----end-----  
  
while (1) ;  
}
```

And the result is as below:

```
riscv64-unknown-elf-objcopy bin/kernel --strip-all -O binary bin/ucore.bin
```

OpenSBI v0.6

```
      /---\          /-----\    \---|  
 |   |   | _--_     --- _--_ | (----| |) || |  
 |   |   | '-\' / _-' \'-' \--- \| < | |  
 |   |   | | ) | _-/ | | |----) | | ) | | |_  
 \---:/ | .../ \---|| | |_____/|___/_---|  
       | |  
       |_|
```

Platform Name : QEMU Virt Machine
Platform HART Features : RV64ACDFIMSU
Platform Max HARTs : 8
Current Hart : 0
Firmware Base : 0x80000000
Firmware Size : 120 KB
Runtime SBI Version : 0.2

MIDELEG : 0x00000000000000222
MEDELEG : 0x0000000000000b109
PMP0 : 0x0000000080000000-0x000000008001ffff (A)
PMP1 : 0x0000000000000000-0xfffffffffffffffff (A,R,W,X)

os is loading ...

The system will close.

```
heza12011323@VM-8-14-ubuntu:~/assign2$
```

```
OpenSBI v0.6
```

```
Platform Name      : QEMU Virt Machine
Platform HART Features : RV64ACDFIMSU
Platform Max HARTs  : 8
Current Hart       : 0
Firmware Base      : 0x80000000
Firmware Size      : 120 KB
Runtime SBI Version : 0.2
```

```
MIDELEG : 0x0000000000000222
MEDELEG : 0x000000000000b109
PMP0    : 0x0000000080000000-0x000000008001ffff (A)
PMP1    : 0x0000000000000000-0xffffffffffffffff (A,R,W,X)
os is loading ...
```

```
The system will close.
```

```
heza12011323@VM-8-14-ubuntu:~/assign2$
```

Required screenshots:

```
int kern_init(void) {
    extern char edata[], end[];
    memset(edata, 0, end - edata);

    const char *message = "os is loading ...\n";
    cputs(message);

    // clock_init();
    // -----start-----

    ...cputs("The system will close.\n");
    ...shutdown();

    // -----end-----

    while (1)
        ;
}
```

C console.h ×

kern > driver > C console.h > shutdown(void)

```
1  ✓ #ifndef __KERN_DRIVER_CONSOLE_H__
2    #define __KERN_DRIVER_CONSOLE_H__
3
4    void cons_init(void);
5    void cons_putc(int c);
6    int cons_getc(void);
7    void serial_intr(void);
8    void kbd_intr(void);
9
10   void shutdown(void); // used by assign2
11
12   #endif /* !__KERN_DRIVER_CONSOLE_H__ */
```


~/Downloads/assign2/kern/driver/intr.h

kern > driver > C console.h > shutdown(void)

```
1  #ifndef __KERN_DRIVER_CONSOLE_H__
2  #define __KERN_DRIVER_CONSOLE_H__
3
4  void cons_init(void);
5  void cons_putc(int c);
6  int cons_getc(void);
7  void serial_intr(void);
8  void kbd_intr(void);
9
10 void shutdown(void); // used by assign2
11
12 #endif /* !__KERN_DRIVER_CONSOLE_H__ */
13
14
```

~/Downloads/assign2/libs

sbi.h

```
libs > C sbi.h > sbi_shutdown(void)
1  #ifndef _ASM_RISCV_SBI_H
2  #define _ASM_RISCV_SBI_H
3
4  typedef struct {
5      unsigned long base;
6      unsigned long size;
7      unsigned long node_id;
8  } memory_block_info;
9
10 unsigned long sbi_query_memory(unsigned long id, memon
11
12 unsigned long sbi_hart_id(void);
13 unsigned long sbi_num_harts(void);
14 unsigned long sbi_timebase(void);
15 void sbi_set_timer(unsigned long long stime_value);
16 void sbi_send_ipi(unsigned long hart_id);
17 unsigned long sbi_clear_ipi(void);
18 void sbi_shutdown(void); // used by assign2
19
```

← →

assign2

C sbi.c

```
libs > C sbi.c > sbi_shutdown()
1  #include <sbi.h>
2  #include <defs.h>
3
4
5  uint64_t SBI_SET_TIMER = 0;
6  uint64_t SBI_CONSOLE_PUTCHAR = 1;
7  uint64_t SBI_CONSOLE_GETCHAR = 2;
8  uint64_t SBI_CLEAR_TDT = 3;
9  uint64_t SBI_REMOTE_FENCE_I
10 uint64_t SBI_REMOTE_FENCE_I = 5;
11 uint64_t SBI_REMOTE_SFENCE_VMA = 6;
12 uint64_t SBI_REMOTE_SFENCE_VMA_ASID = 7;
```

```

13  uint64_t SBI_SHUTDOWN = 8;
14
15  uint64_t sbi_call(uint64_t sbi_type, uint64_t arg0, uint64_t arg1,
16      uint64_t ret_val;
17      __asm__ volatile (
18          "mv x17, %[sbi_type]\n"
19          "mv x10, %[arg0]\n"
20          "mv x11, %[arg1]\n"
21          "mv x12, %[arg2]\n"
22          "ecall\n"
23          "mv %[ret_val], x10"
24          : [ret_val] "=r" (ret_val)
25          : [sbi_type] "r" (sbi_type), [arg0] "r" (arg0), [arg1] "r"
26          : "memory"
27      );
28      return ret_val;
29  }
30
31  int sbi_console_getchar(void) {
32      return sbi_call(SBI_CONSOLE_GETCHAR, 0, 0, 0);
33  }
34  void sbi_console_putchar(unsigned char ch) {
35      sbi_call(SBI_CONSOLE_PUTCHAR, ch, 0, 0);
36  }
37
38  void sbi_set_timer(unsigned long long stime_value) {
39      sbi_call(SBI_SET_TIMER, stime_value, 0, 0);
40  }
41
42  // assign2
43  void sbi_shutdown() {
44      sbi_call(SBI_SHUTDOWN, 0, 0, 0);
45  }

```